

# Machine Learning

CS 539

Worcester Polytechnic Institute

Department of Computer Science

Instructor: Prof. Kyumin Lee


# HW1

- <https://canvas.wpi.edu/courses/58900/assignments/355140>
- Due date is June 6th 11:59pm.


# Training Data and Test Data

- Training data: data used to build the model
- Test data: new data, not used in the training process
- Training performance is often a poor indicator of generalization performance
  - Generalization is what we really care about in ML
  - Easy to overfit the training data
  - Performance on test data is a good indicator of generalization performance
  - i.e., test accuracy is more important than training accuracy


# Simple Decision Boundary




# More Complex Decision Boundary




# Example: The Overfitting Phenomenon



# A Complex Model




# The True (simpler) Model






# How Overfitting Affects Prediction



# How Overfitting Affects Prediction




# How Overfitting Affects Prediction



# Comparing Classifiers

- Say we have two classifiers,  $C1$  and  $C2$ , and want to choose the best one to use for future predictions
- Can we use training accuracy to choose between them?
- No!
- Instead, choose based on test accuracy...

# Training and Test Data



## **Idea:**

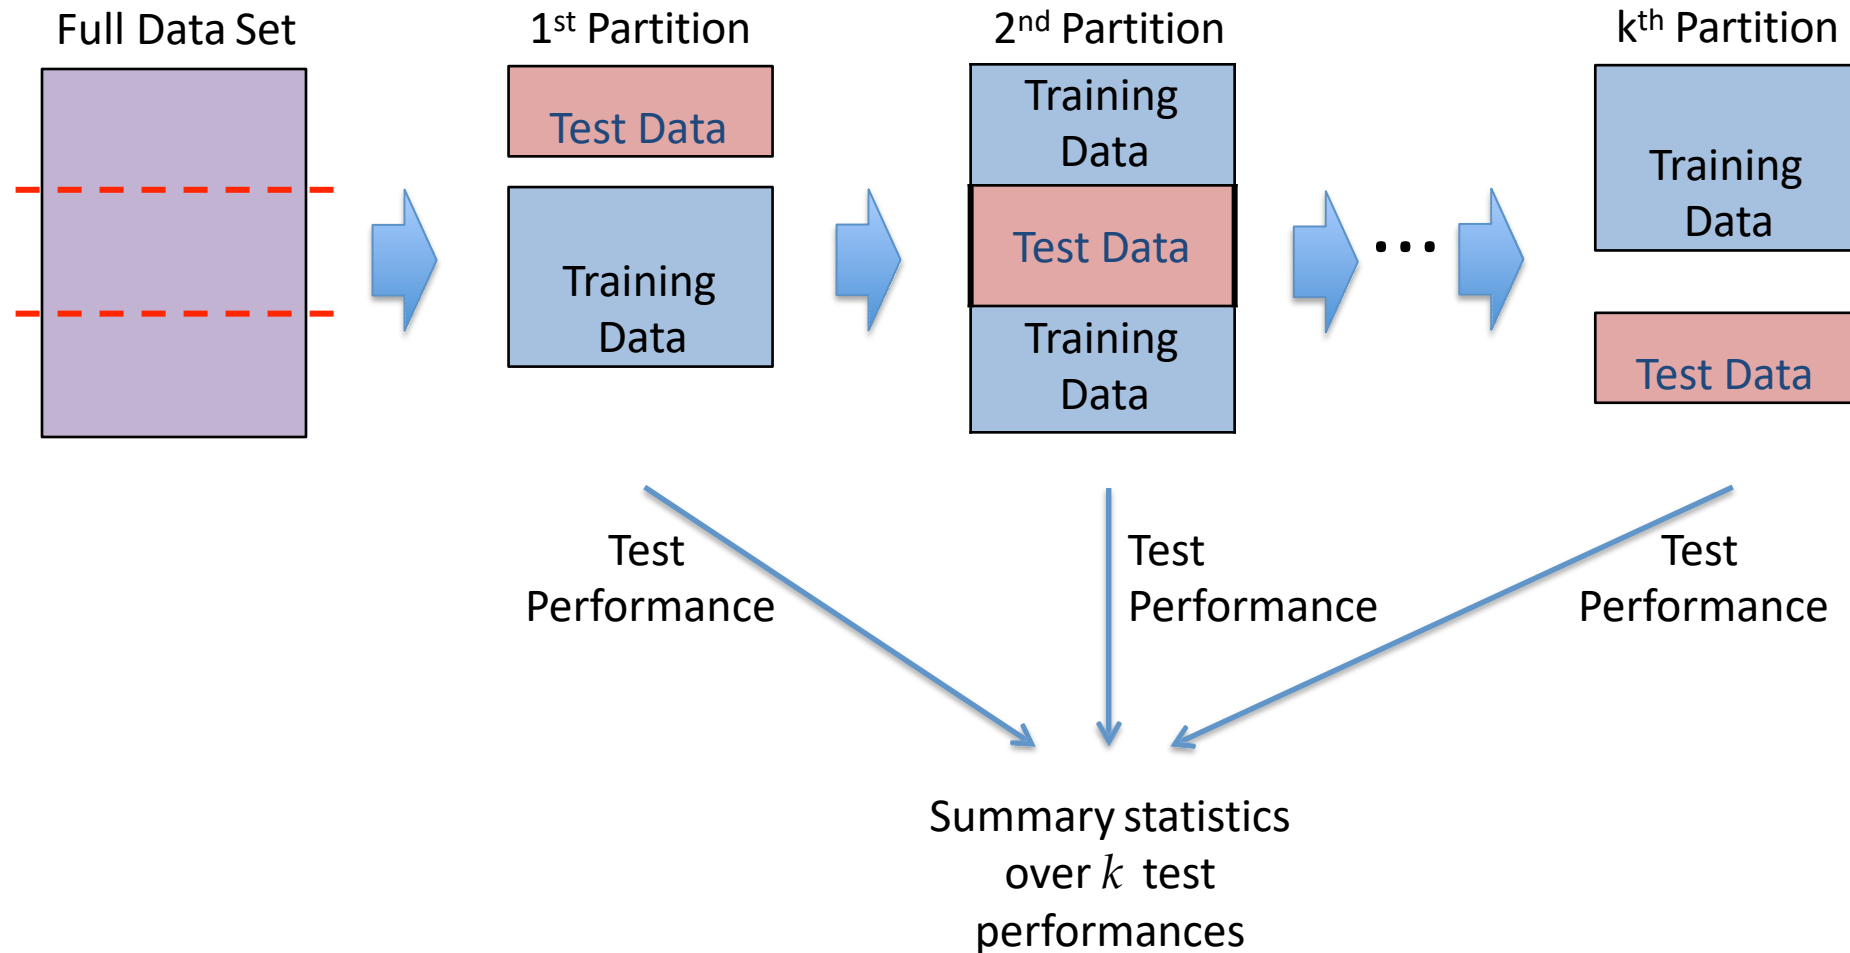
Train each model on the “training data” ...

...and then test each model's accuracy on the test data

# $k$ -Fold Cross-Validation

- Why just choose one particular “split” of the data?
  - In principle, we should do this multiple times since performance may be different for each split
- $k$ -Fold Cross-Validation (e.g.,  $k=10$ )
  - randomly partition full data set of  $n$  instances into  $k$  disjoint subsets (each roughly of size  $n/k$ )
  - Choose each fold in turn as the test set; train model on the other folds and evaluate
  - Compute statistics over  $k$  test performances, or choose best of the  $k$  models

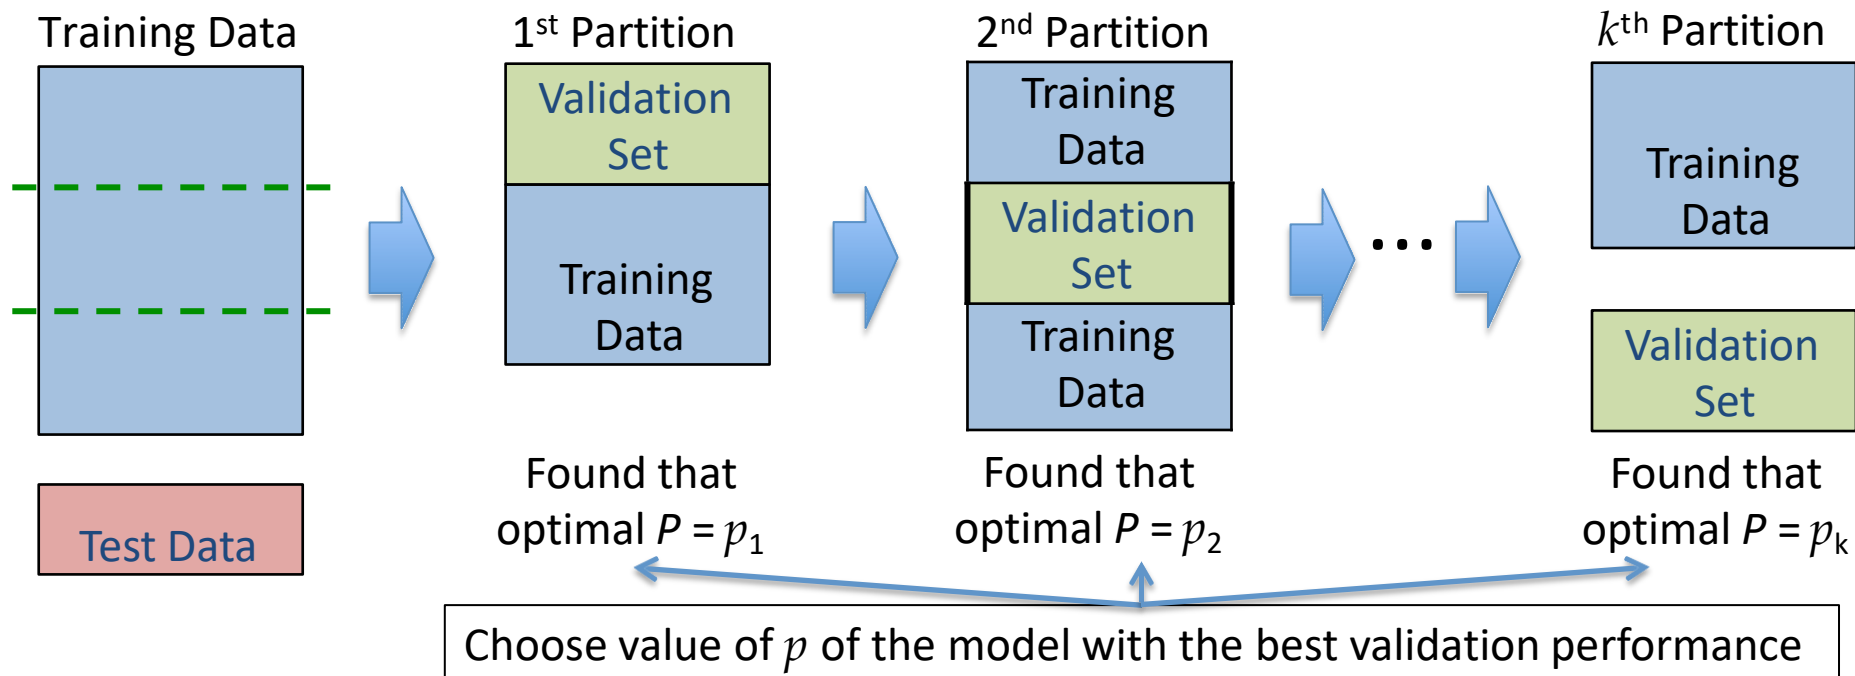
# Example 3-Fold CV



# Optimizing Model Parameters

Can also use CV to choose value of model parameter  $P$

- Search over space of parameter values  $p \in \text{values}(P)$ 
  - Evaluate model with  $P = p$  on validation set
- Choose value  $p'$  with highest validation performance
- Learn model on full training set with  $P = p'$






# $k$ -Nearest Neighbor & Instance-based Learning


# 1-Nearest Neighbor

- One of the simplest of all machine learning classifiers
- Simple idea: label a new point the same as the closest known point




# 1-Nearest Neighbor

- A type of instance-based learning
  - Also known as “memory-based” learning
- Forms a Voronoi tessellation of the instance space




# Distance Metrics

- Different metrics can change the decision surface



$$\text{Dist}(\mathbf{a}, \mathbf{b}) = (a_1 - b_1)^2 + (a_2 - b_2)^2$$



$$\text{Dist}(\mathbf{a}, \mathbf{b}) = (a_1 - b_1)^2 + (3a_2 - 3b_2)^2$$

- Standard Euclidean distance metric:
  - Two-dimensional:  $\text{Dist}(\mathbf{a}, \mathbf{b}) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2}$
  - Multivariate:  $\text{Dist}(\mathbf{a}, \mathbf{b}) = \sqrt{\sum (a_i - b_i)^2}$

## Four Aspects of an Instance-Based Learner:

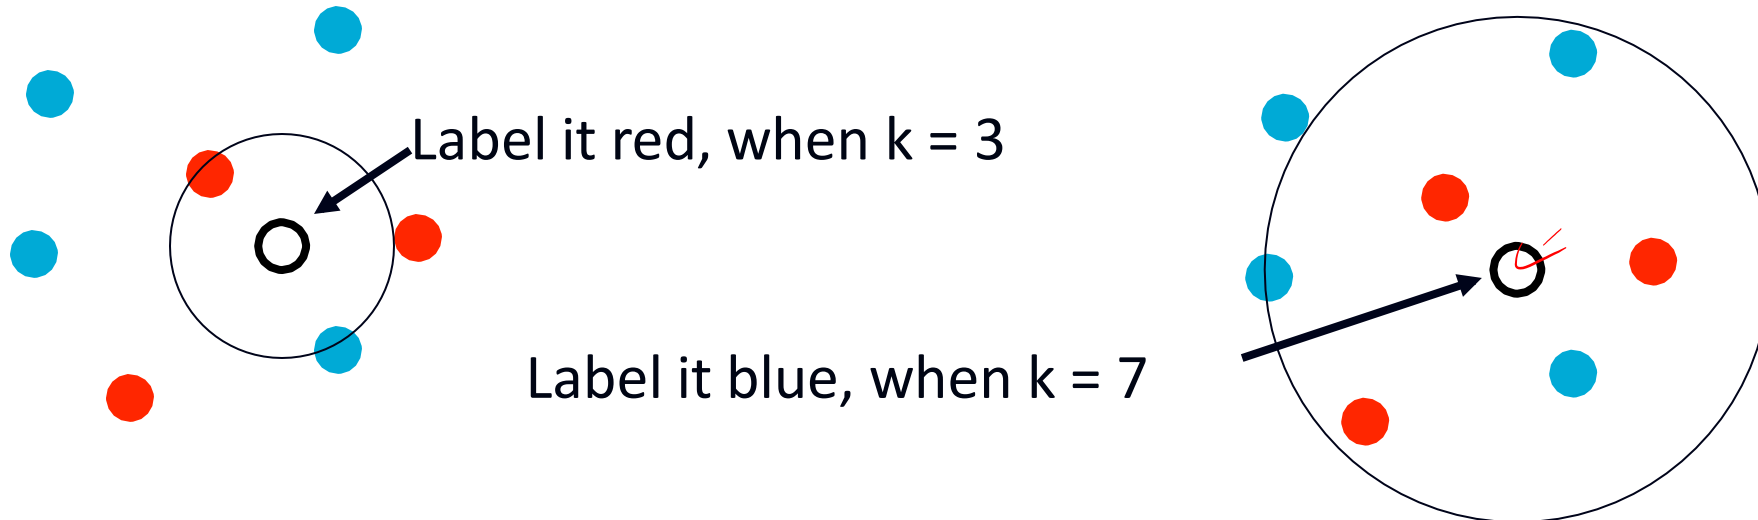
1. A distance metric
2. How many nearby neighbors to look at?
3. A weighting function (optional)
4. How to fit with the local points?

# 1-NN's Four Aspects as an Instance-Based Learner:

1. A distance metric
  - *Euclidean*
2. How many nearby neighbors to look at?
  - *One*
3. A weighting function (optional)
  - *Unused*
4. How to fit with the local points?
  - *Just predict the same output as the nearest neighbor.*

# $k$ – Nearest Neighbor

- Generalizes 1-NN to smooth away noise in the labels
- A new point is now assigned the most frequent label of its  $k$  nearest neighbors



# k-NN

- Instance-based learning & lazy learning
- Memorize training data, and measure all each pair of instance in the training set and new instance in the test set
- However, computationally expensive
  - Require  $N$  comparison for the prediction
  - Refer to <https://machinelearningmastery.com/tutorial-to-implement-k-nearest-neighbors-in-python-from-scratch/>
- To reduce the search time (i.e., reduce  $O(n)$ )
  - We may use some data structure
  - e.g., k-d tree (k- dimensional tree)
    - [https://en.wikipedia.org/wiki/K-d\\_tree](https://en.wikipedia.org/wiki/K-d_tree)
- k-NN regression
  - k nearest neighbors' average target class value
  - [http://www.saedsayad.com/k\\_nearest\\_neighbors\\_reg.htm](http://www.saedsayad.com/k_nearest_neighbors_reg.htm)
  - <http://scikit-learn.org/stable/modules/neighbors.html>




# Quiz1

- Quiz1 will be taken on June 4 – it will be available only during the day
- The coverage will be from content of the first lecture to K-nearest neighbors

# Linear Regression

# Regression

- Given:
  - Data  $X = \{x^{(1)}, \dots, x^{(n)}\}$  where  $x^{(i)} \in \mathbb{R}^d$
  - Corresponding labels  $y = \{y^{(1)}, \dots, y^{(n)}\}$  where  $y^{(i)} \in \mathbb{R}$




# Linear Regression

- Hypothesis:

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d = \sum_{j=0}^d \theta_j x_j$$

Assume  $x_0 = 1$

- Fit model by minimizing sum of squared errors





# Least Squares Linear Regression

- Cost Function

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Fit by solving  $\min_{\theta} J(\theta)$



# Intuition Behind Cost Function



$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

For insight on  $J()$ , let's assume  $x \in \mathbb{R}$  so  $\theta = [\theta_0, \theta_1]$

# Intuition Behind Cost Function

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

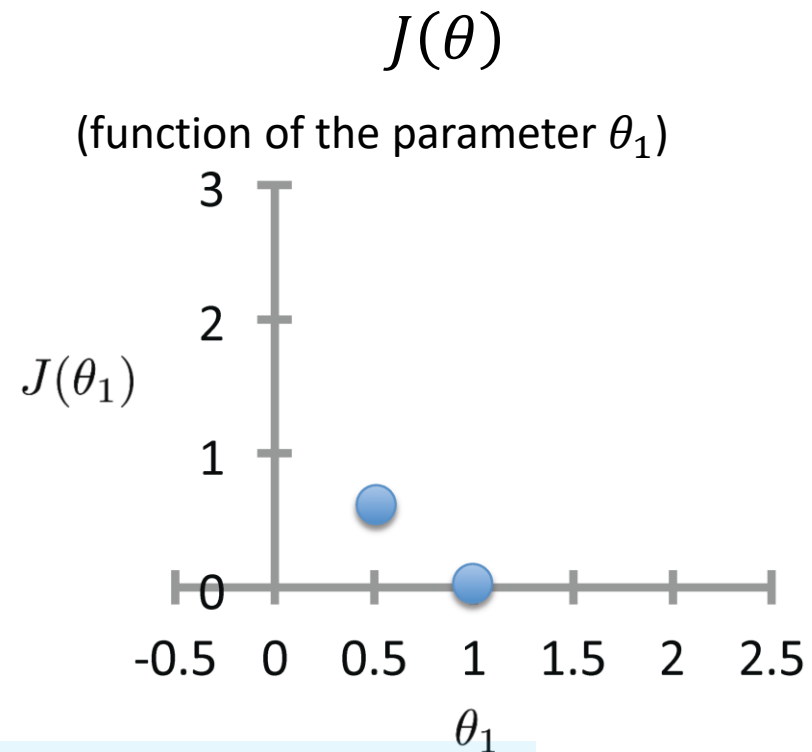

For insight on  $J()$ , let's assume  $x \in \mathbb{R}$  so  $\theta = [\theta_0, \theta_1] \rightarrow \theta_0 = 0$



# Intuition Behind Cost Function

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

For insight on  $J()$ , let's assume  $x \in \mathbb{R}$  so  $\theta = [\theta_0, \theta_1] \rightarrow \theta_0 = 0$





$$J([0, 0.5]) = \frac{1}{2 \times 3} [(0.5 - 1)^2 + (1 - 2)^2 + (1.5 - 3)^2] \approx 0.58$$



# Intuition Behind Cost Function

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2$$


For insight on  $J()$ , let's assume  $x \in \mathbb{R}$  so  $\theta = [\theta_0, \theta_1]$



<http://mathworld.wolfram.com/ConvexFunction.html>

<https://www.desmos.com/calculator/kreo2ssqj8>


# Intuition Behind Cost Function (3-D surface plot)



# Intuition Behind Cost Function


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$

(function of the parameter  $\theta_0, \theta_1$ )




Contour Figure

# Intuition Behind Cost Function


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameter  $\theta_0, \theta_1$ )



# Intuition Behind Cost Function


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameter  $\theta_0, \theta_1$ )



# Intuition Behind Cost Function


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameter  $\theta_0, \theta_1$ )




# Basic Search Procedure

- Choose initial value for  $\theta$
- Until we reach a minimum:
  - Choose a new value for  $\theta$  to reduce  $J(\theta)$



# Basic Search Procedure


- Choose initial value for  $\theta$
- Until we reach a minimum:
  - Choose a new value for  $\theta$  to reduce  $J(\theta)$





# Basic Search Procedure

- Choose initial value for  $\theta$
- Until we reach a minimum:
  - Choose a new value for  $\theta$  to reduce  $J(\theta)$



Since the least squares objective function is convex, we don't need to worry about local minima