

Machine Learning

CS 539

Worcester Polytechnic Institute
Department of Computer Science
Instructor: Prof. Kyumin Lee

Upcoming Schedule

- HW 3
 - <https://canvas.wpi.edu/courses/58900/assignments/356655>
 - Due date is July 2

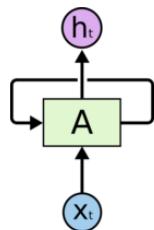
Project Website & Presentation

Item		4	3	2	1	0
Effective use of visual aids	Clear and interesting?					
Clarity of voice and pronunciation	Clear and loud? Good use of vocabulary?					
Completeness	Did they complete what they proposed?					
Demo/Evaluation						
Total						

Recurrent Neural Network (RNN)

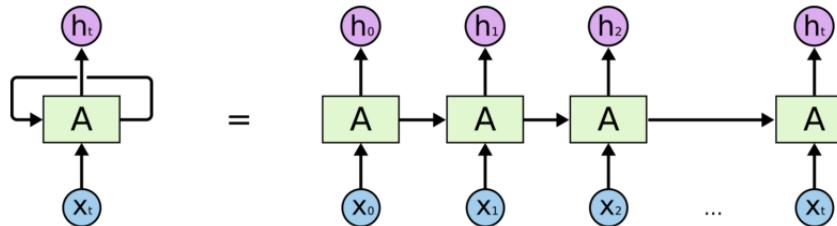
Recurrent Neural Networks

- Recurrent Neural Networks are networks with loops in them, allowing information to persist.



Recurrent Neural Networks have loops.

In the above diagram, a chunk of neural network, A , looks at some input x_t and outputs a value h_t . A loop allows information to be passed from one step of the network to the next.

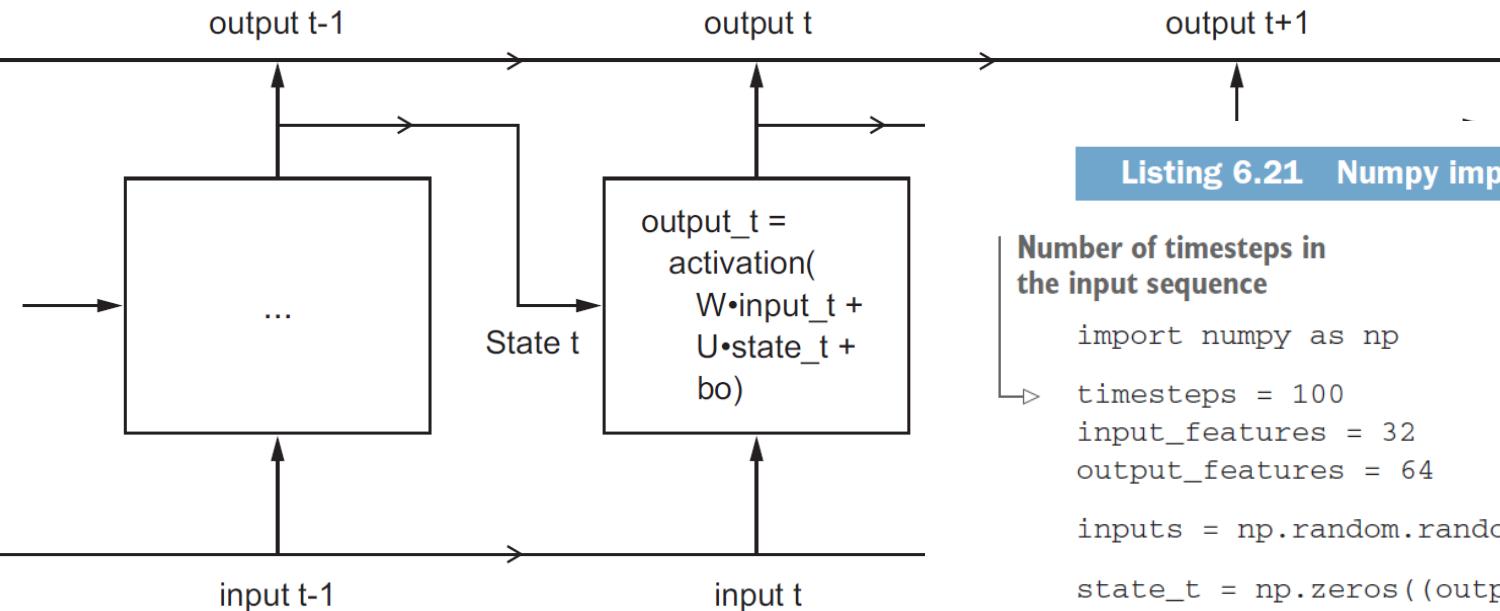


An unrolled recurrent neural network.

A recurrent neural network can be thought of as multiple copies of the same network, each passing a message to a successor. The diagram above shows what happens if we unroll the loop.

Recurrent Neural Networks

- Intuition of Recurrent Neural Networks
 - Human thoughts have persistence; humans don't start their thinking from scratch every second.
 - As you read this sentence, you understand each word based on your understanding of previous words.
 - One of the appeals of RNNs is the idea that they are able to connect previous information to the present task
 - E.g., using previous video frames to inform the understanding of the present frame.
 - E.g., a language model tries to predict the next word based on the previous ones.



Listing 6.21 Numpy implementation of a simple RNN

Number of timesteps in the input sequence

```
import numpy as np
timesteps = 100
input_features = 32
output_features = 64
```

Dimensionality of the input feature space

Dimensionality of the output feature space

Input data: random noise for the sake of the example

Initial state: an all-zero vector

Creates random weight matrices

```
inputs = np.random.random((timesteps, input_features))
state_t = np.zeros((output_features,))
```

```
W = np.random.random((output_features, input_features))
U = np.random.random((output_features, output_features))
b = np.random.random((output_features,))

successive_outputs = []
for input_t in inputs:
    output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)
    successive_outputs.append(output_t)
```

input_t is a vector of shape (input_features,).

```
state_t = output_t
```

```
final_output_sequence = np.concatenate(successive_outputs, axis=0)
```

The final output is a 2D tensor of shape (timesteps, output_features).

Stores this output in a list

Combines the input with the current state (the previous output) to obtain the current output

Updates the state of the network for the next timestep

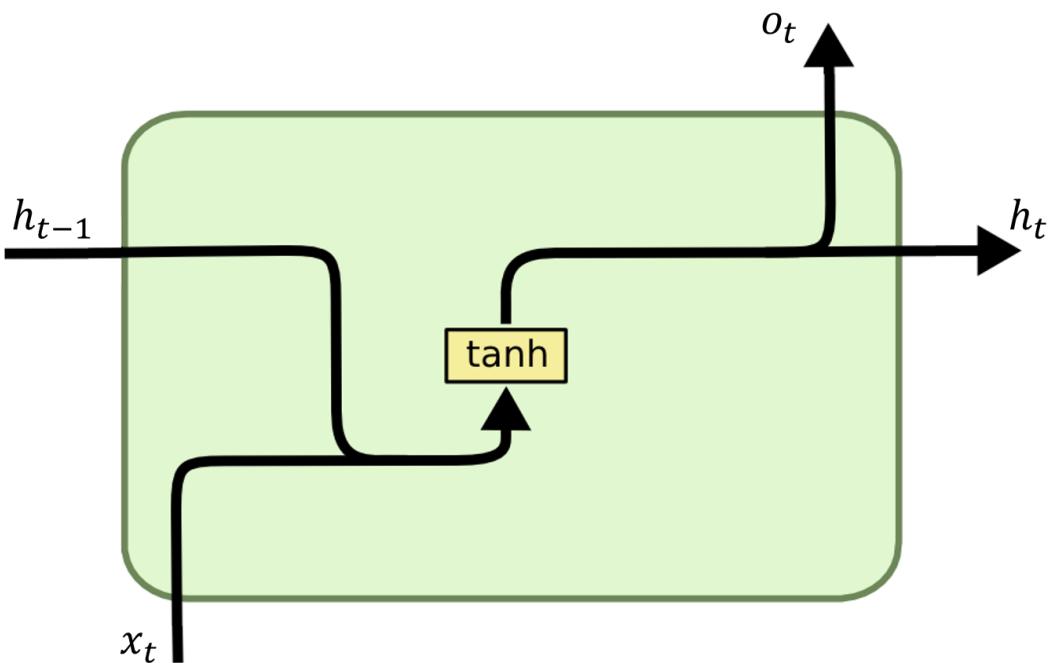
RNN/SimpleRNN in PyTorch

```
CLASS torch.nn.RNN(self, input_size, hidden_size, num_layers=1, nonlinearity='tanh', bias=True,  
batch_first=False, dropout=0.0, bidirectional=False, device=None, dtype=None) [SOURCE]
```



Apply a multi-layer Elman RNN with tanh or ReLU non-linearity to an input sequence. For each element in the input sequence, each layer computes the following function:

$$h_t = \tanh(x_t W_{ih}^T + b_{ih} + h_{t-1} W_{hh}^T + b_{hh})$$



```
class RNNTagger(torch.nn.Module):  
  
    def __init__(self, embedding_dim, hidden_dim, vocab_size, tagset_size):  
        super(RNNTagger, self).__init__()  
        self.hidden_dim = hidden_dim  
  
        self.word_embeddings = torch.nn.Embedding(vocab_size, embedding_dim)  
  
        # The RNN takes word embeddings as inputs, and outputs hidden states  
        # with dimensionality hidden_dim.  
        self.rnn = torch.nn.RNN(embedding_dim, hidden_dim)  
  
        # The linear layer that maps from hidden state space to tag space  
        self.hidden2tag = torch.nn.Linear(hidden_dim, tagset_size)  
  
    def forward(self, sentence):  
        embeds = self.word_embeddings(sentence)  
        rnn_out, _ = self.rnn(embeds.view(len(sentence), 1, -1))  
        tag_space = self.hidden2tag(rnn_out.view(len(sentence), -1))  
        tag_scores = F.log_softmax(tag_space, dim=1)  
        return tag_scores
```

Pytorch reference: <https://pytorch.org/docs/stable/generated/torch.nn.RNN.html>

RNN/SimpleRNN in PyTorch

```
class RNNTagger(torch.nn.Module):

    def __init__(self, embedding_dim, hidden_dim, vocab_size, tagset_size):
        super(RNNTagger, self).__init__()
        self.hidden_dim = hidden_dim

        self.word_embeddings = torch.nn.Embedding(vocab_size, embedding_dim)

        # The RNN takes word embeddings as inputs, and outputs hidden states
        # with dimensionality hidden_dim.
        self.rnn = torch.nn.RNN(embedding_dim, hidden_dim)

        # The linear layer that maps from hidden state space to tag space
        self.hidden2tag = torch.nn.Linear(hidden_dim, tagset_size)

    def forward(self, sentence):
        embeds = self.word_embeddings(sentence)
        rnn_out, _ = self.rnn(embeds.view(len(sentence), 1, -1))
        tag_space = self.hidden2tag(rnn_out.view(len(sentence), -1))
        tag_scores = F.log_softmax(tag_space, dim=1)
        return tag_scores
```

vocab_size: # of words/tokens in the vocabulary/dictionary

embedding_dim: embedding size of each word/token

torch.nn.Embedding:

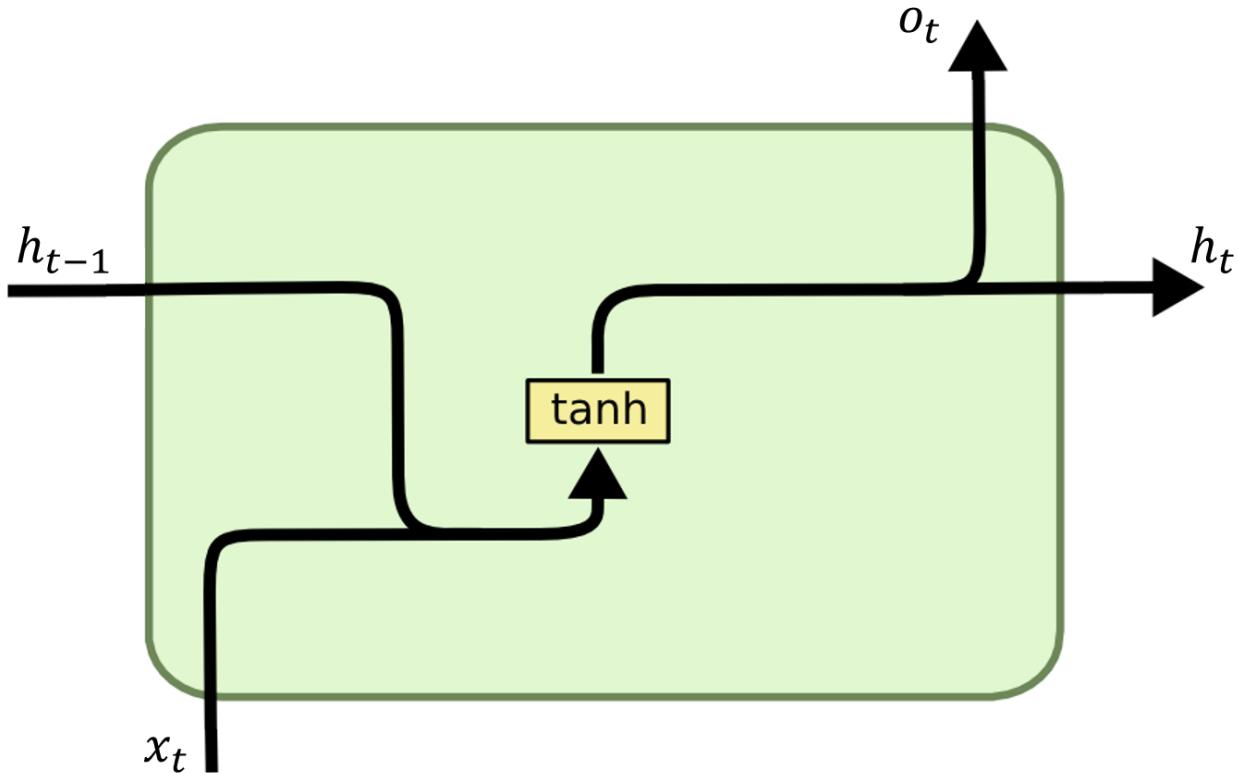
A simple lookup table that stores embeddings of a fixed dictionary and size.

This module is often used to store word embeddings and retrieve them using indices. The input to the module is a list of indices, and the output is the corresponding word embeddings.

Turns positive integers (indexes) into dense vectors of fixed size.

e.g. [4, 20] -> [[0.25, 0.1], [0.6, -0.2]]

RNN/SimpleRNN in Keras

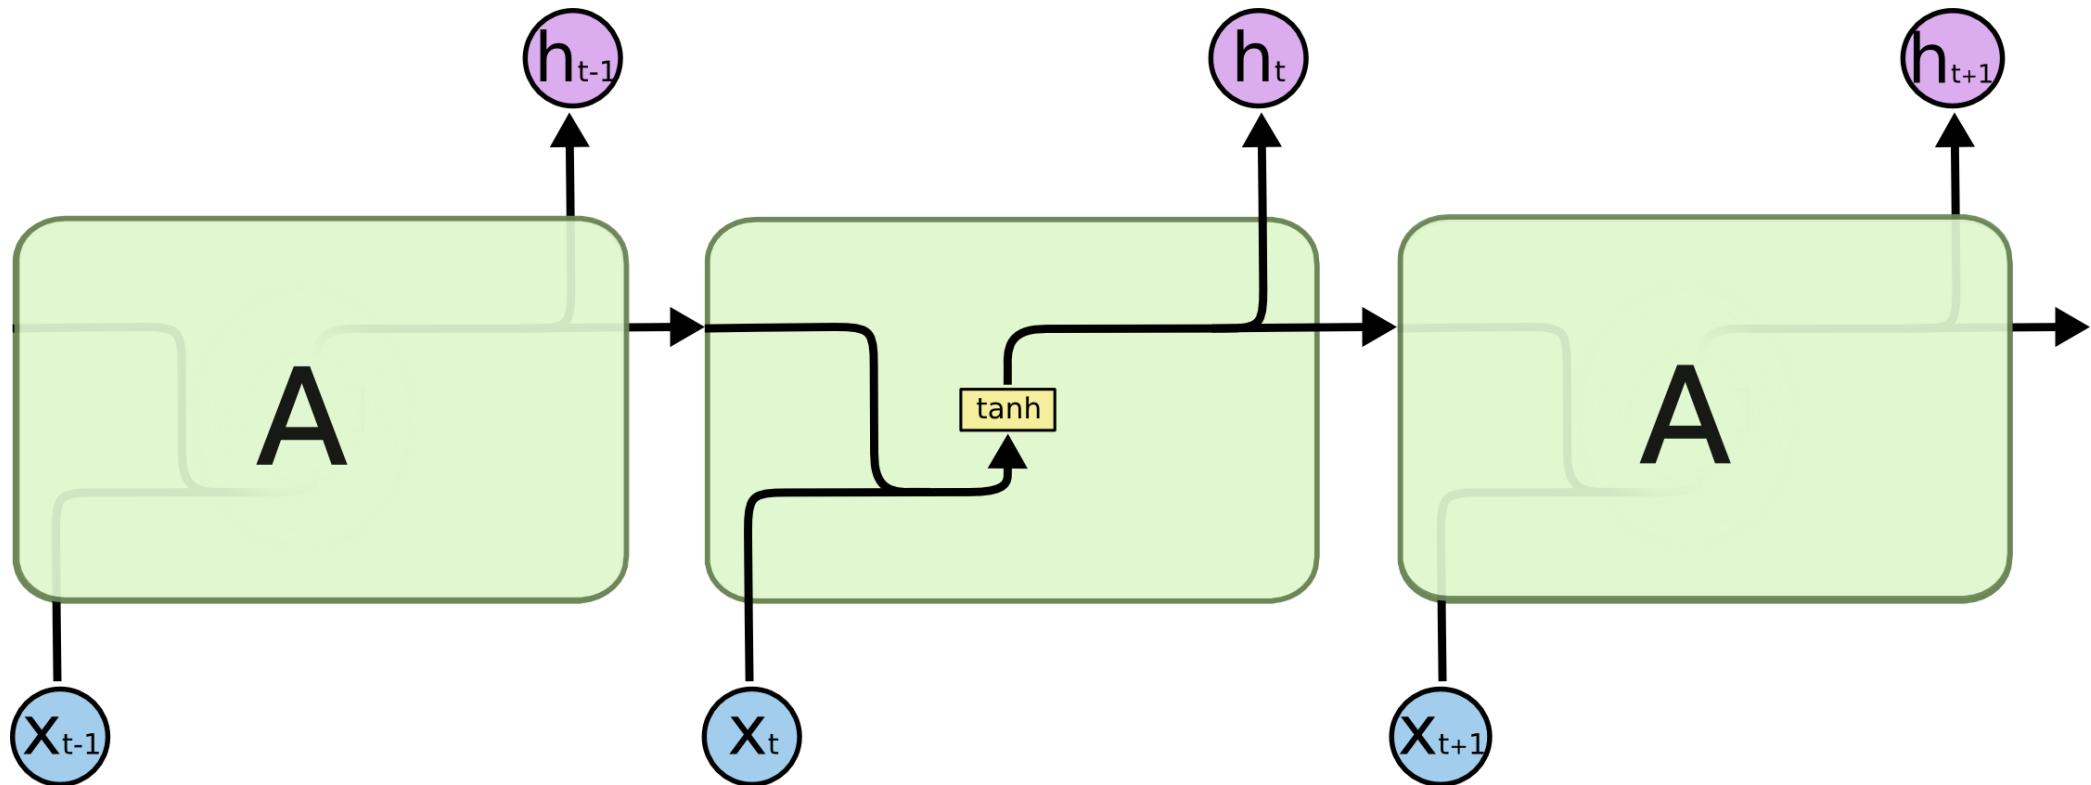


```
| from keras.layers import Dense  
|  
| max_features = 10000 in this example  
model = Sequential()  
model.add(Embedding(max_features, 32))  
model.add(SimpleRNN(32)) ← dimensionality of the output space  
model.add(Dense(1, activation='sigmoid'))  
model.summary()
```

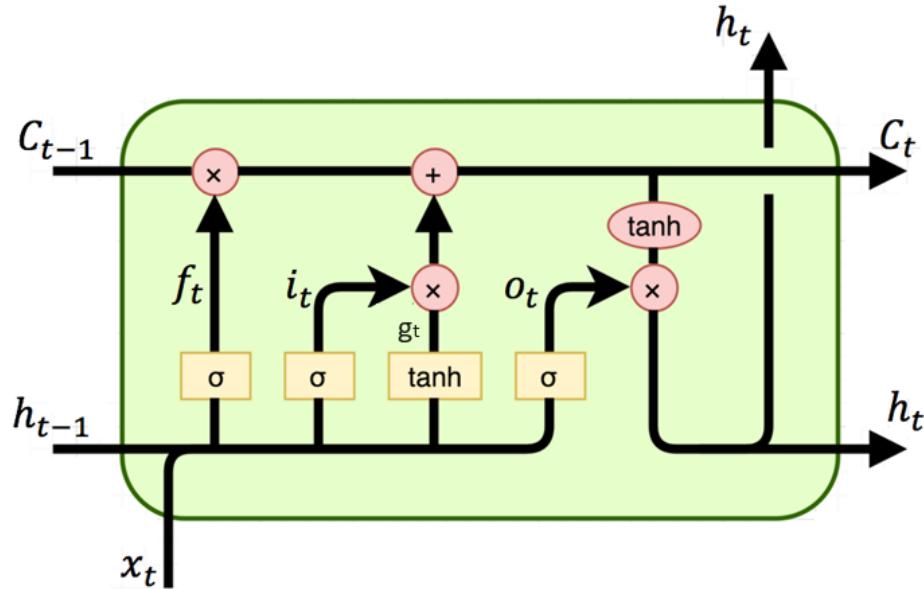
Model: "sequential_11"

Layer (type)	Output Shape	Param
=====		
embedding_11 (Embedding)	(None, None, 32)	320000
simple_rnn_12 (SimpleRNN)	(None, 32)	2080
dense_5 (Dense)	(None, 1)	33
=====		
Total params: 322,113		
Trainable params: 322,113		
Non-trainable params: 0		

RNN/SimpleRNN



Long-Short Term Memory (LSTM) in PyTorch



$$i_t = \sigma(W_{ii}x_t + b_{ii} + W_{hi}h_{t-1} + b_{hi})$$

$$f_t = \sigma(W_{if}x_t + b_{if} + W_{hf}h_{t-1} + b_{hf})$$

$$g_t = \tanh(W_{ig}x_t + b_{ig} + W_{hg}h_{t-1} + b_{hg})$$

$$o_t = \sigma(W_{io}x_t + b_{io} + W_{ho}h_{t-1} + b_{ho})$$

$$c_t = f_t \odot c_{t-1} + i_t \odot g_t$$

$$h_t = o_t \odot \tanh(c_t)$$

```
▶ class LSTMTagger(torch.nn.Module):

    def __init__(self, embedding_dim, hidden_dim, vocab_size, tagset_size):
        super(LSTMTagger, self).__init__()
        self.hidden_dim = hidden_dim

        self.word_embeddings = torch.nn.Embedding(vocab_size, embedding_dim)

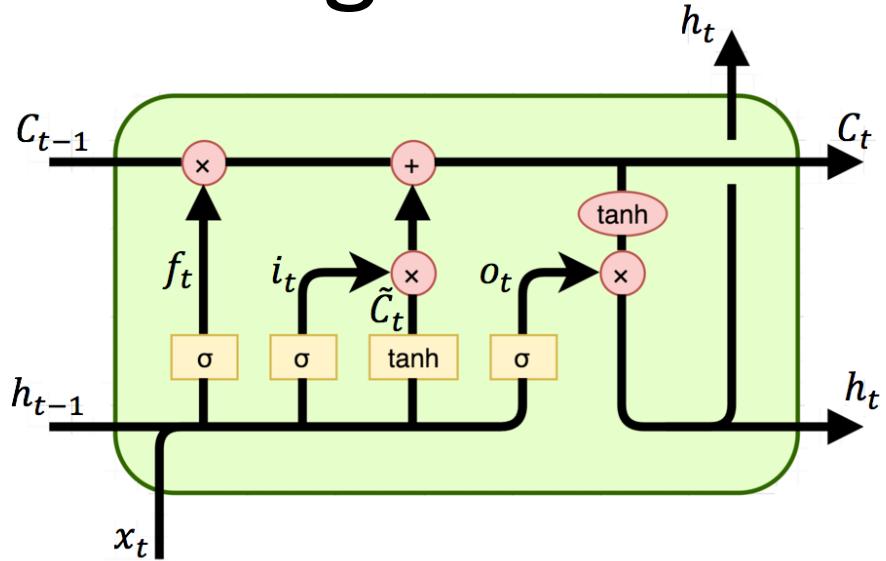
        # The LSTM takes word embeddings as inputs, and outputs hidden states
        # with dimensionality hidden_dim.
        self.lstm = torch.nn.LSTM(embedding_dim, hidden_dim)

        # The linear layer that maps from hidden state space to tag space
        self.hidden2tag = torch.nn.Linear(hidden_dim, tagset_size)

    def forward(self, sentence):
        embeds = self.word_embeddings(sentence)
        lstm_out, _ = self.lstm(embeds.view(len(sentence), 1, -1))
        tag_space = self.hidden2tag(lstm_out.view(len(sentence), -1))
        tag_scores = F.log_softmax(tag_space, dim=1)
        return tag_scores
```

PyTorch reference: <https://pytorch.org/docs/stable/generated/torch.nn.LSTM.html>
Example: https://pytorch.org/tutorials/beginner/nlp/sequence_models_tutorial.html

Long-Short Term Memory (LSTM) in Keras

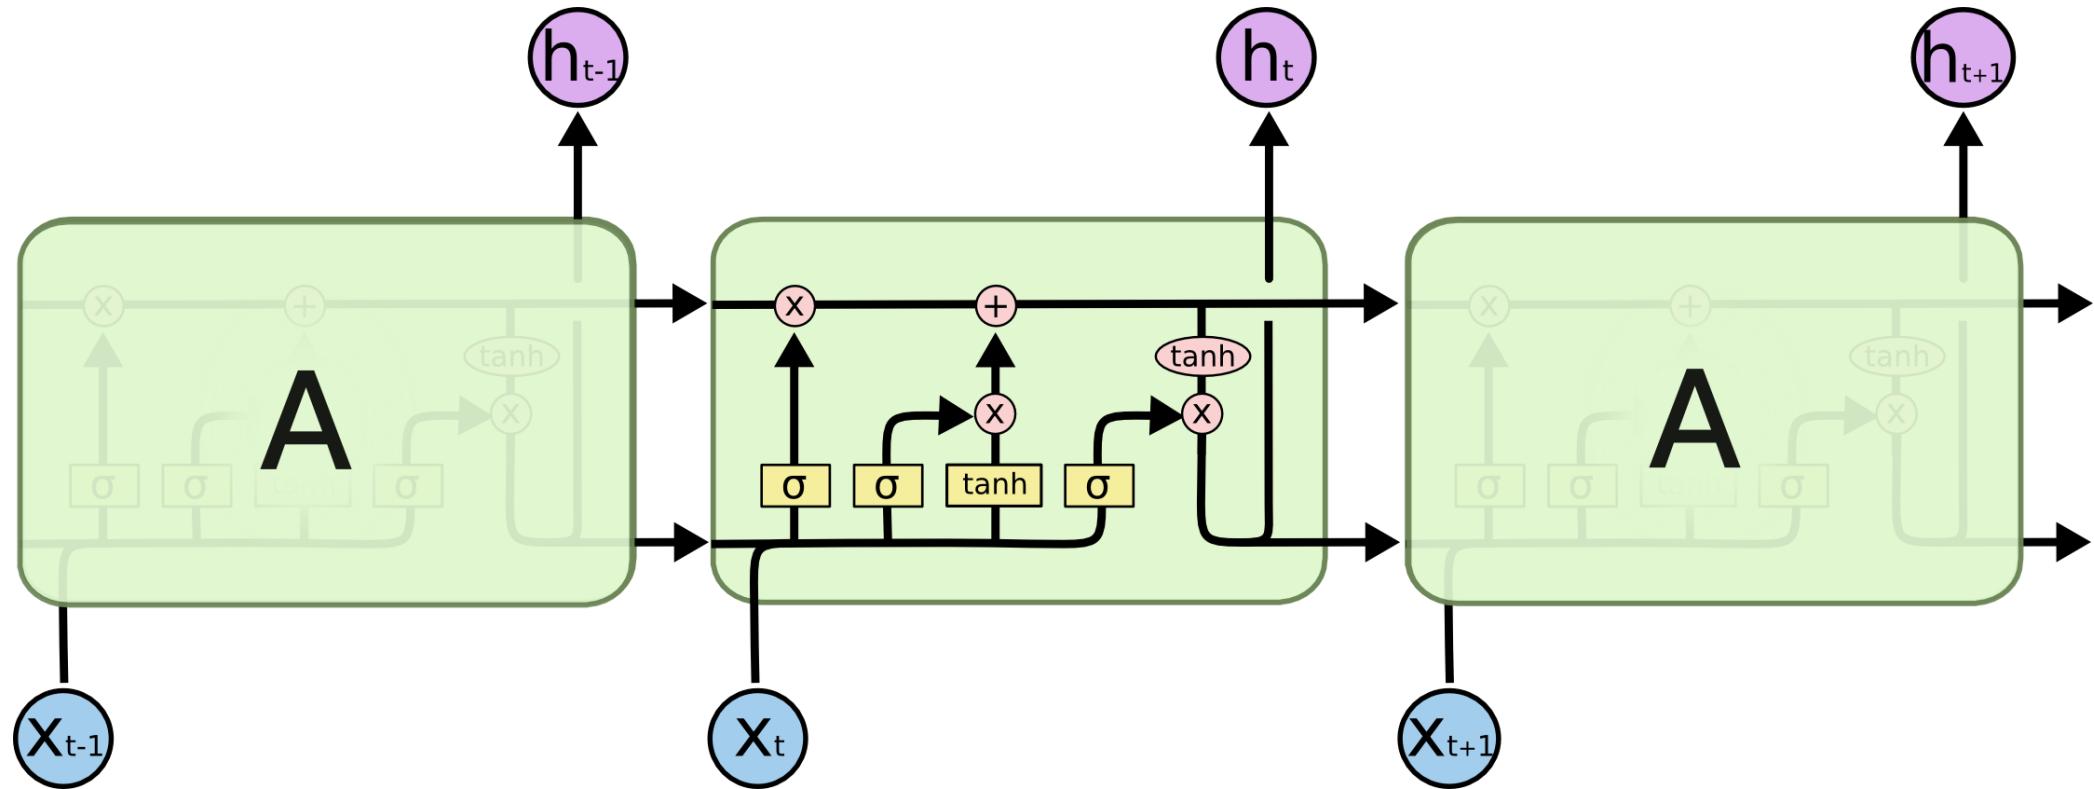


```
from keras.layers import LSTM  
  
model = Sequential()  
model.add(Embedding(max_features, 32))  
model.add(LSTM(32))  
model.add(Dense(1, activation='sigmoid'))  
model.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
=====		
embedding_13 (Embedding)	(None, None, 32)	320000
=====		
lstm_4 (LSTM)	(None, 32)	8320
=====		
dense_7 (Dense)	(None, 1)	33
=====		
Total params:	328,353	
Trainable params:	328,353	
Non-trainable params:	0	

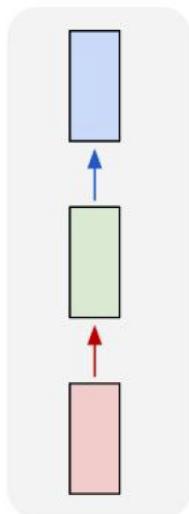
Long-Short Term Memory (LSTM)



Recurrent Neural Network

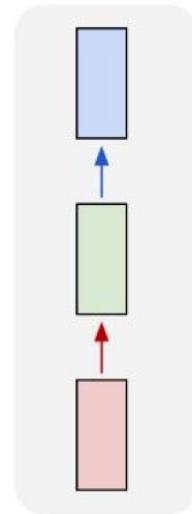
“Vanilla” Neural Network

one to one

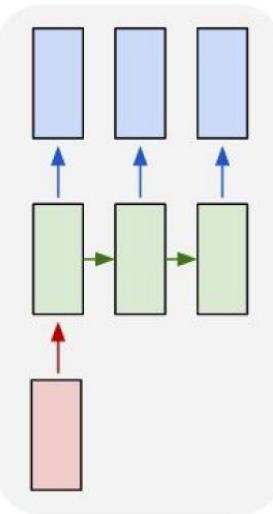


Recurrent Neural Networks: Process Sequences

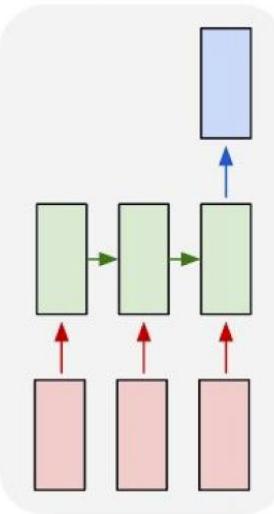
one to one



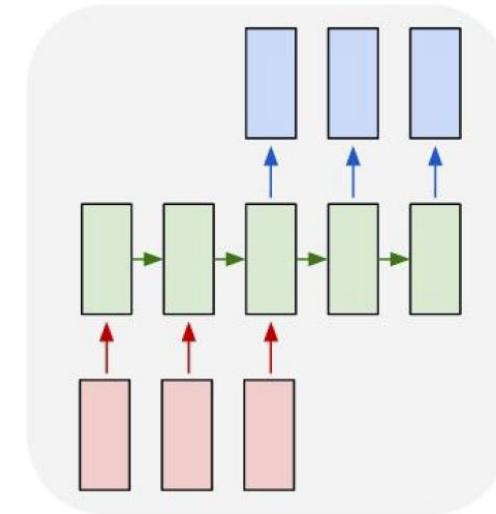
one to many



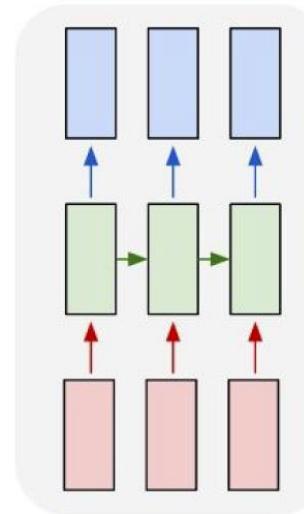
many to one



many to many



many to many



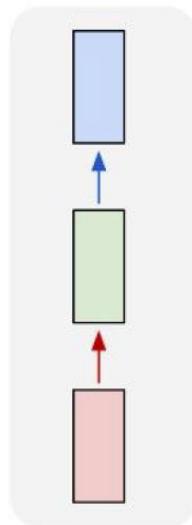
e.g. **Image Captioning**
image -> sequence of words



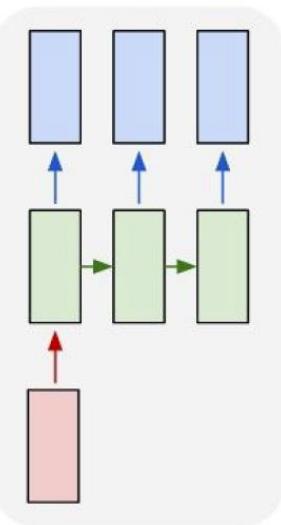
→ A person riding a motorbike on dirt road

Recurrent Neural Networks: Process Sequences

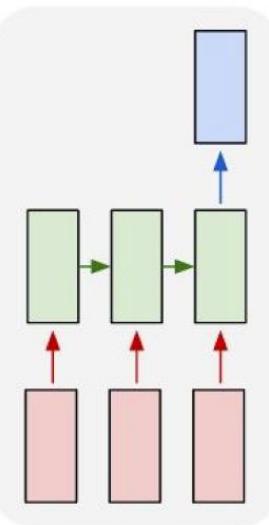
one to one



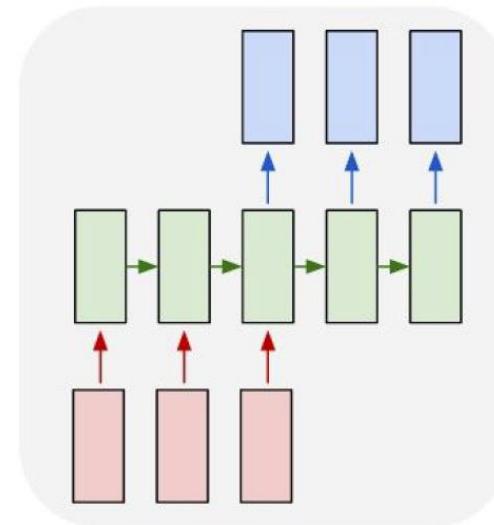
one to many



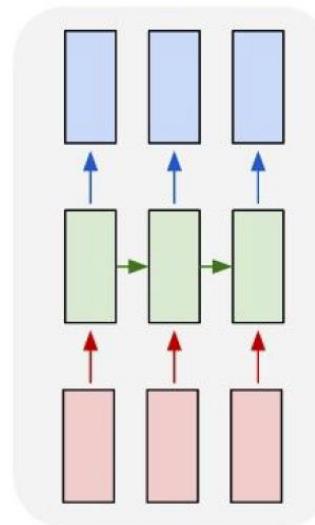
many to one



many to many



many to many

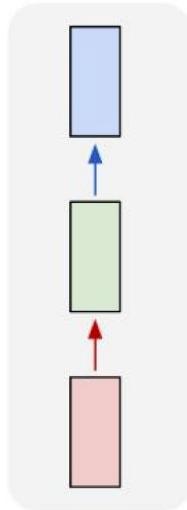


e.g. **Sentiment Classification**
sequence of words -> sentiment

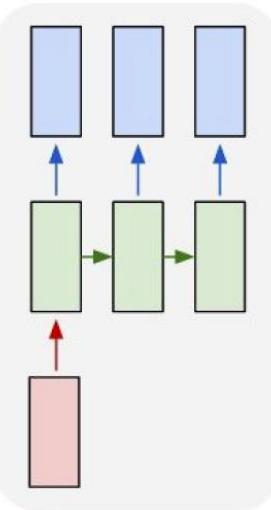
Awesome lecture! → Positive

Recurrent Neural Networks: Process Sequences

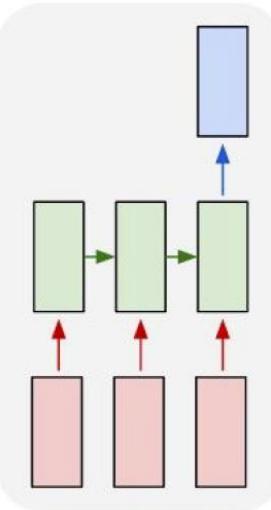
one to one



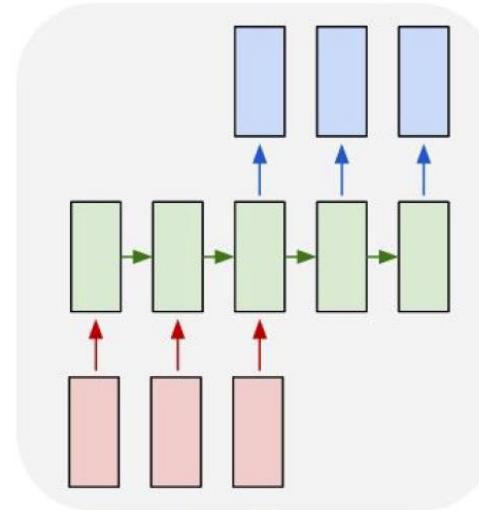
one to many



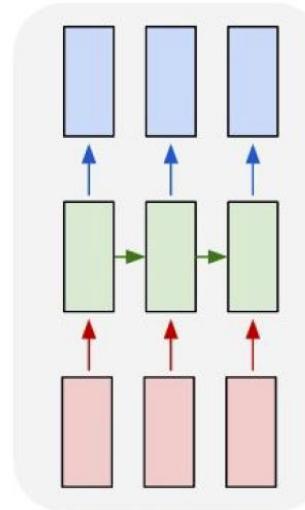
many to one



many to many



many to many

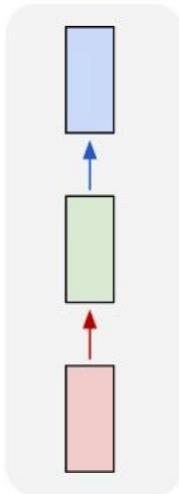


e.g. **Machine Translation**
seq of words -> seq of words

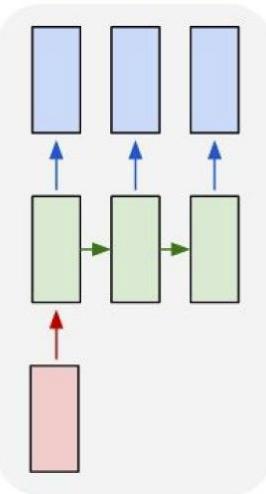
happy thanksgiving → 행복한 추수 감사절

Recurrent Neural Networks: Process Sequences

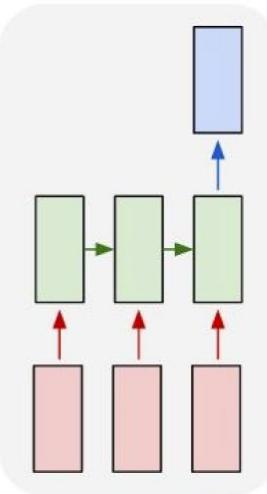
one to one



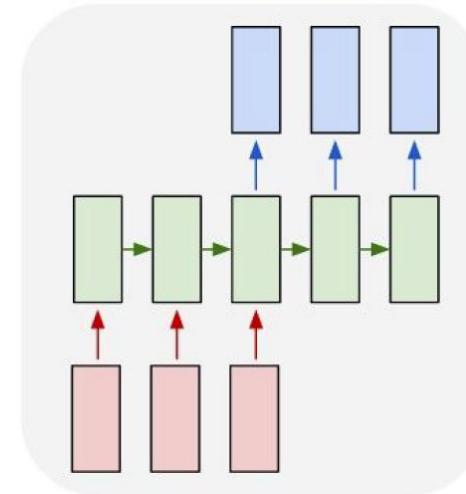
one to many



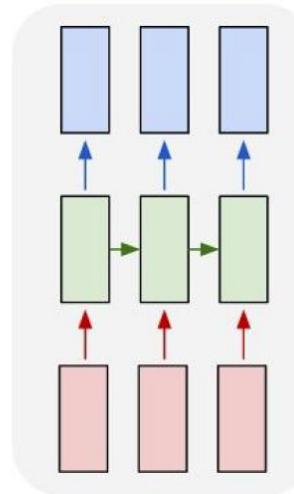
many to one



many to many

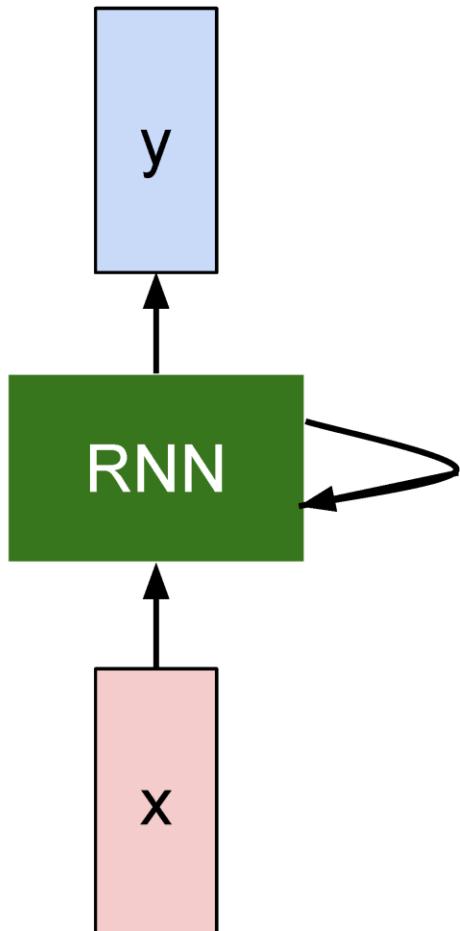


many to many

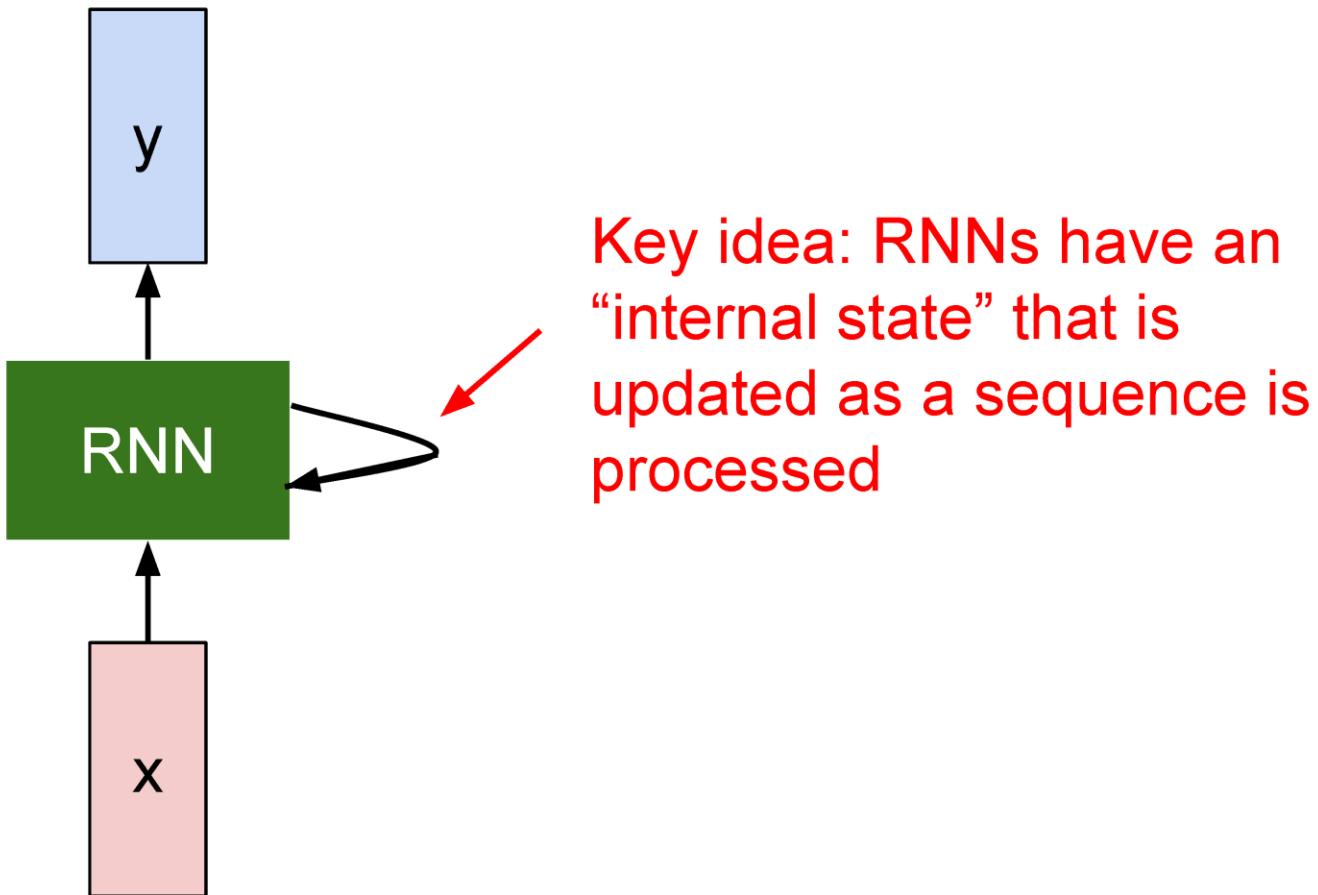


e.g. Video classification on frame level

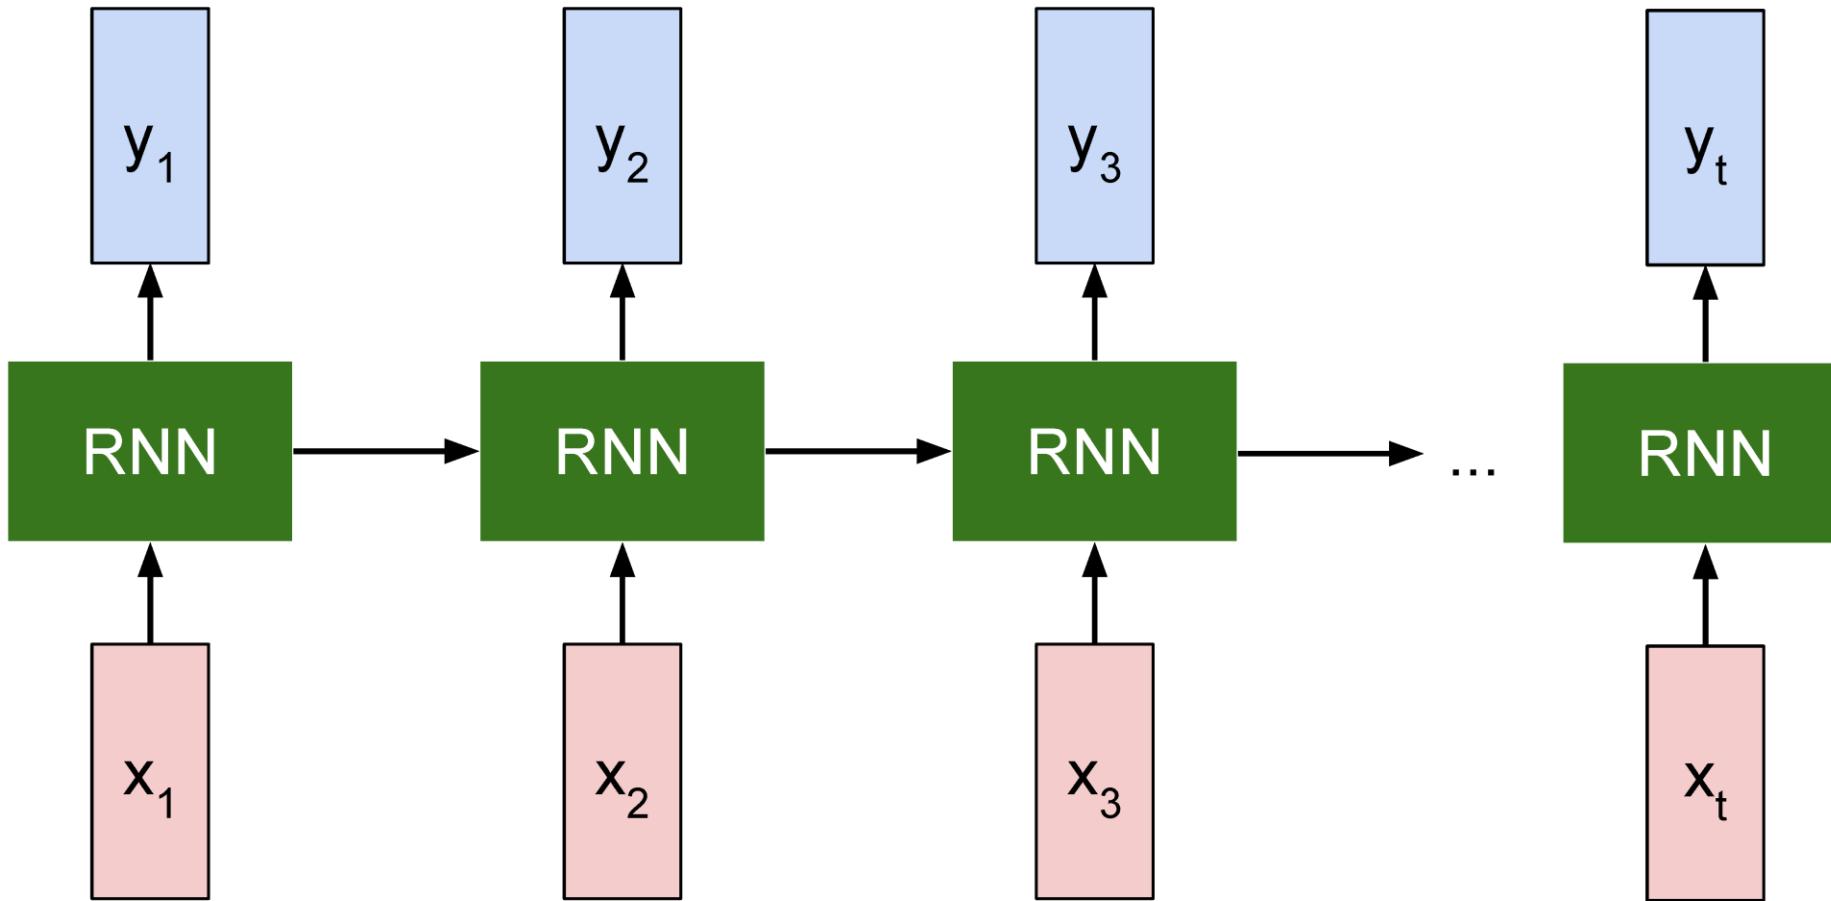
Recurrent Neural Network



Recurrent Neural Network



Unrolled RNN

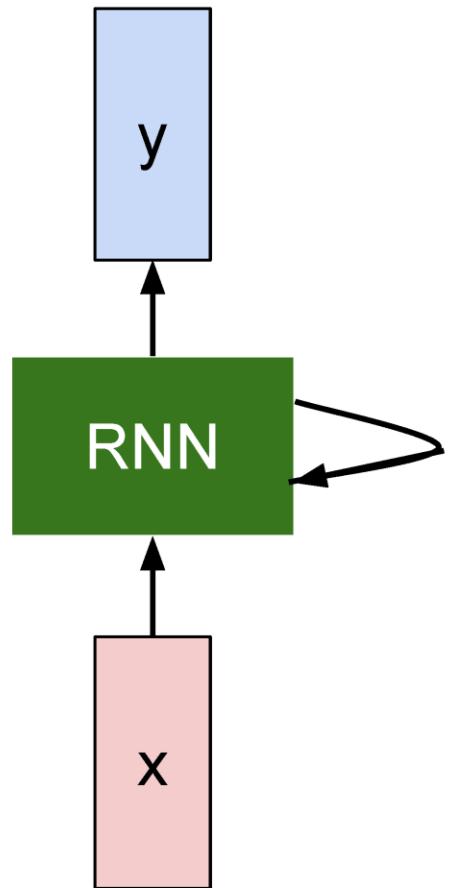


RNN hidden state update

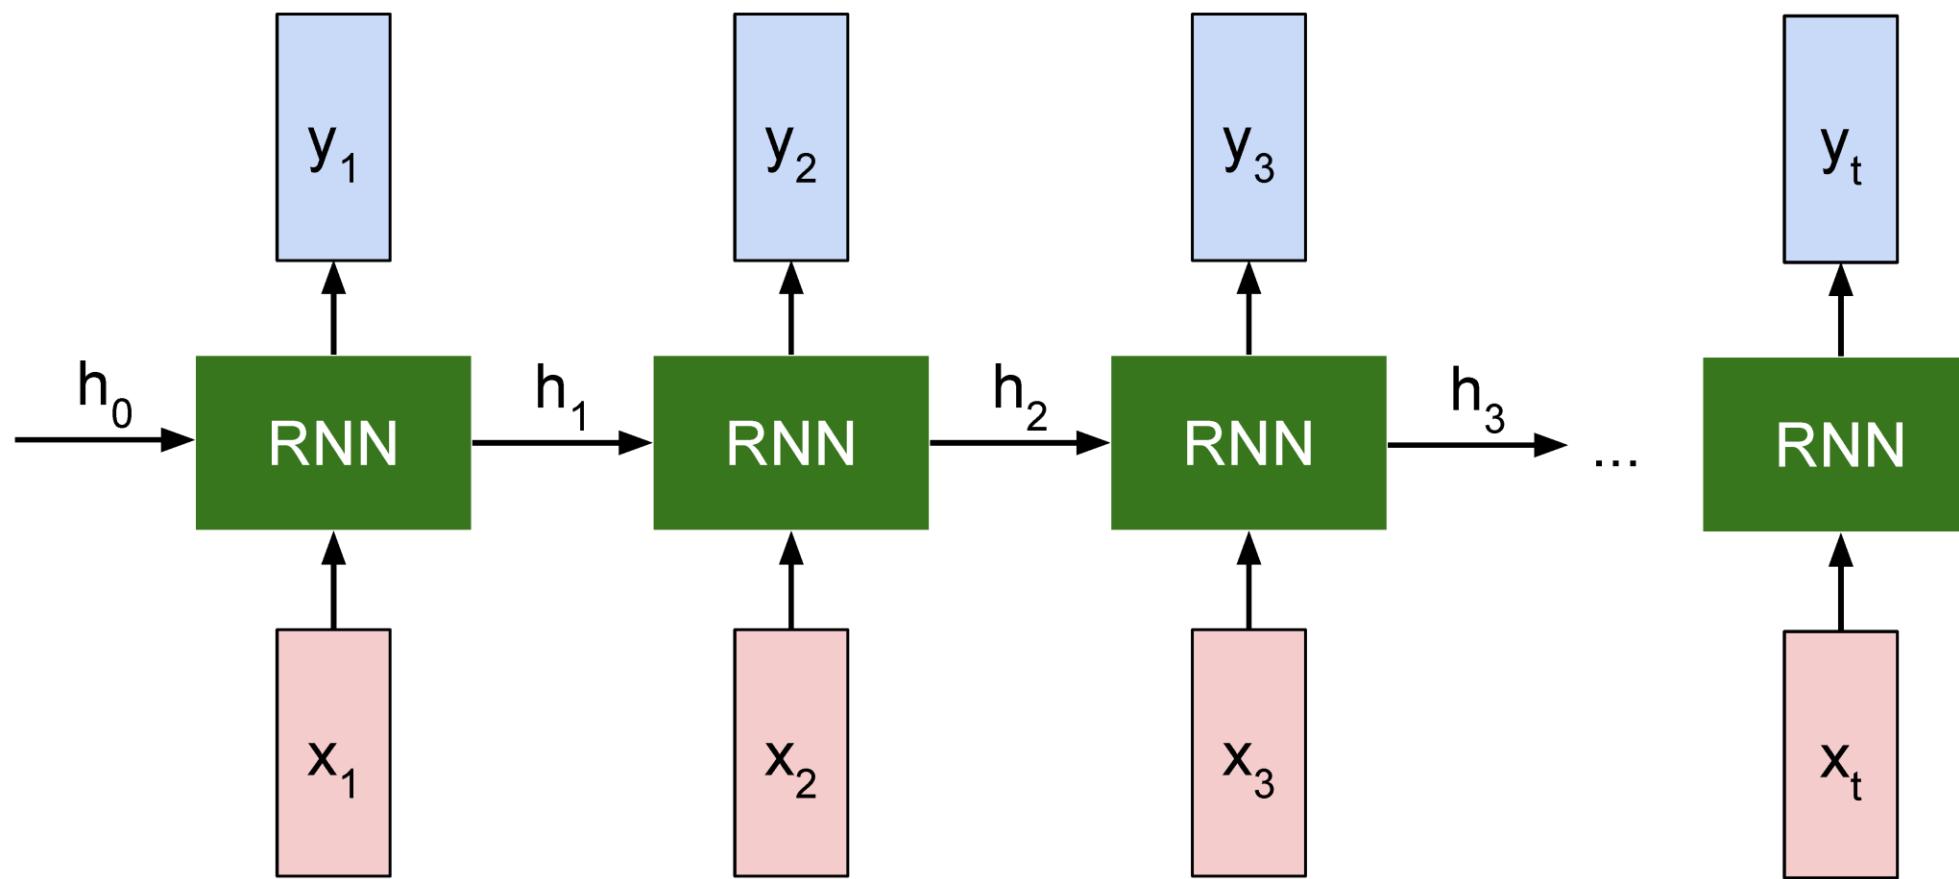
We can process a sequence of vectors \mathbf{x} by applying a **recurrence formula** at every time step:

$$h_t = f_W(h_{t-1}, x_t)$$

new state old state input vector at
 / some time step
 some function
 with parameters W



Recurrent Neural Network

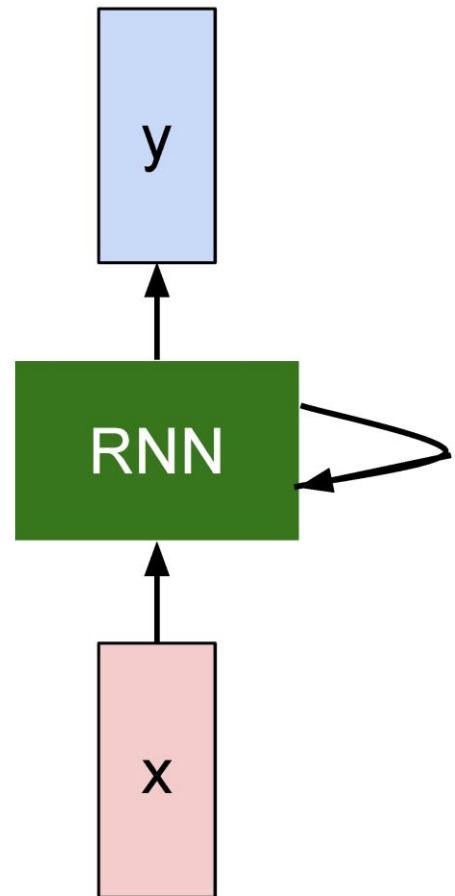


Recurrent Neural Network

We can process a sequence of vectors \mathbf{x} by applying a **recurrence formula** at every time step:

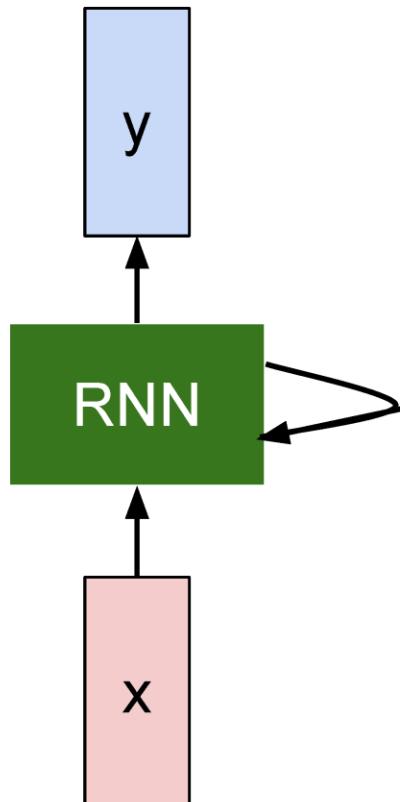
$$h_t = f_W(h_{t-1}, x_t)$$

Notice: the same function and the same set of parameters are used at every time step.



(Vanilla) Recurrent Neural Network

The state consists of a single “*hidden*” vector \mathbf{h} :



$$\mathbf{h}_t = f_W(\mathbf{h}_{t-1}, \mathbf{x}_t)$$

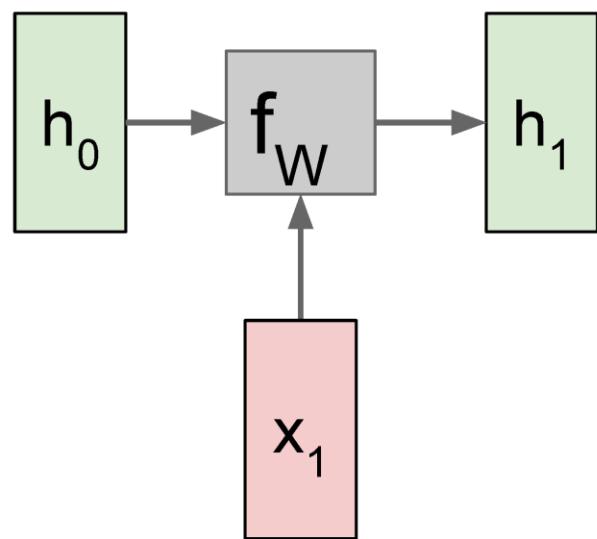


$$\mathbf{h}_t = \tanh(\mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{W}_{xh}\mathbf{x}_t)$$

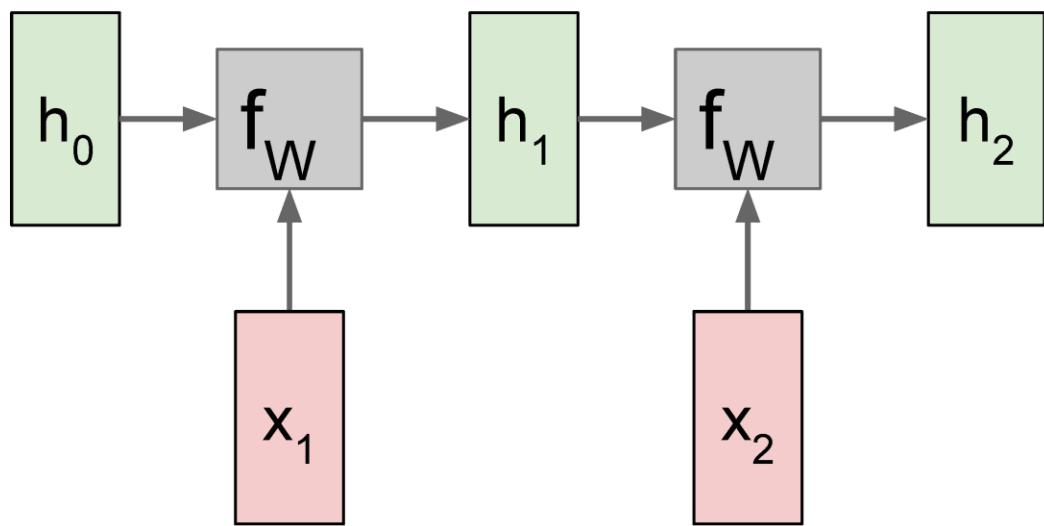
$$y_t = \mathbf{W}_{hy}\mathbf{h}_t$$

Sometimes called a “Vanilla RNN” or an
“Elman RNN” after Prof. Jeffrey Elman

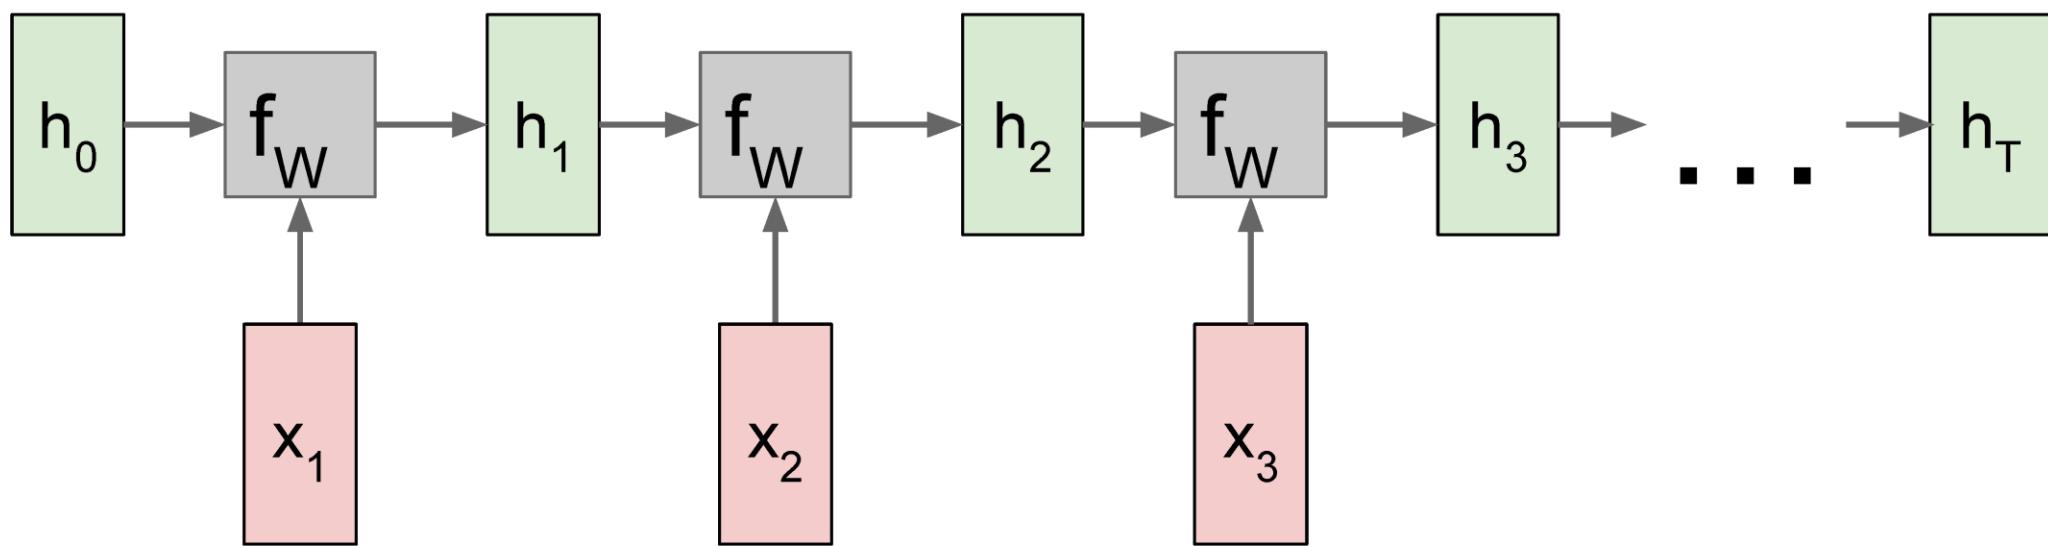
RNN: Computational Graph



RNN: Computational Graph

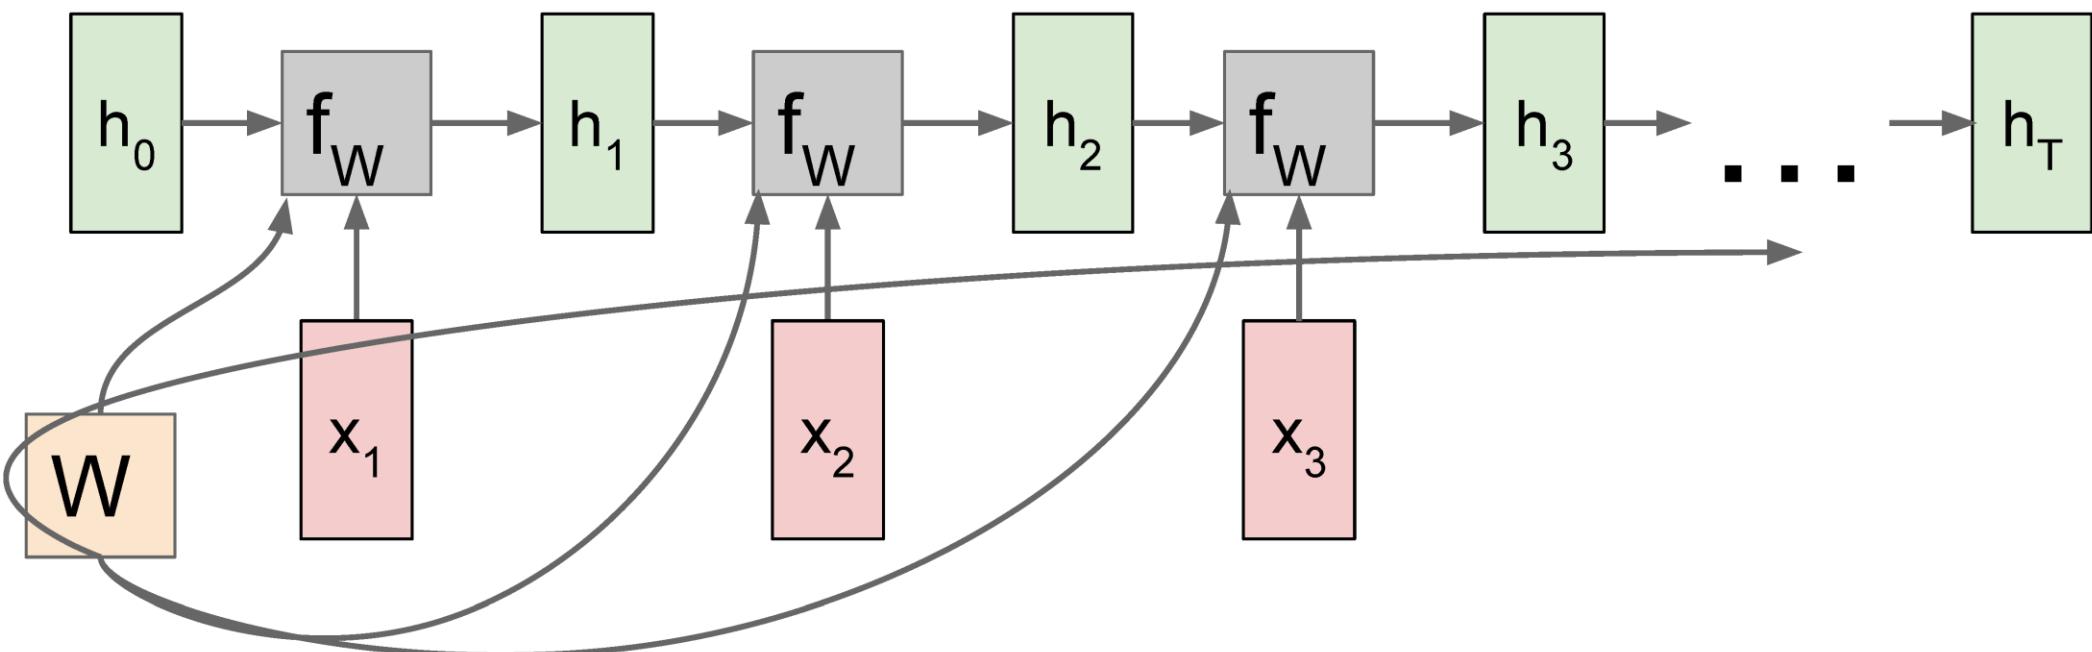


RNN: Computational Graph

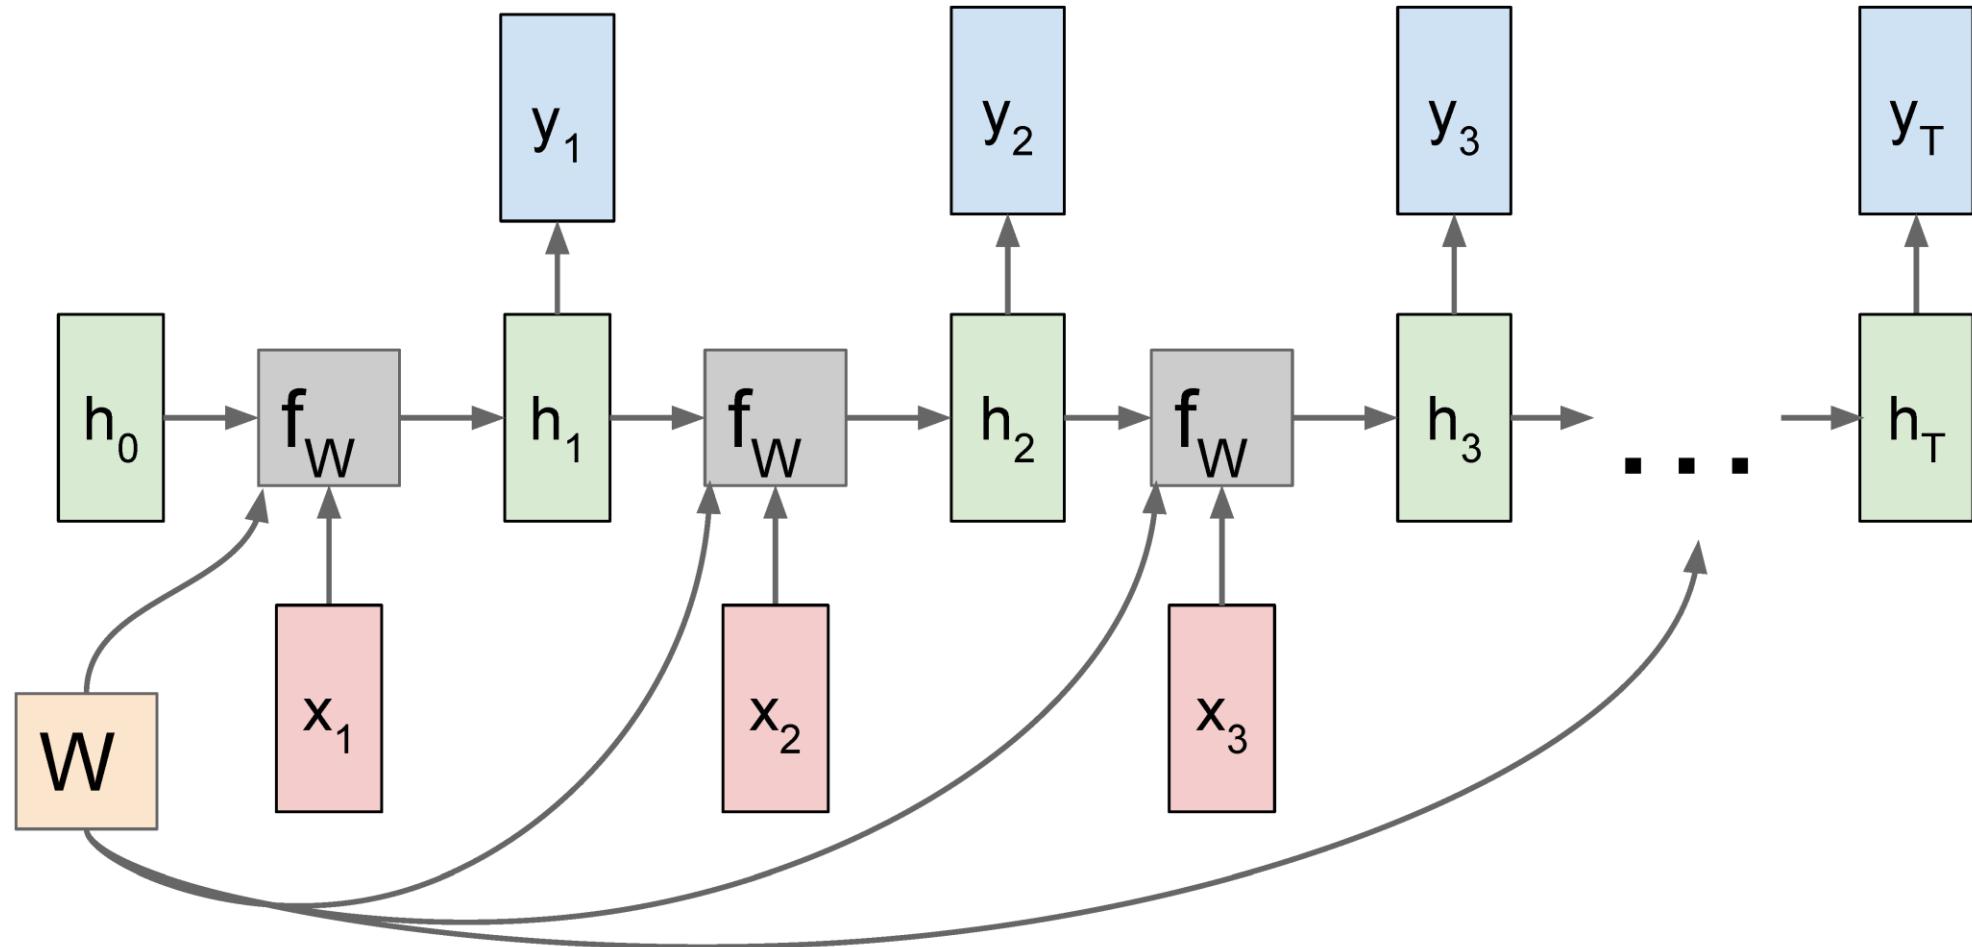


RNN: Computational Graph

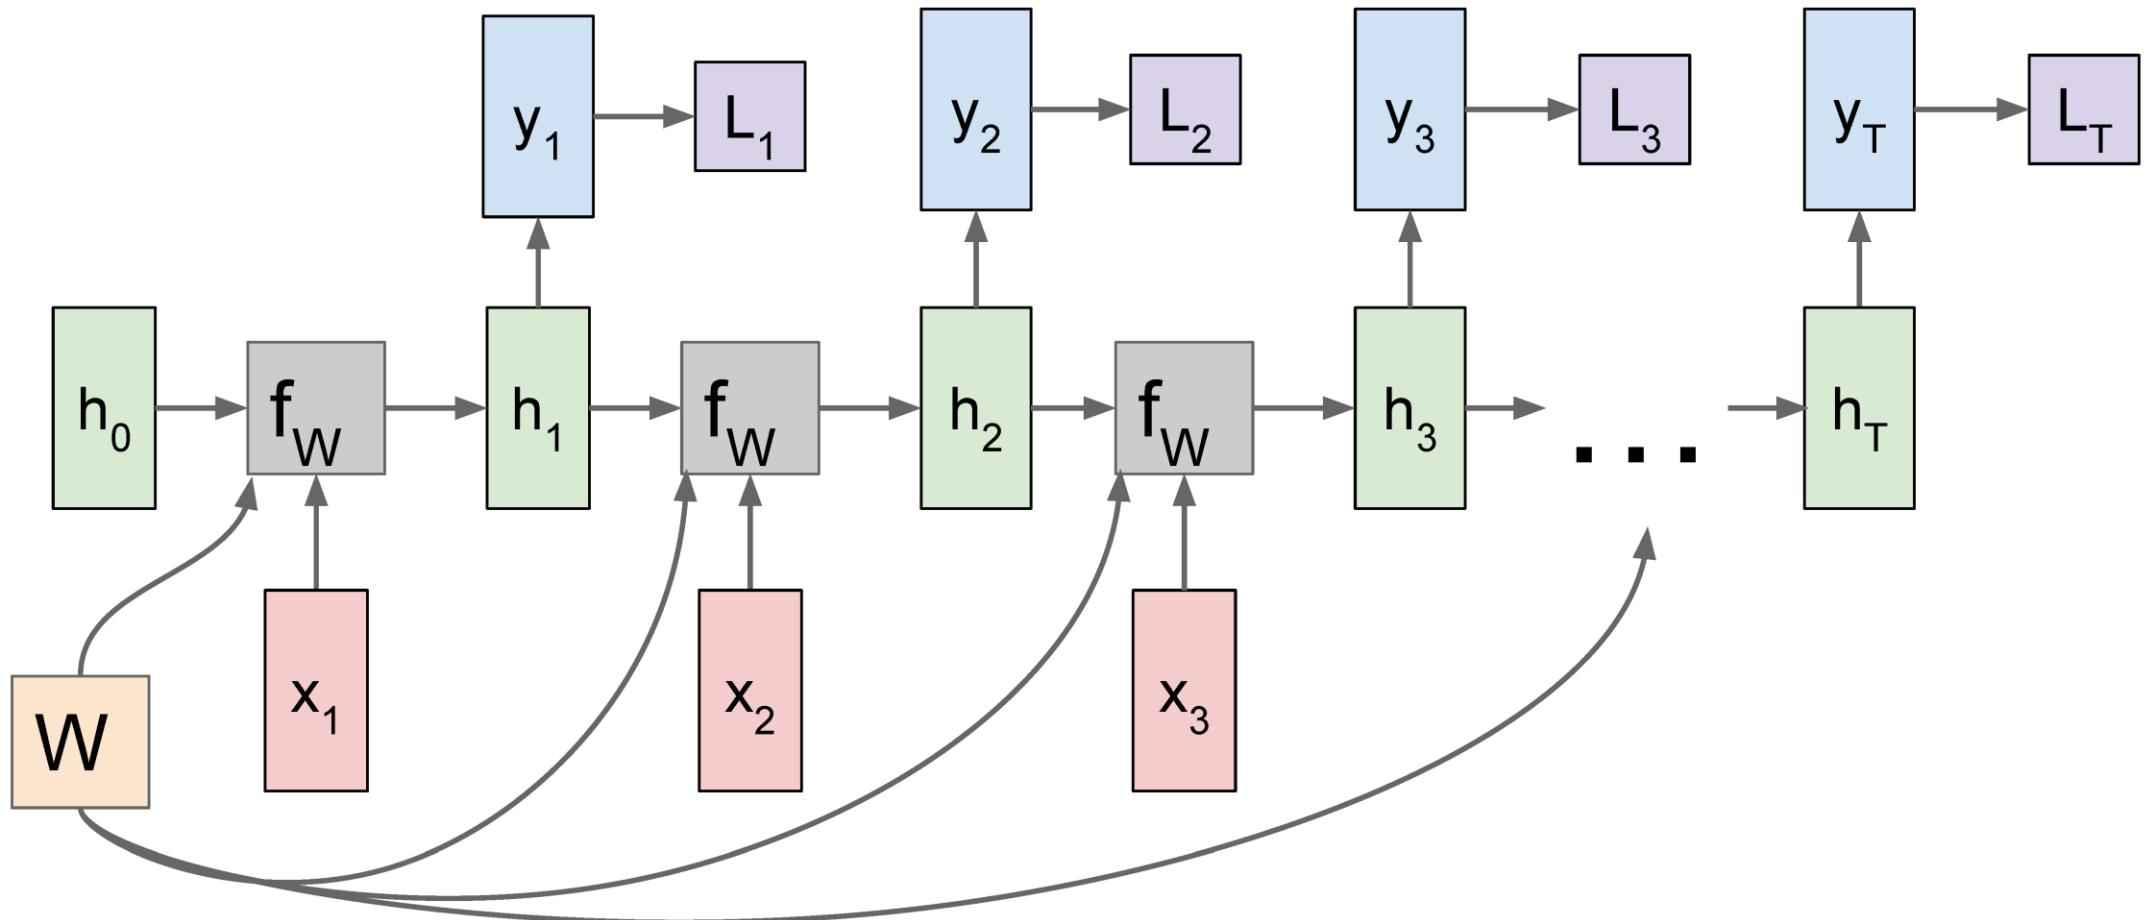
Re-use the same weight matrix at every time-step



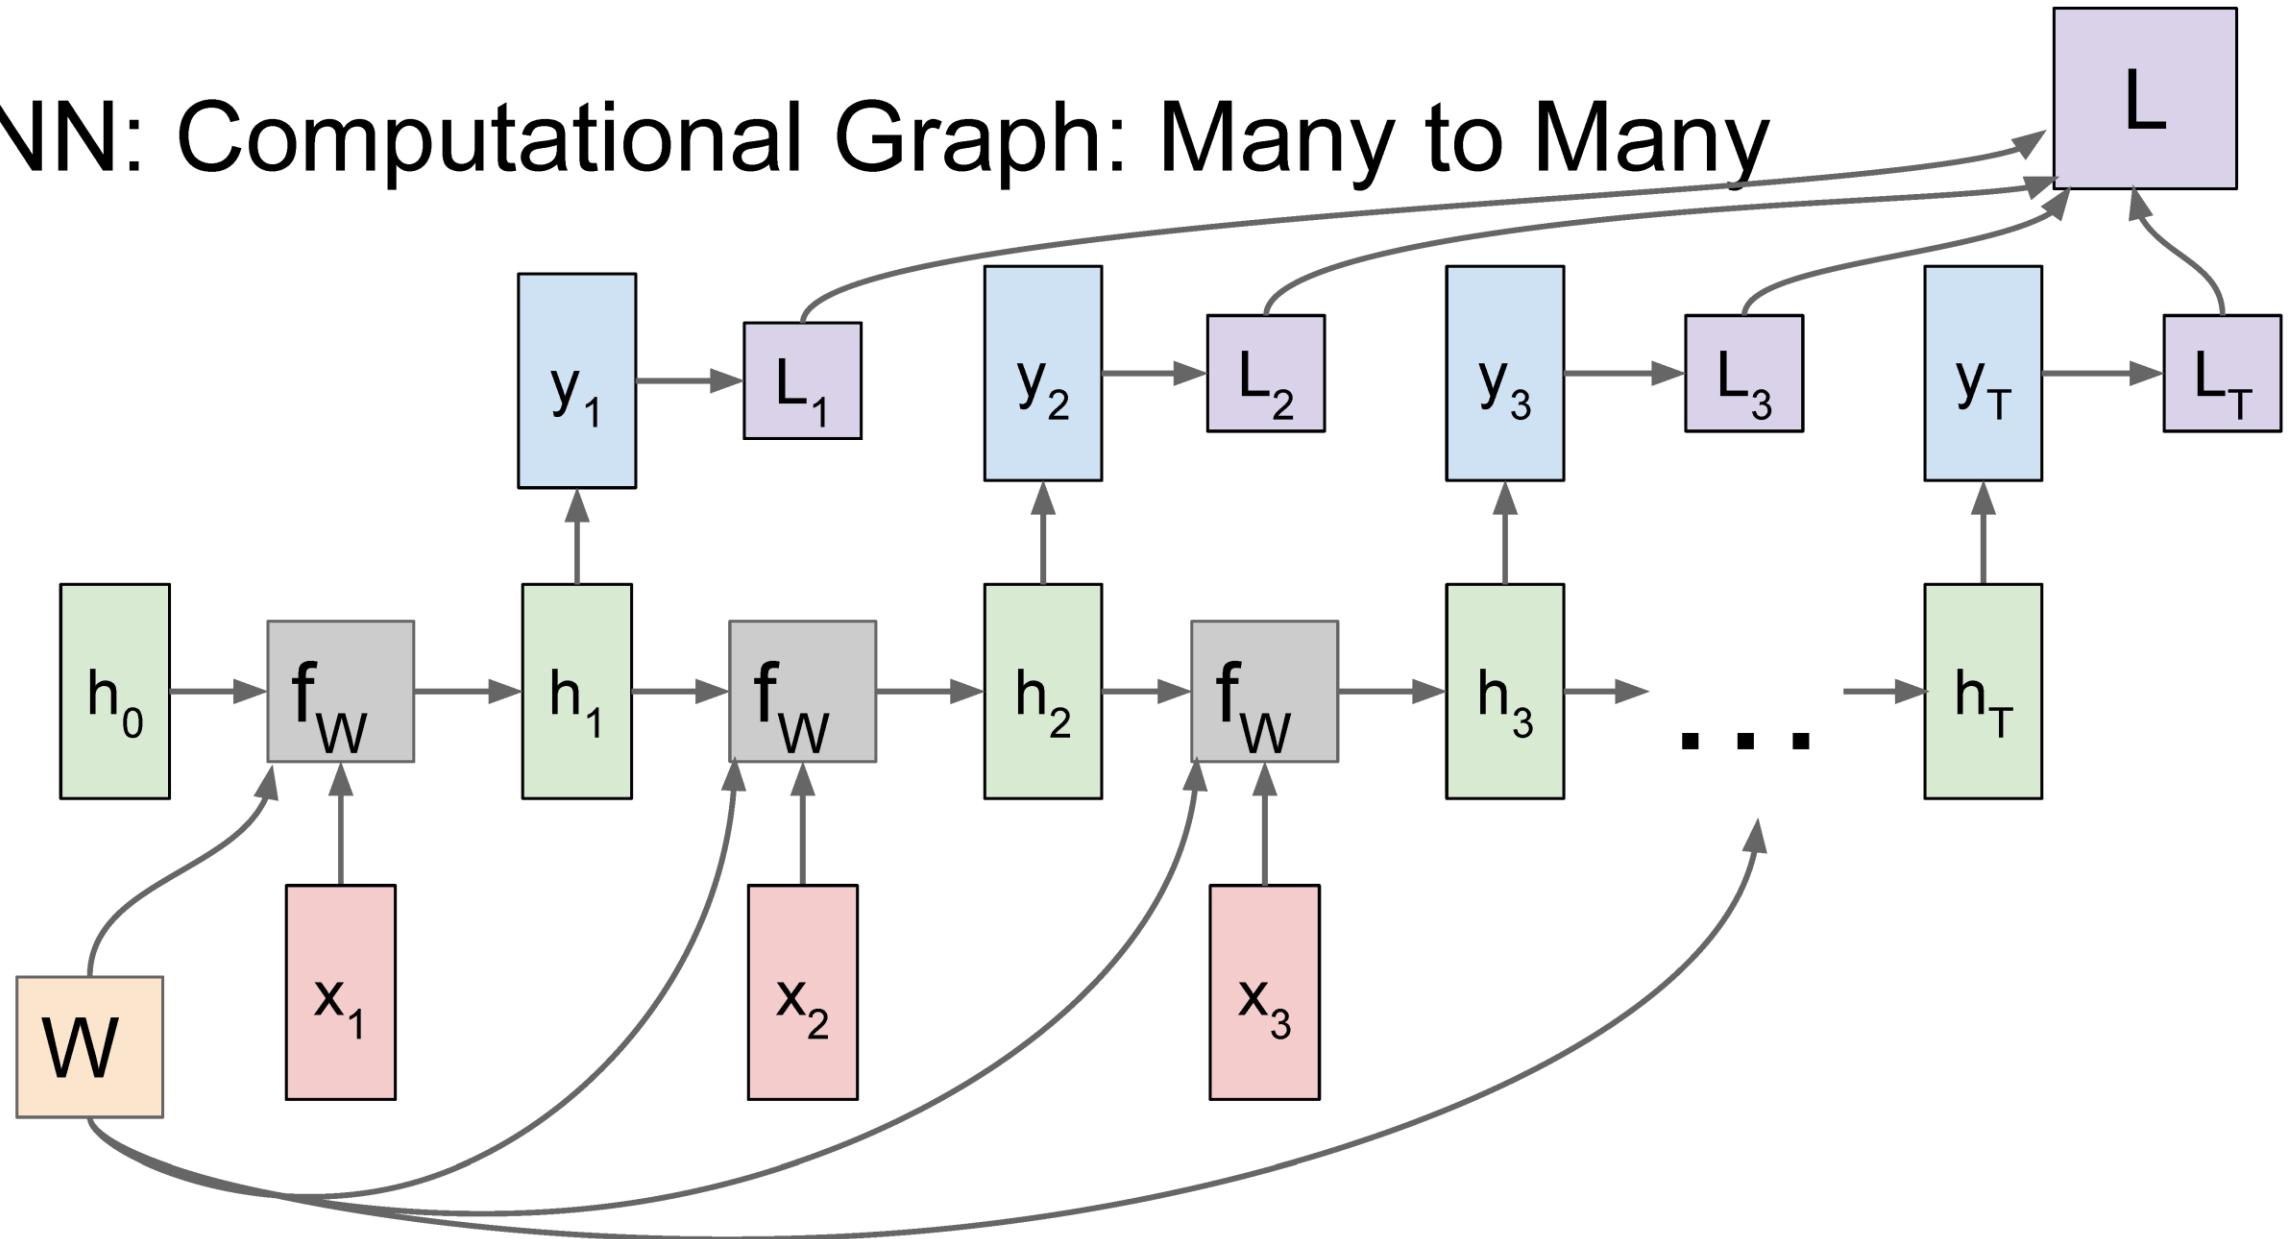
RNN: Computational Graph: Many to Many



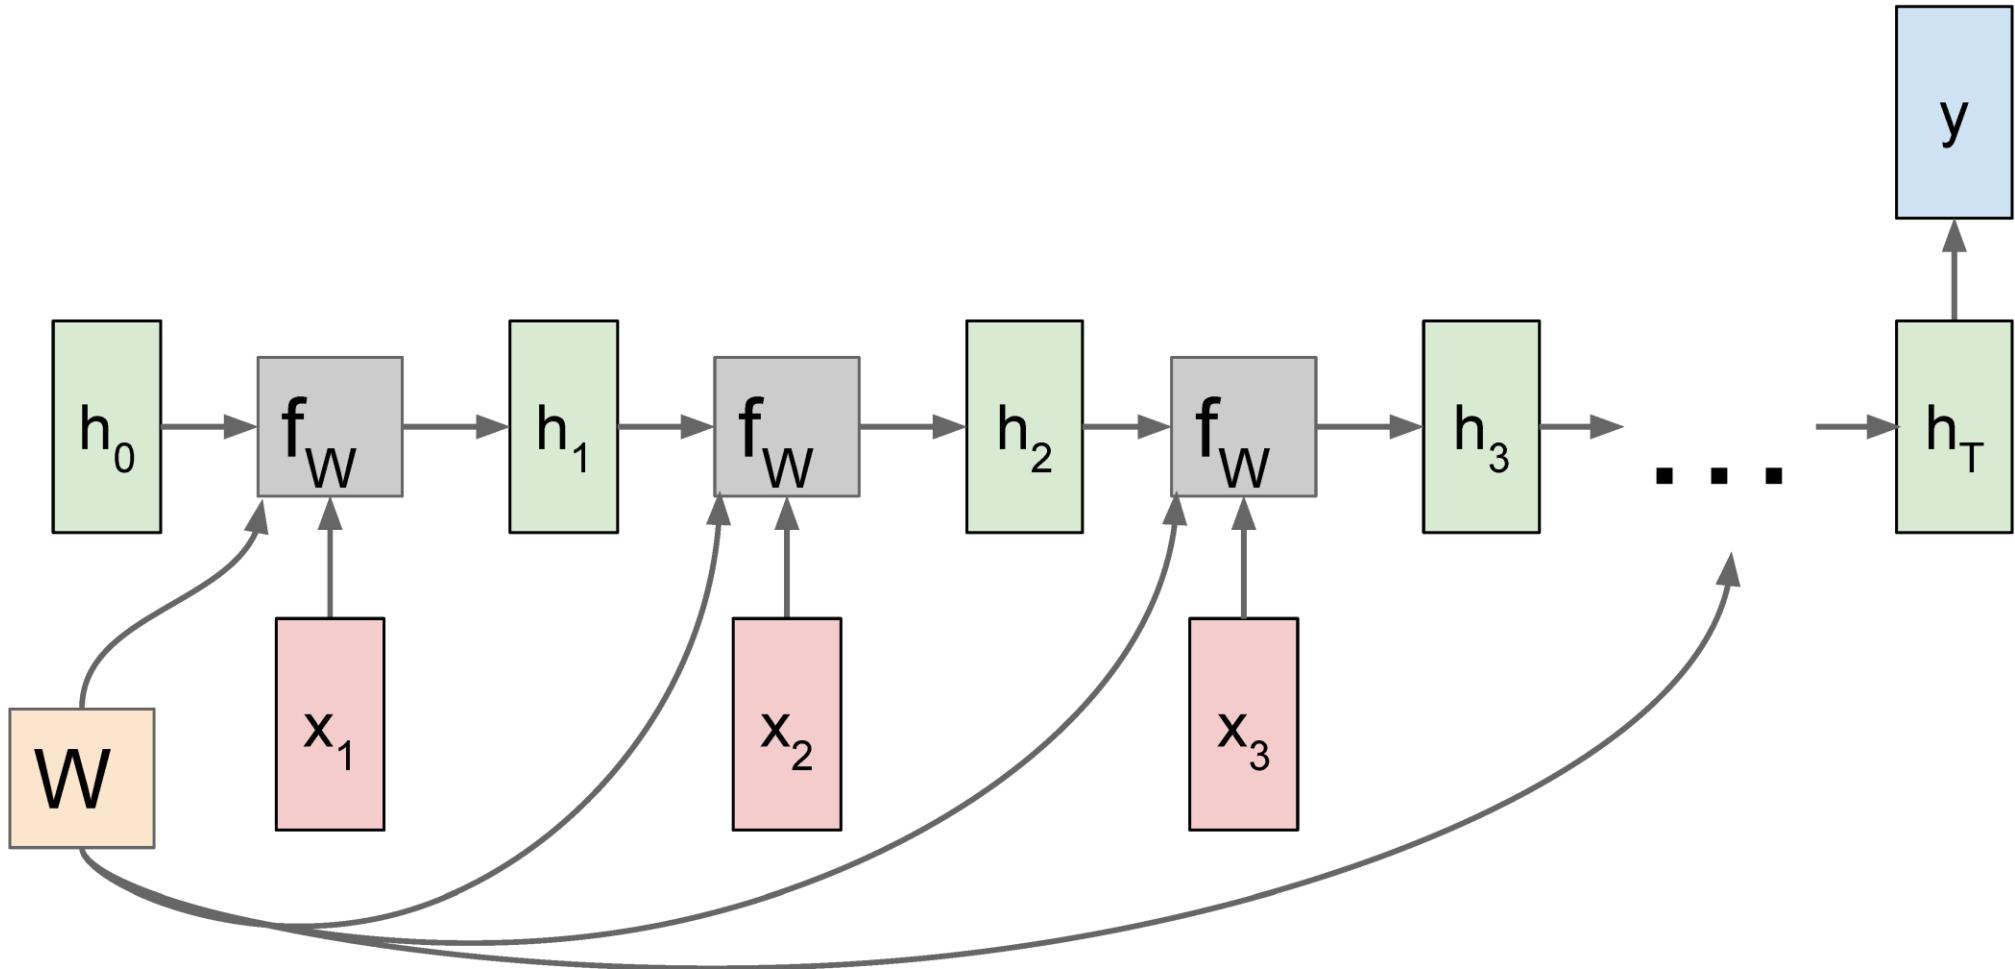
RNN: Computational Graph: Many to Many



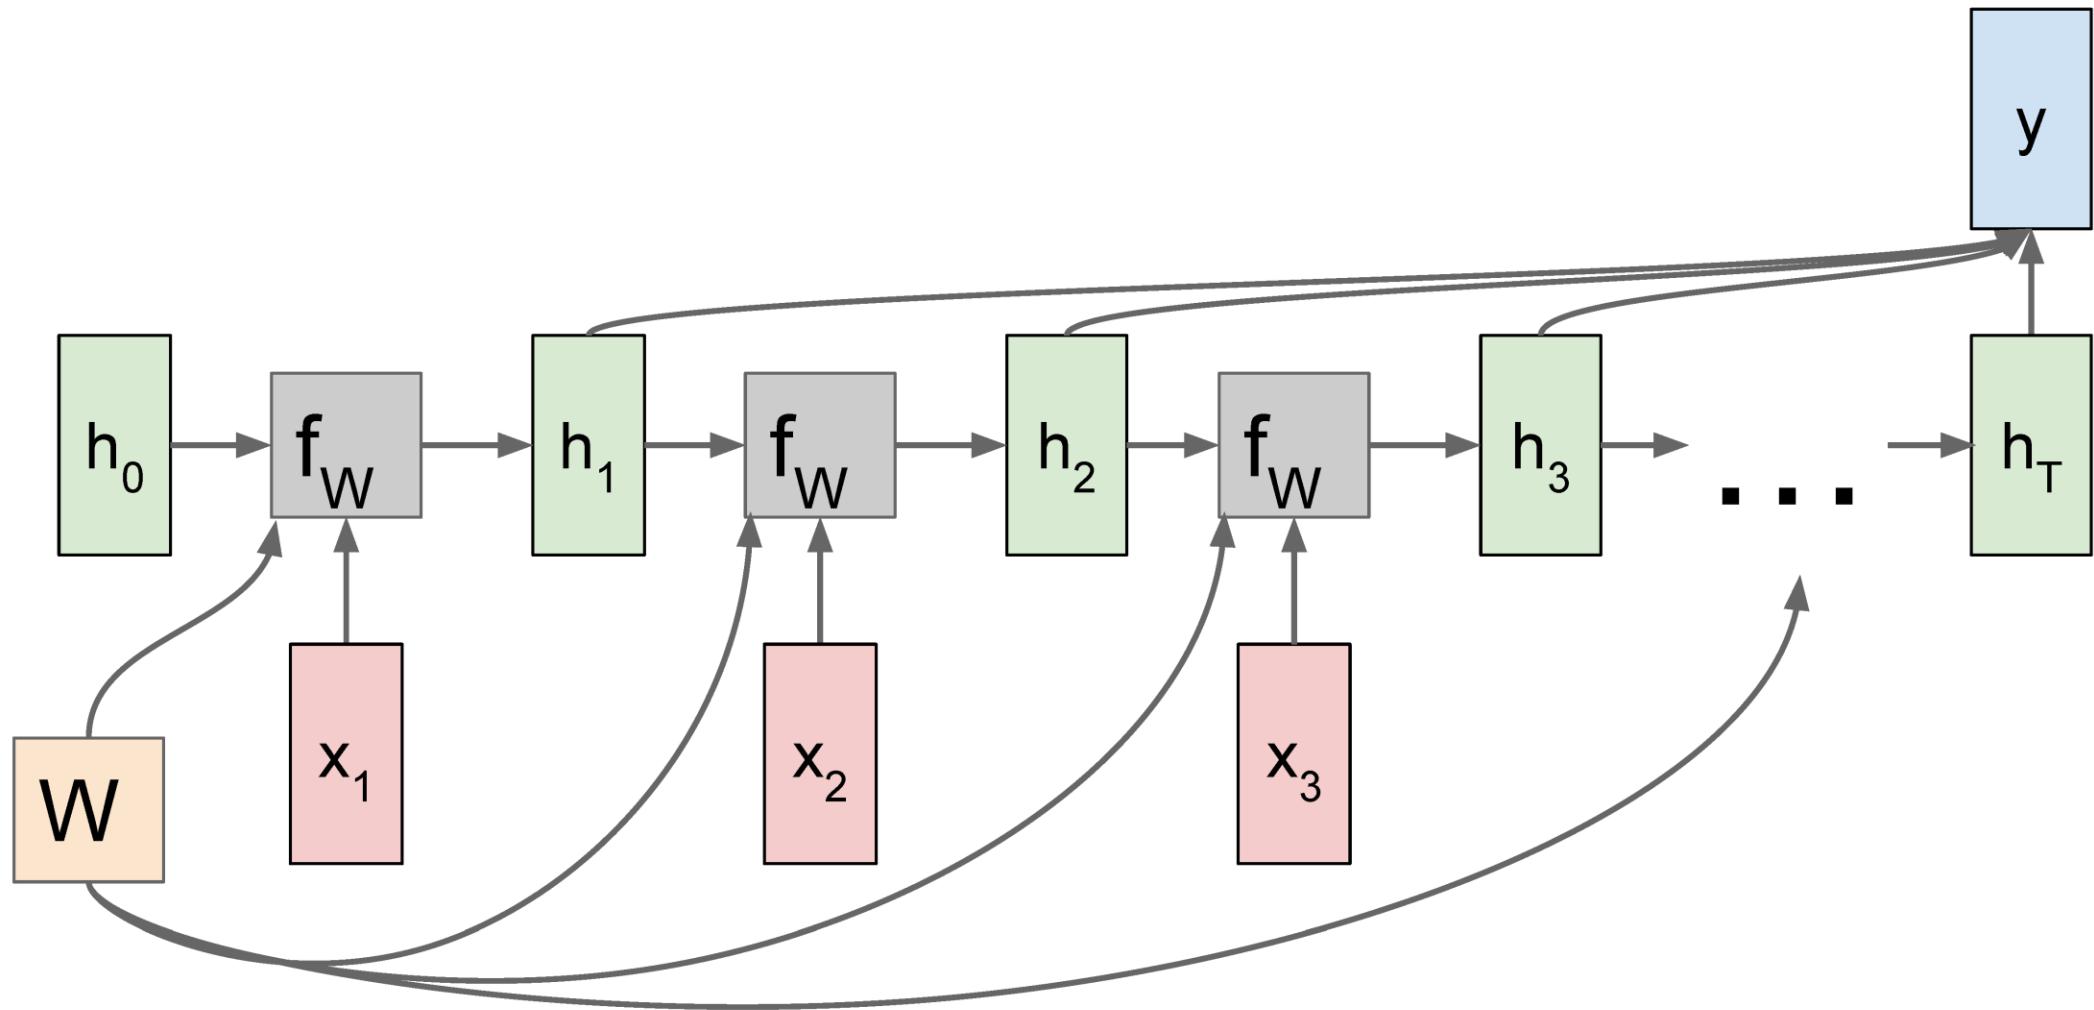
RNN: Computational Graph: Many to Many



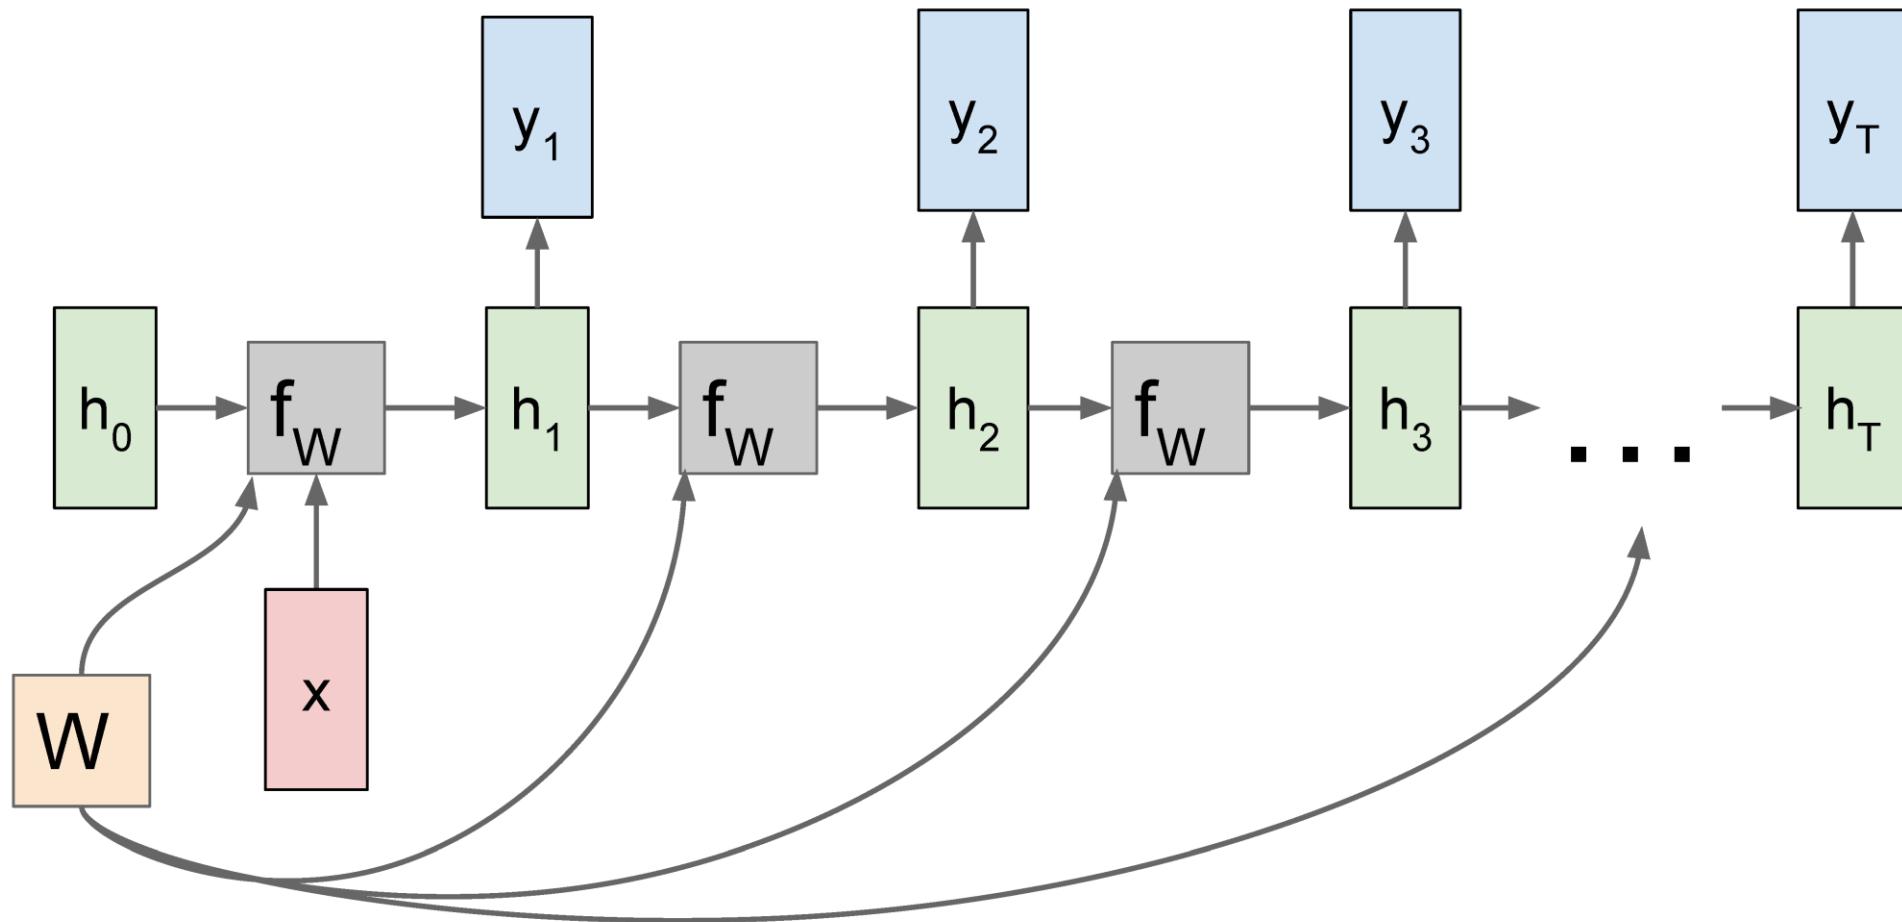
RNN: Computational Graph: Many to One



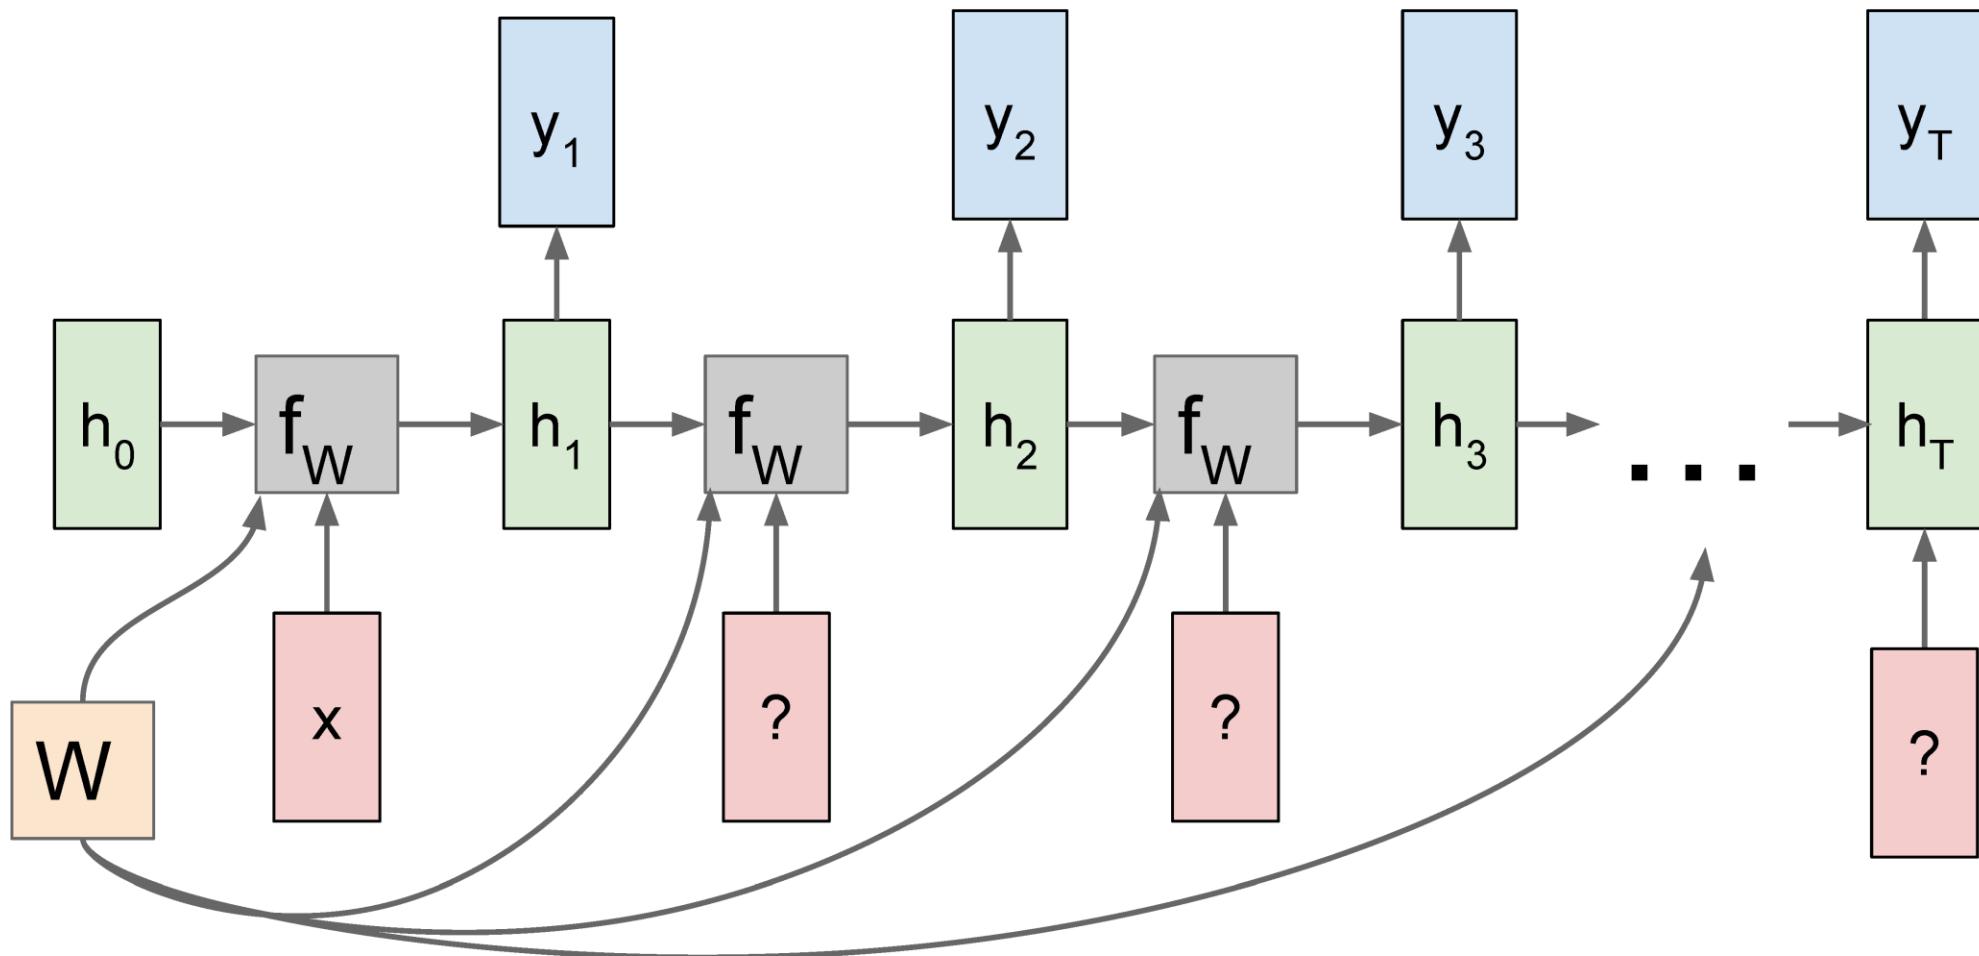
RNN: Computational Graph: Many to One



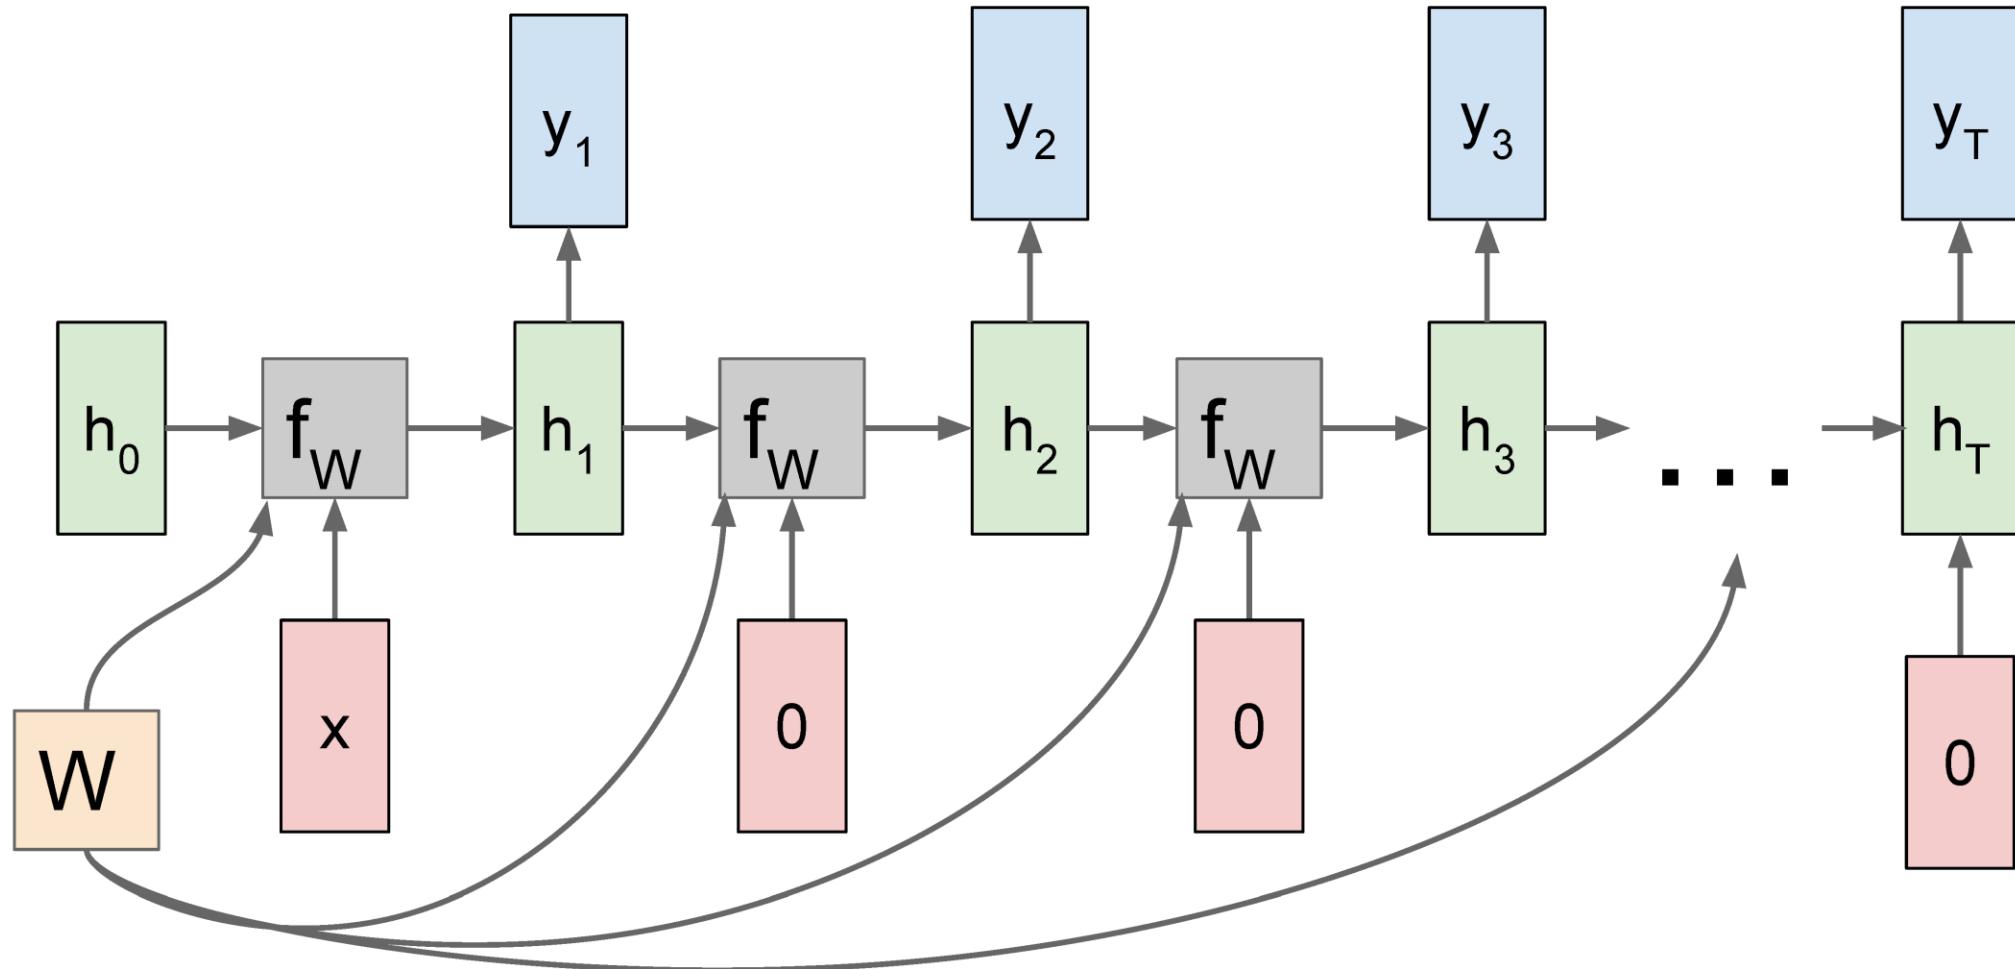
RNN: Computational Graph: One to Many



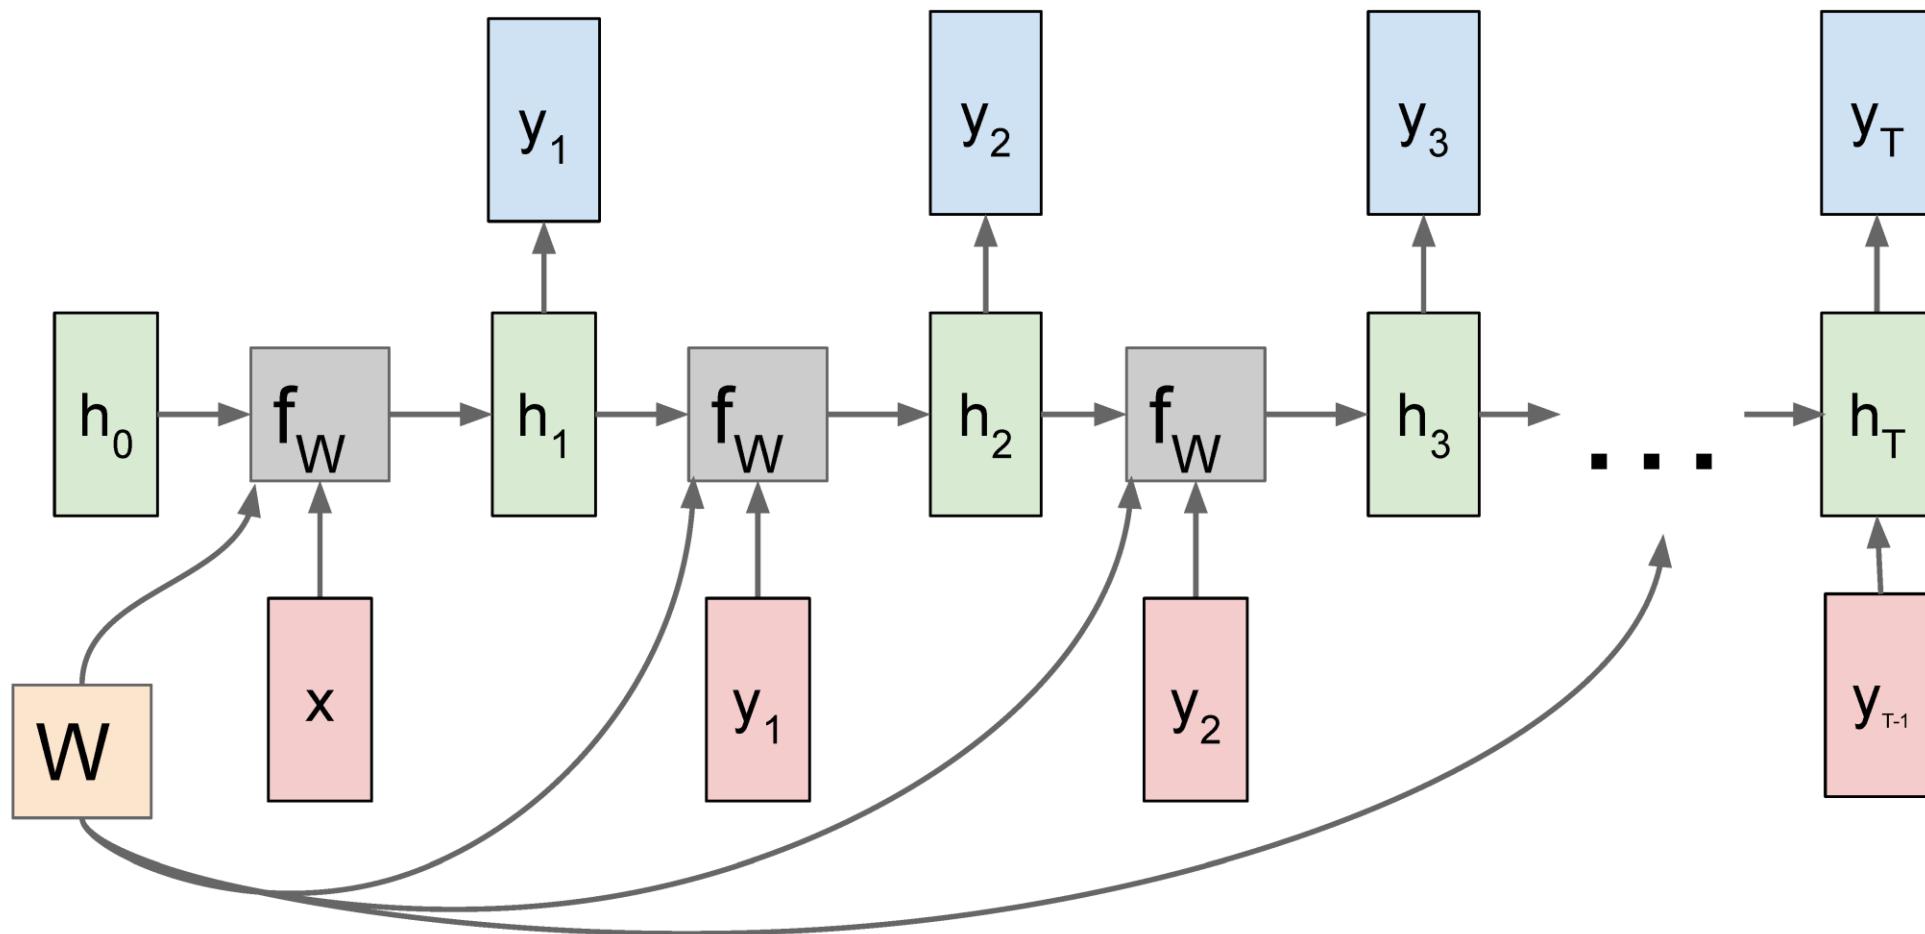
RNN: Computational Graph: One to Many



RNN: Computational Graph: One to Many

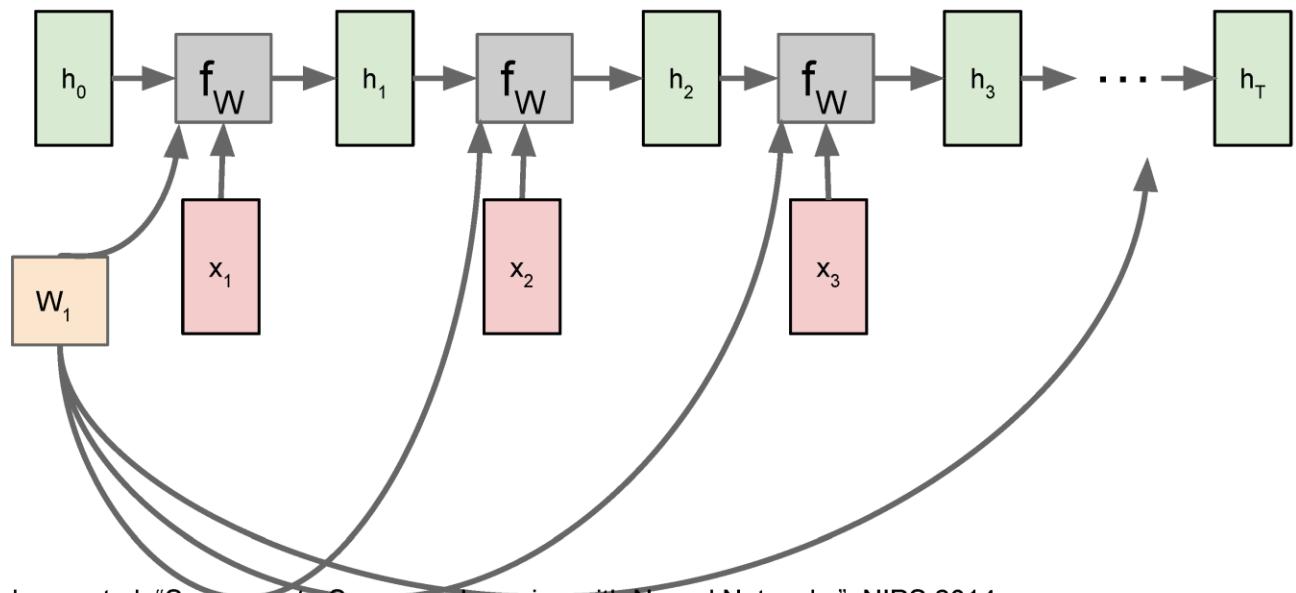


RNN: Computational Graph: One to Many



Sequence to Sequence: Many-to-one + one-to-many

Many to one: Encode input sequence in a single vector

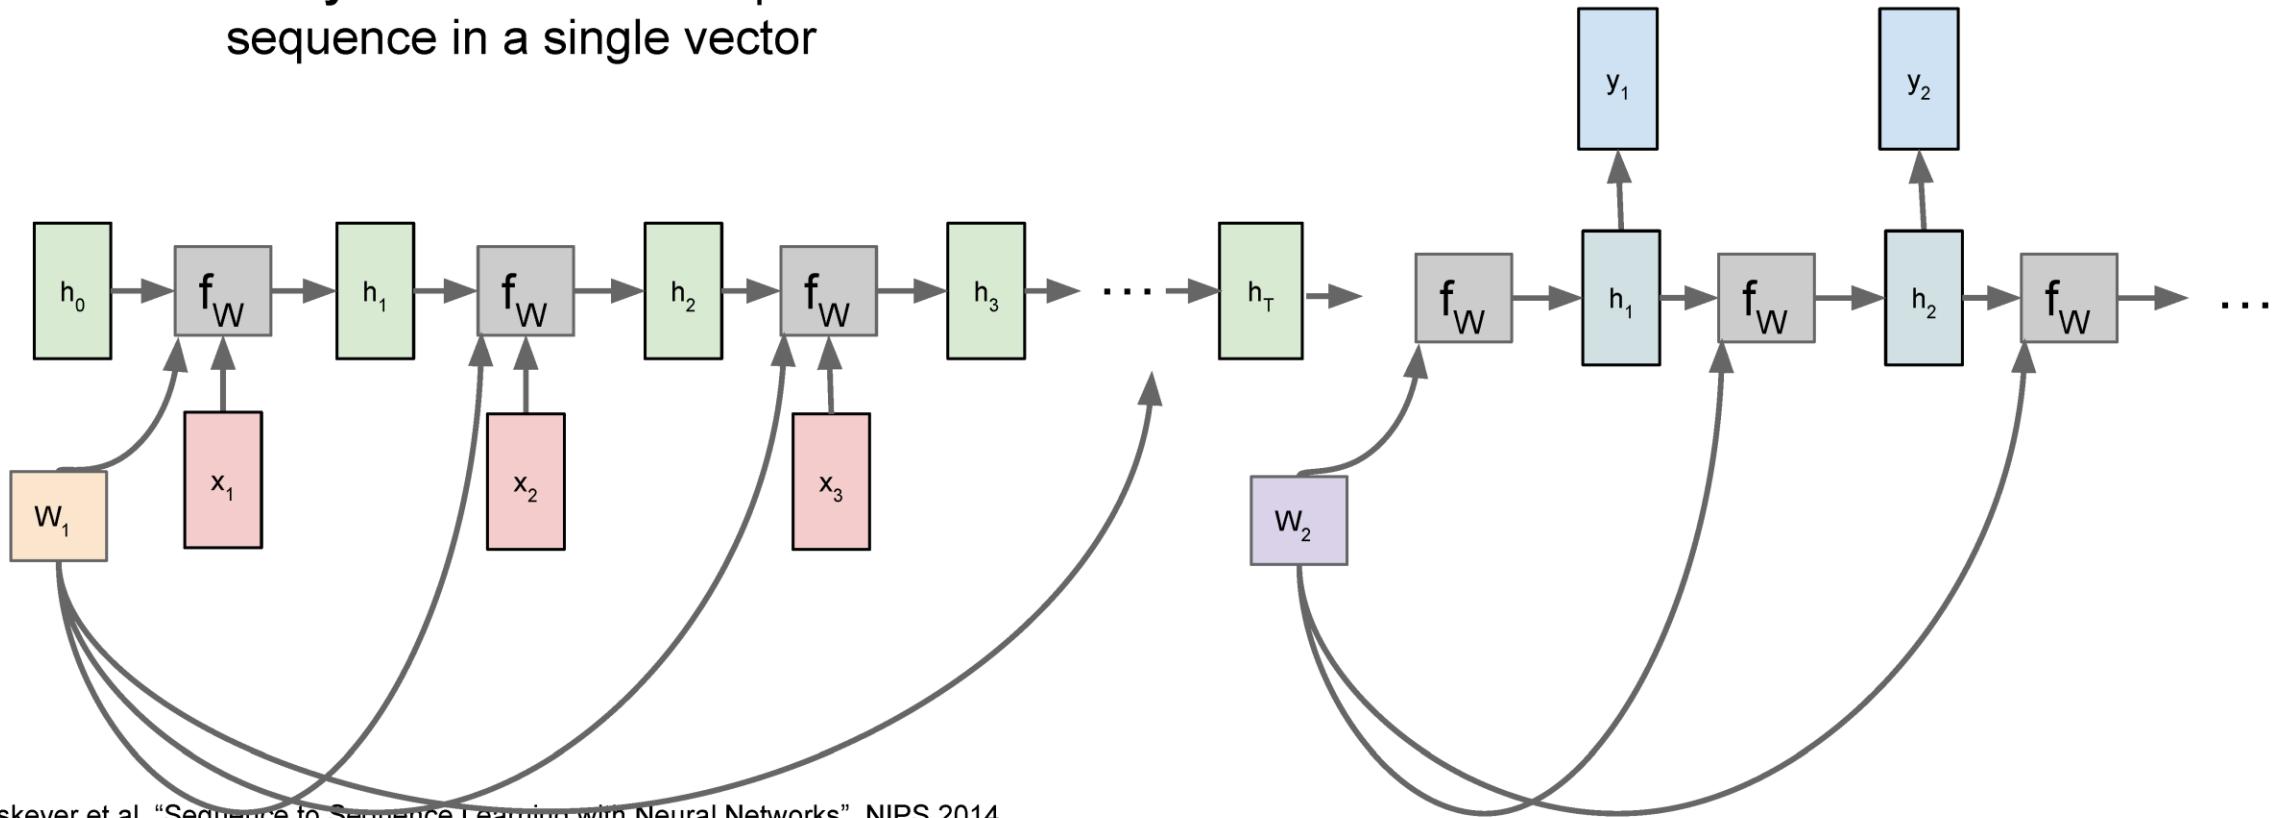


Sutskever et al, "Sequence to Sequence Learning with Neural Networks", NIPS 2014

Sequence to Sequence: Many-to-one + one-to-many

Many to one: Encode input sequence in a single vector

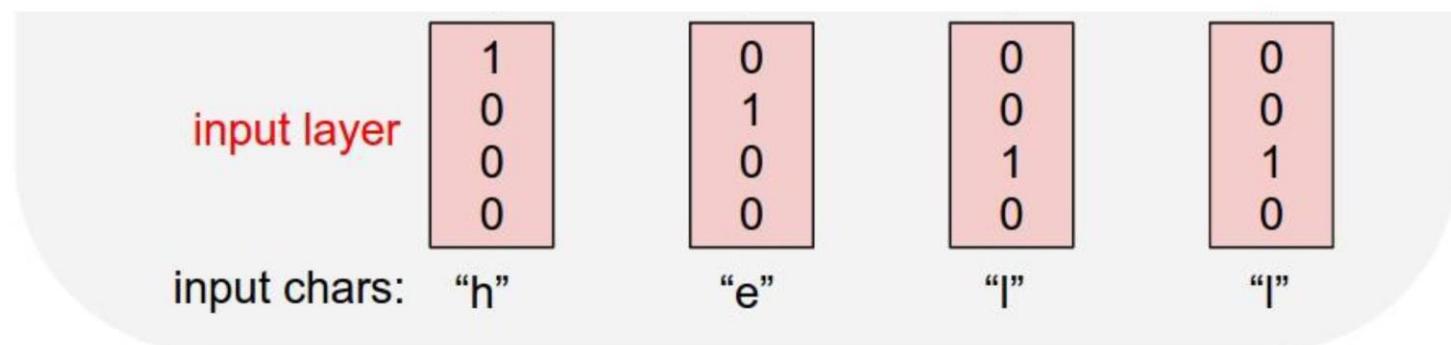
One to many: Produce output sequence from single input vector



Example: Character-level Language Model

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

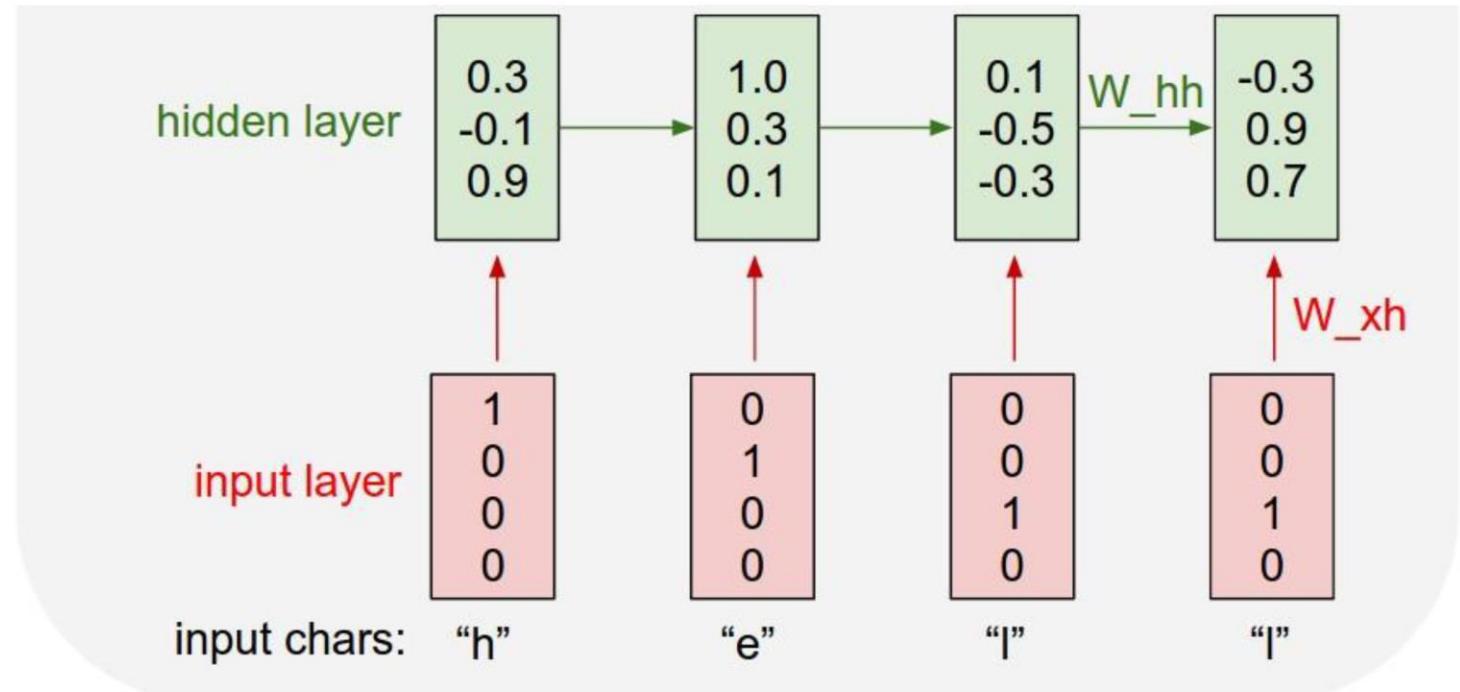


Example: Character-level Language Model

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

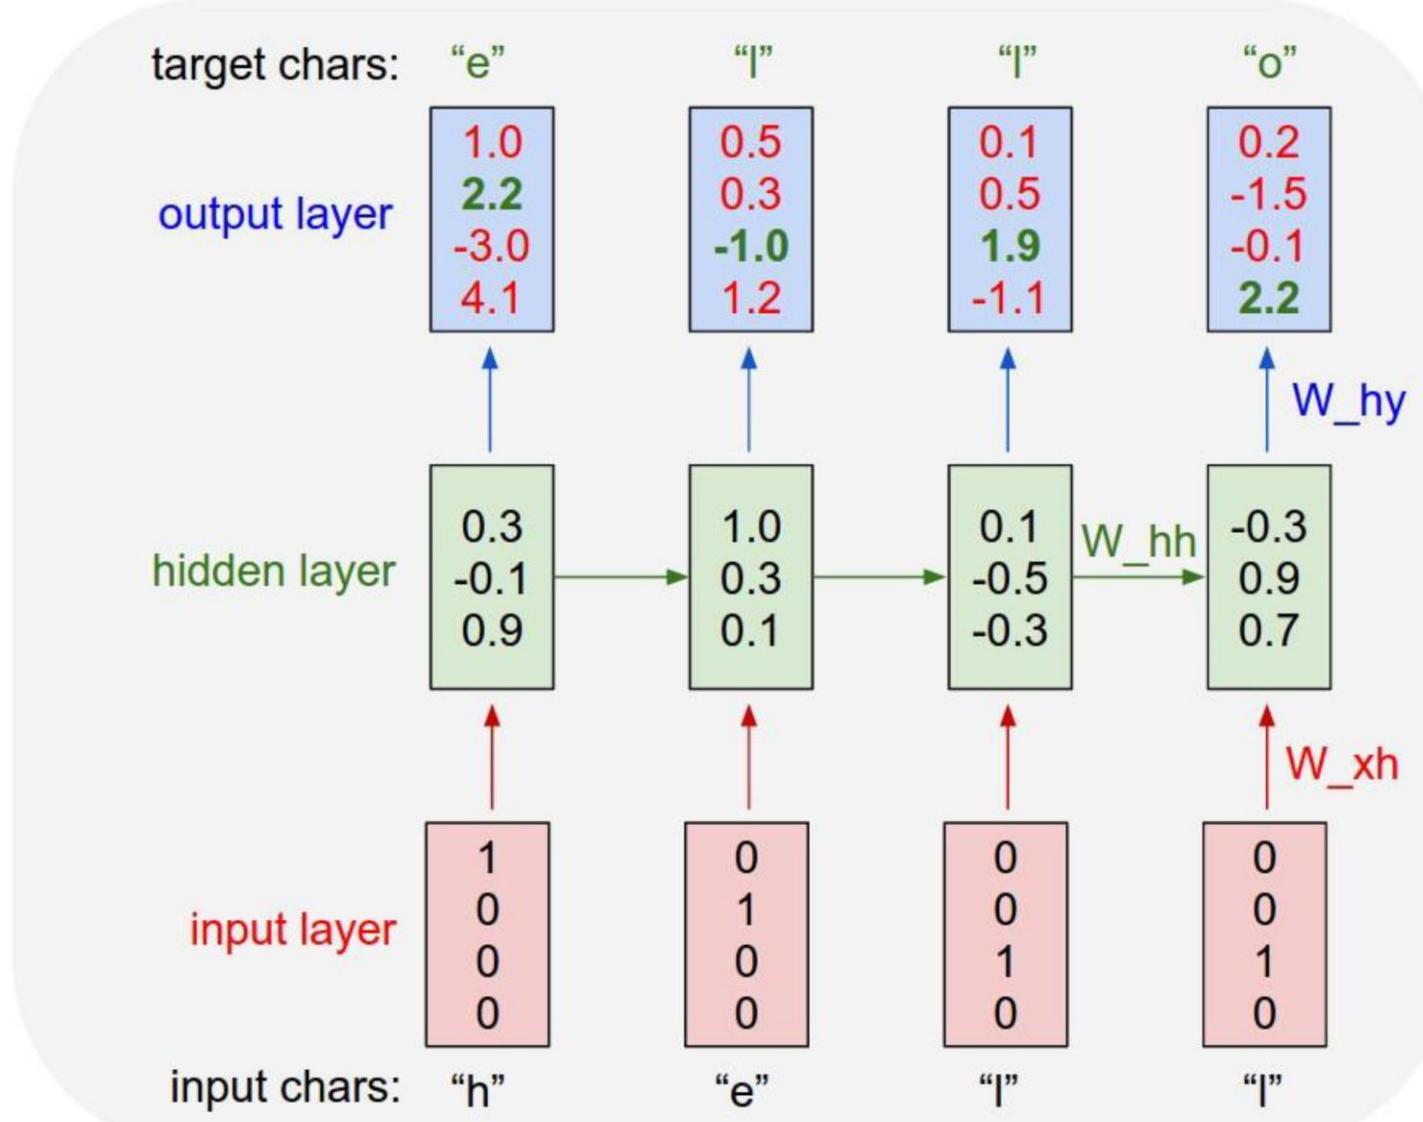
$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$



Example: Character-level Language Model

Vocabulary:
[h,e,l,o]

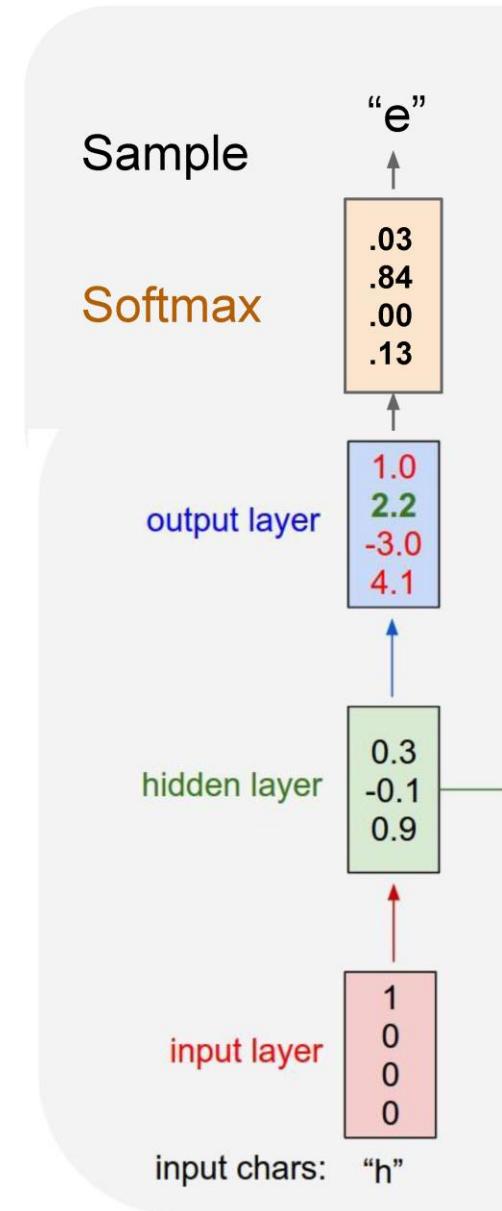
Example training
sequence:
“hello”



Example: Character-level Language Model Sampling

Vocabulary:
[h,e,l,o]

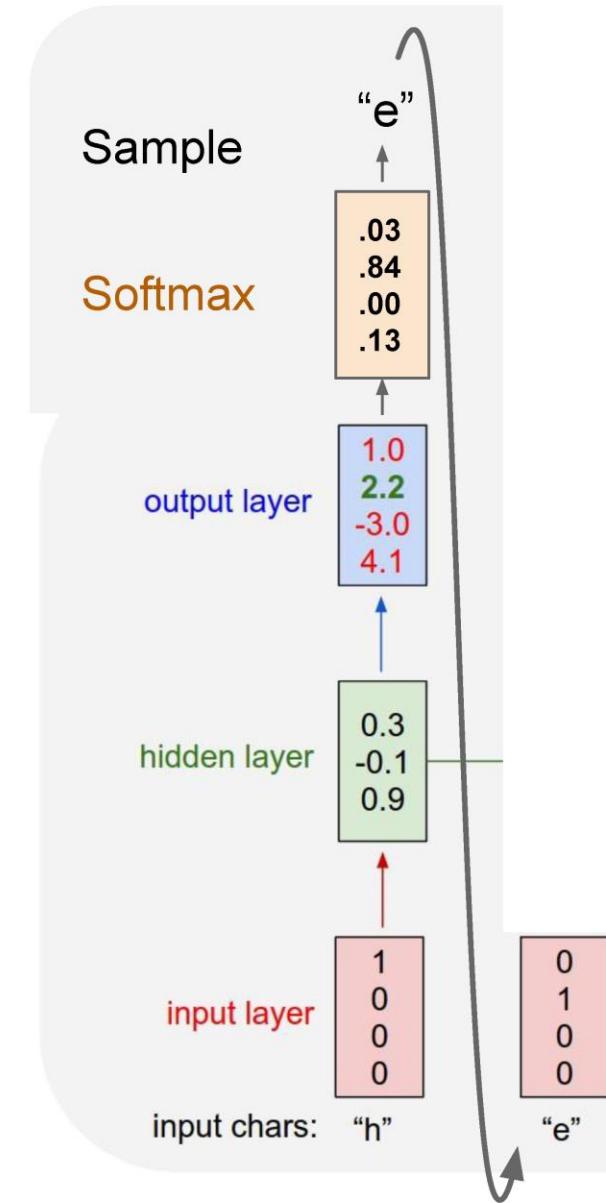
At test-time sample characters one at a time, feed back to model



Example: Character-level Language Model Sampling

Vocabulary:
[h,e,l,o]

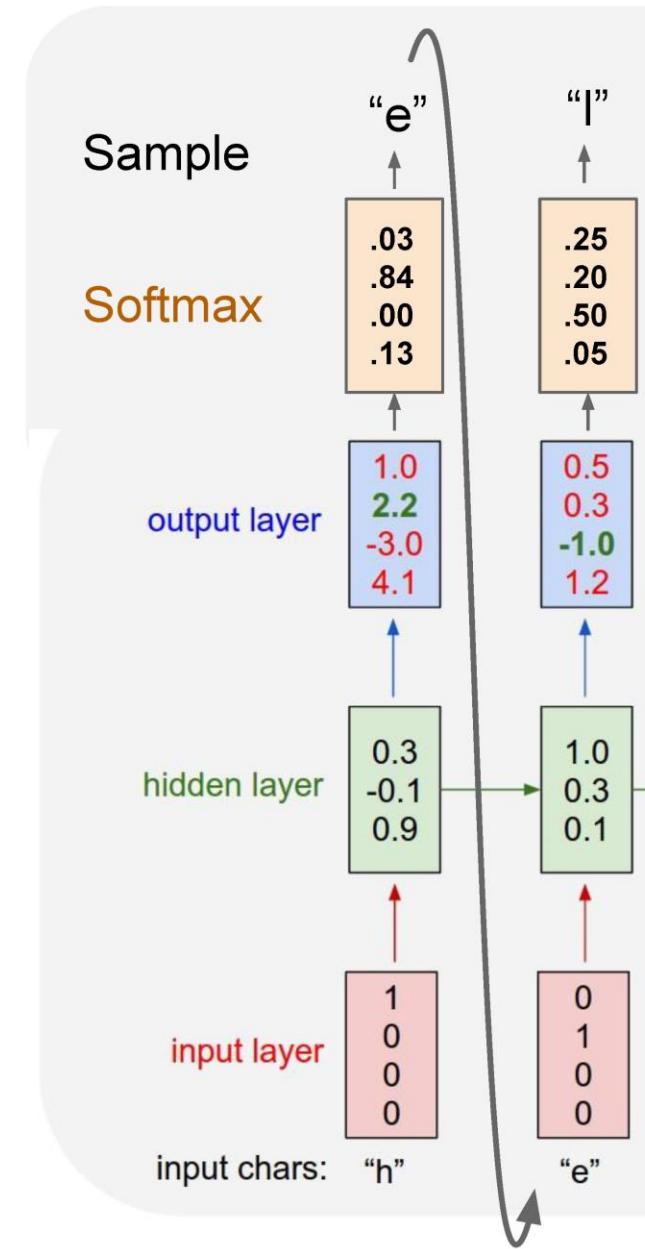
At test-time sample characters one at a time, feed back to model



Example: Character-level Language Model Sampling

Vocabulary:
[h,e,l,o]

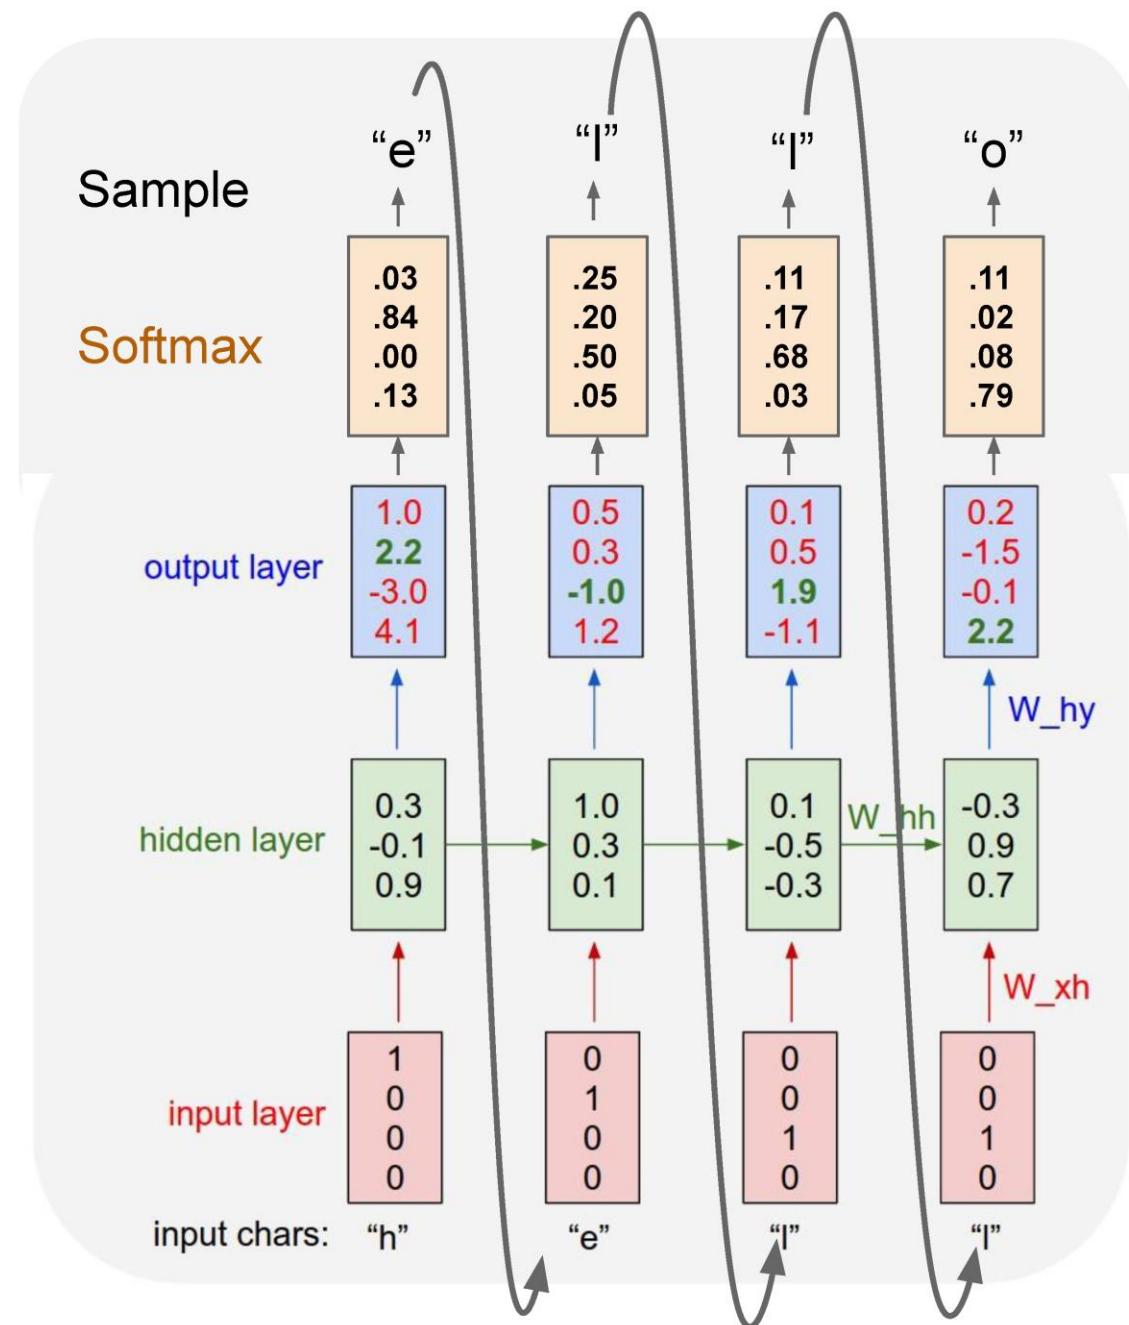
At test-time sample characters one at a time, feed back to model



Example: Character-level Language Model Sampling

Vocabulary:
[h,e,l,o]

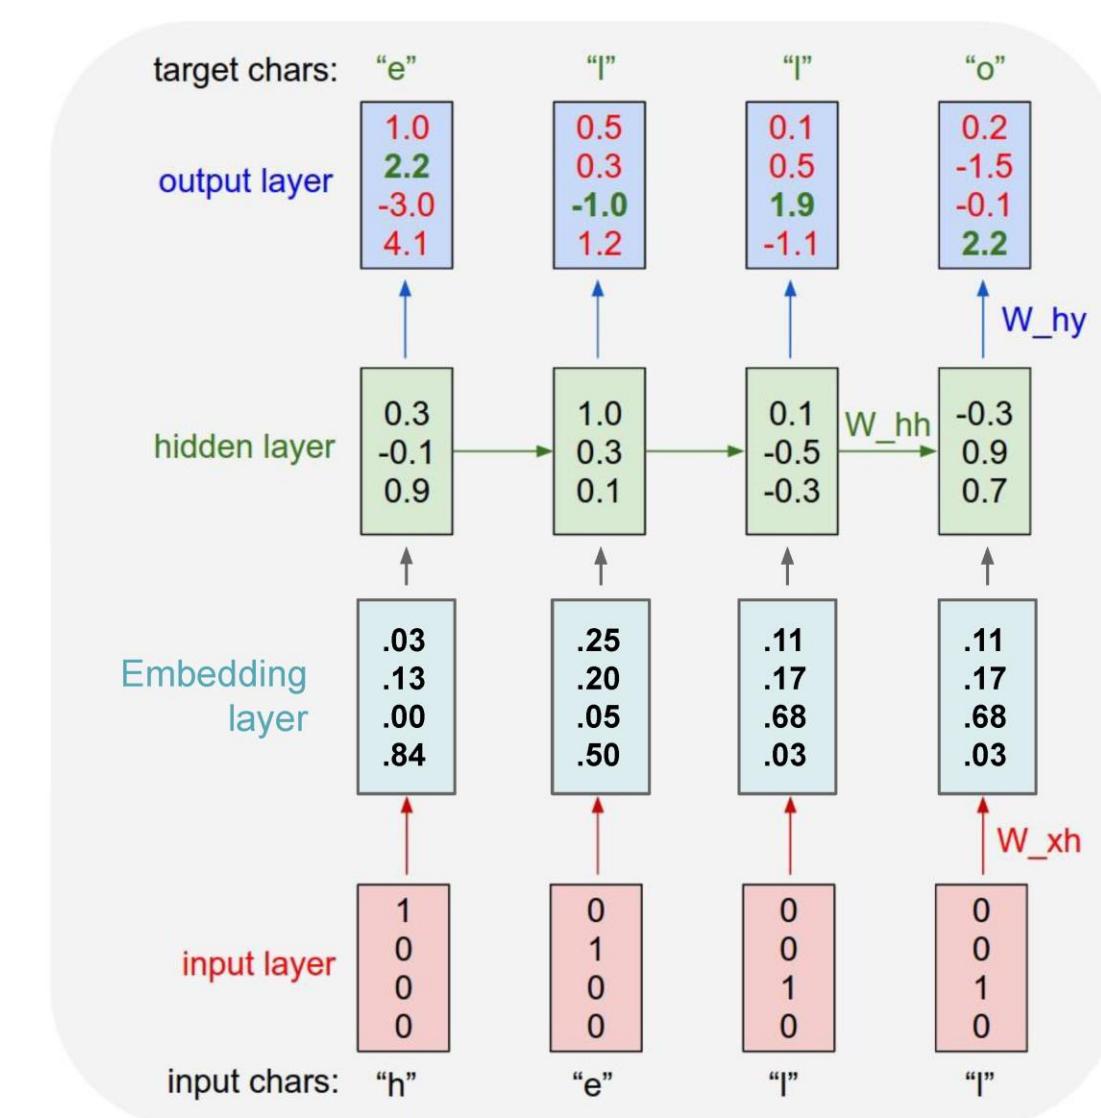
At test-time sample
characters one at a time, feed
back to model



Example: Character-level Language Model Sampling

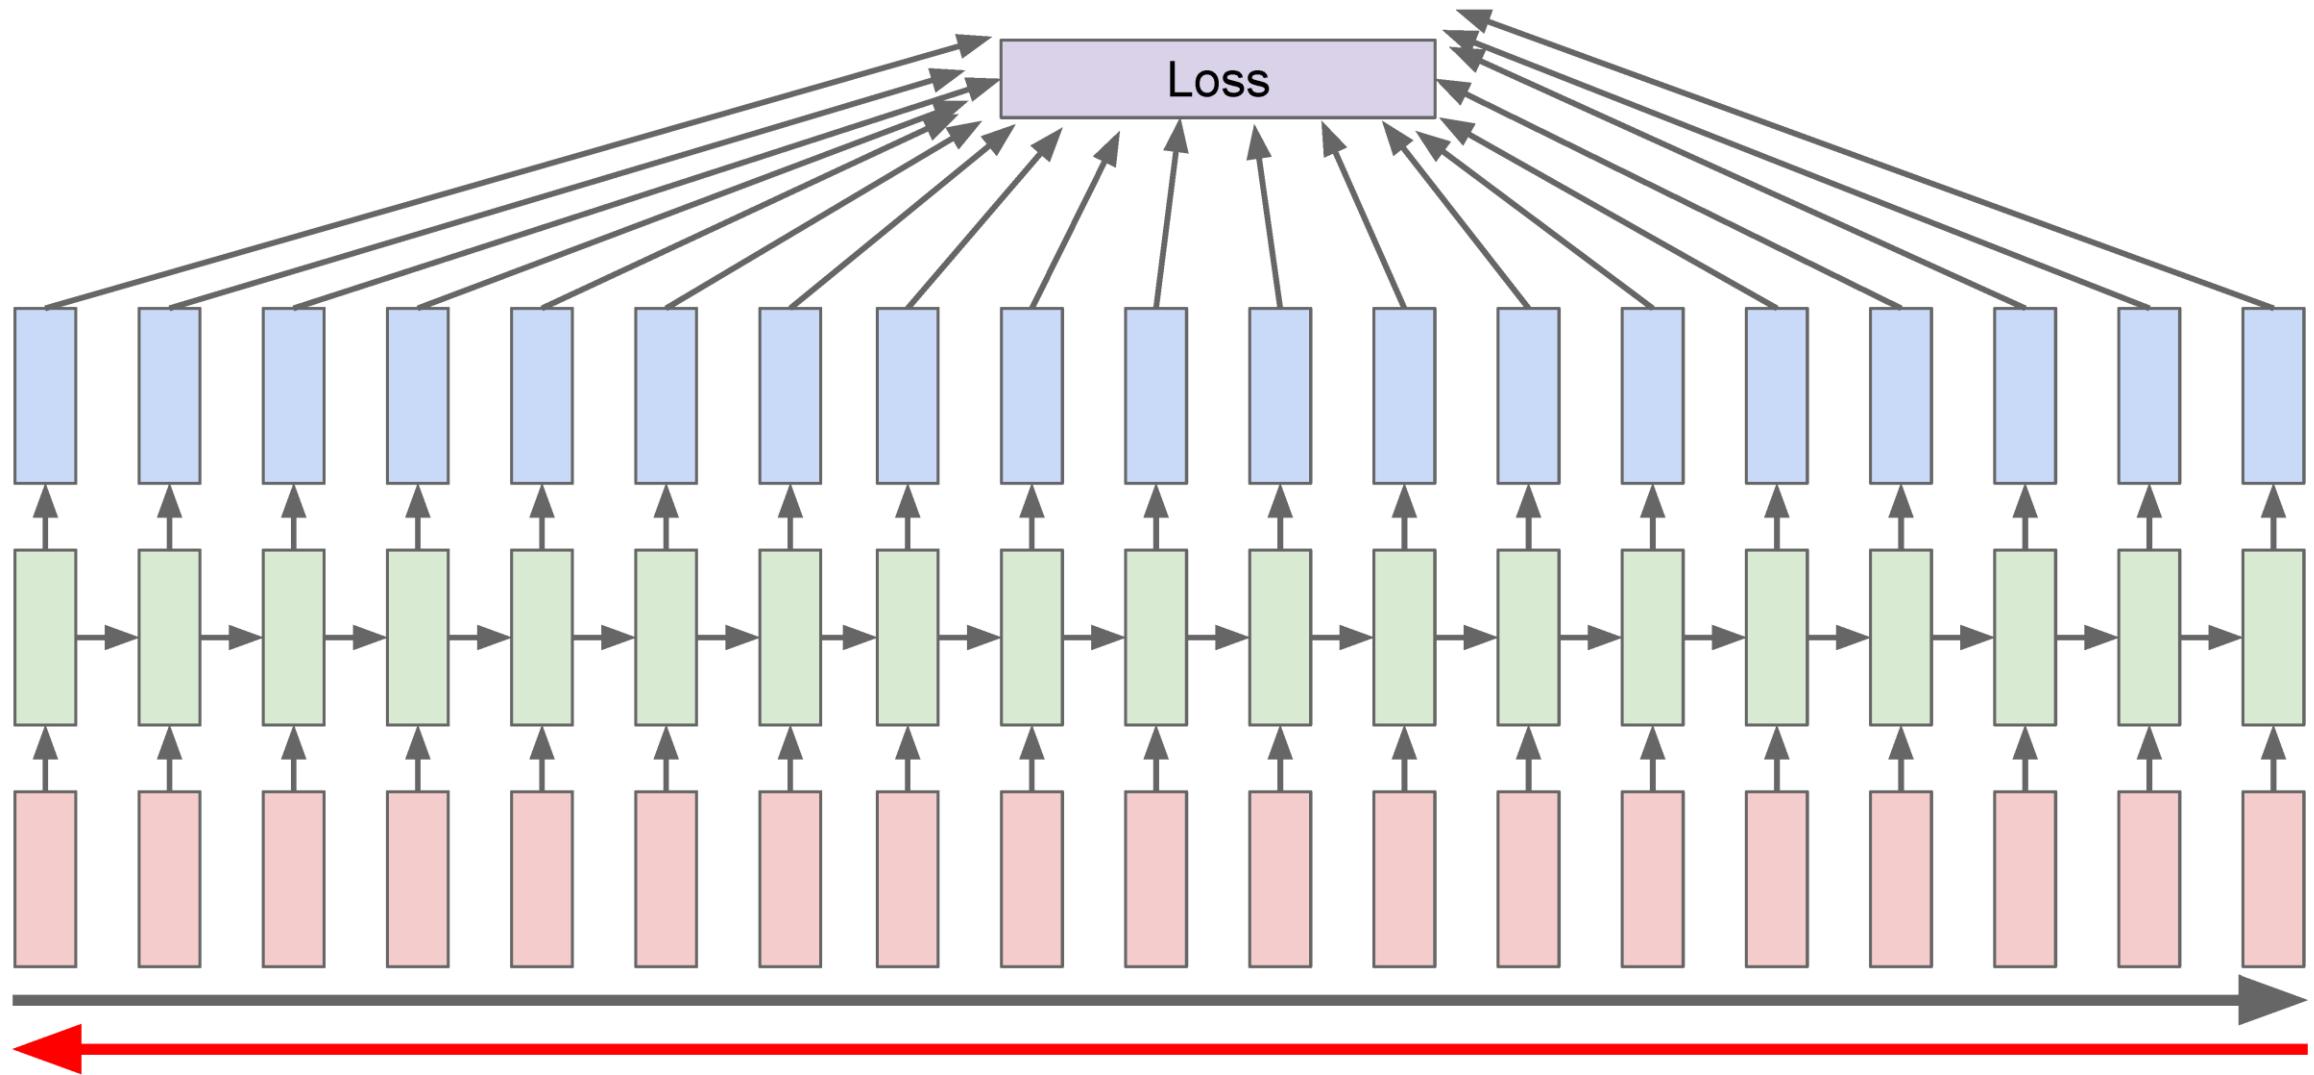
$$\begin{bmatrix} w_{11} & w_{12} & w_{13} & w_{14} \end{bmatrix} [1] \quad [w_{11}] \\ \begin{bmatrix} w_{21} & w_{22} & w_{23} & w_{14} \end{bmatrix} [0] = [w_{21}] \\ \begin{bmatrix} w_{31} & w_{32} & w_{33} & w_{14} \end{bmatrix} [0] \quad [w_{31}] \\ [0] \end{math>$$

Matrix multiply with a one-hot vector just extracts a column from the weight matrix. We often put a separate **embedding** layer between input and hidden layers.

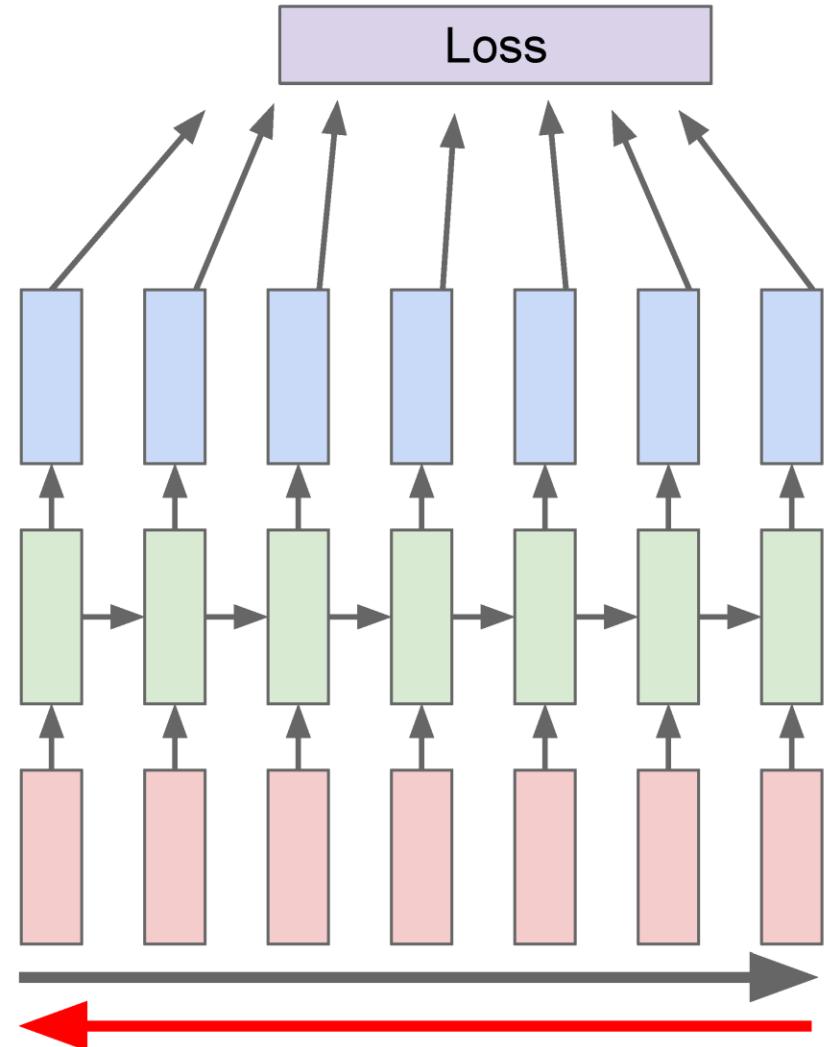


Backpropagation through time

Forward through entire sequence to compute loss, then backward through entire sequence to compute gradient

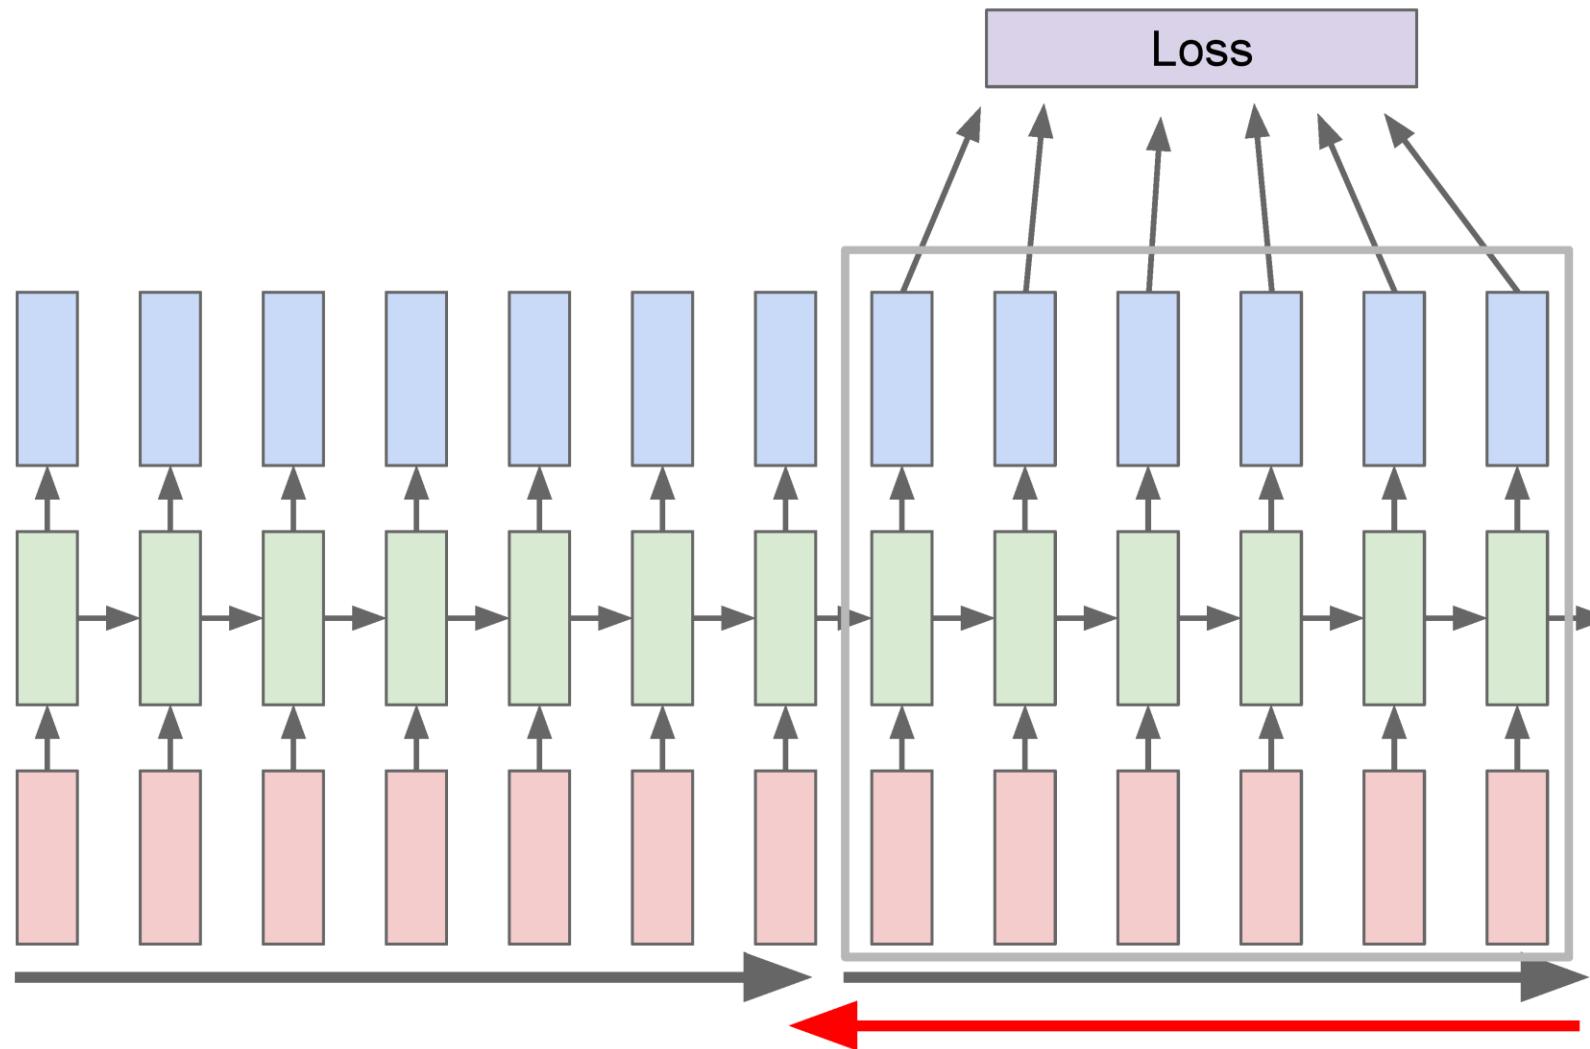


Truncated Backpropagation through time



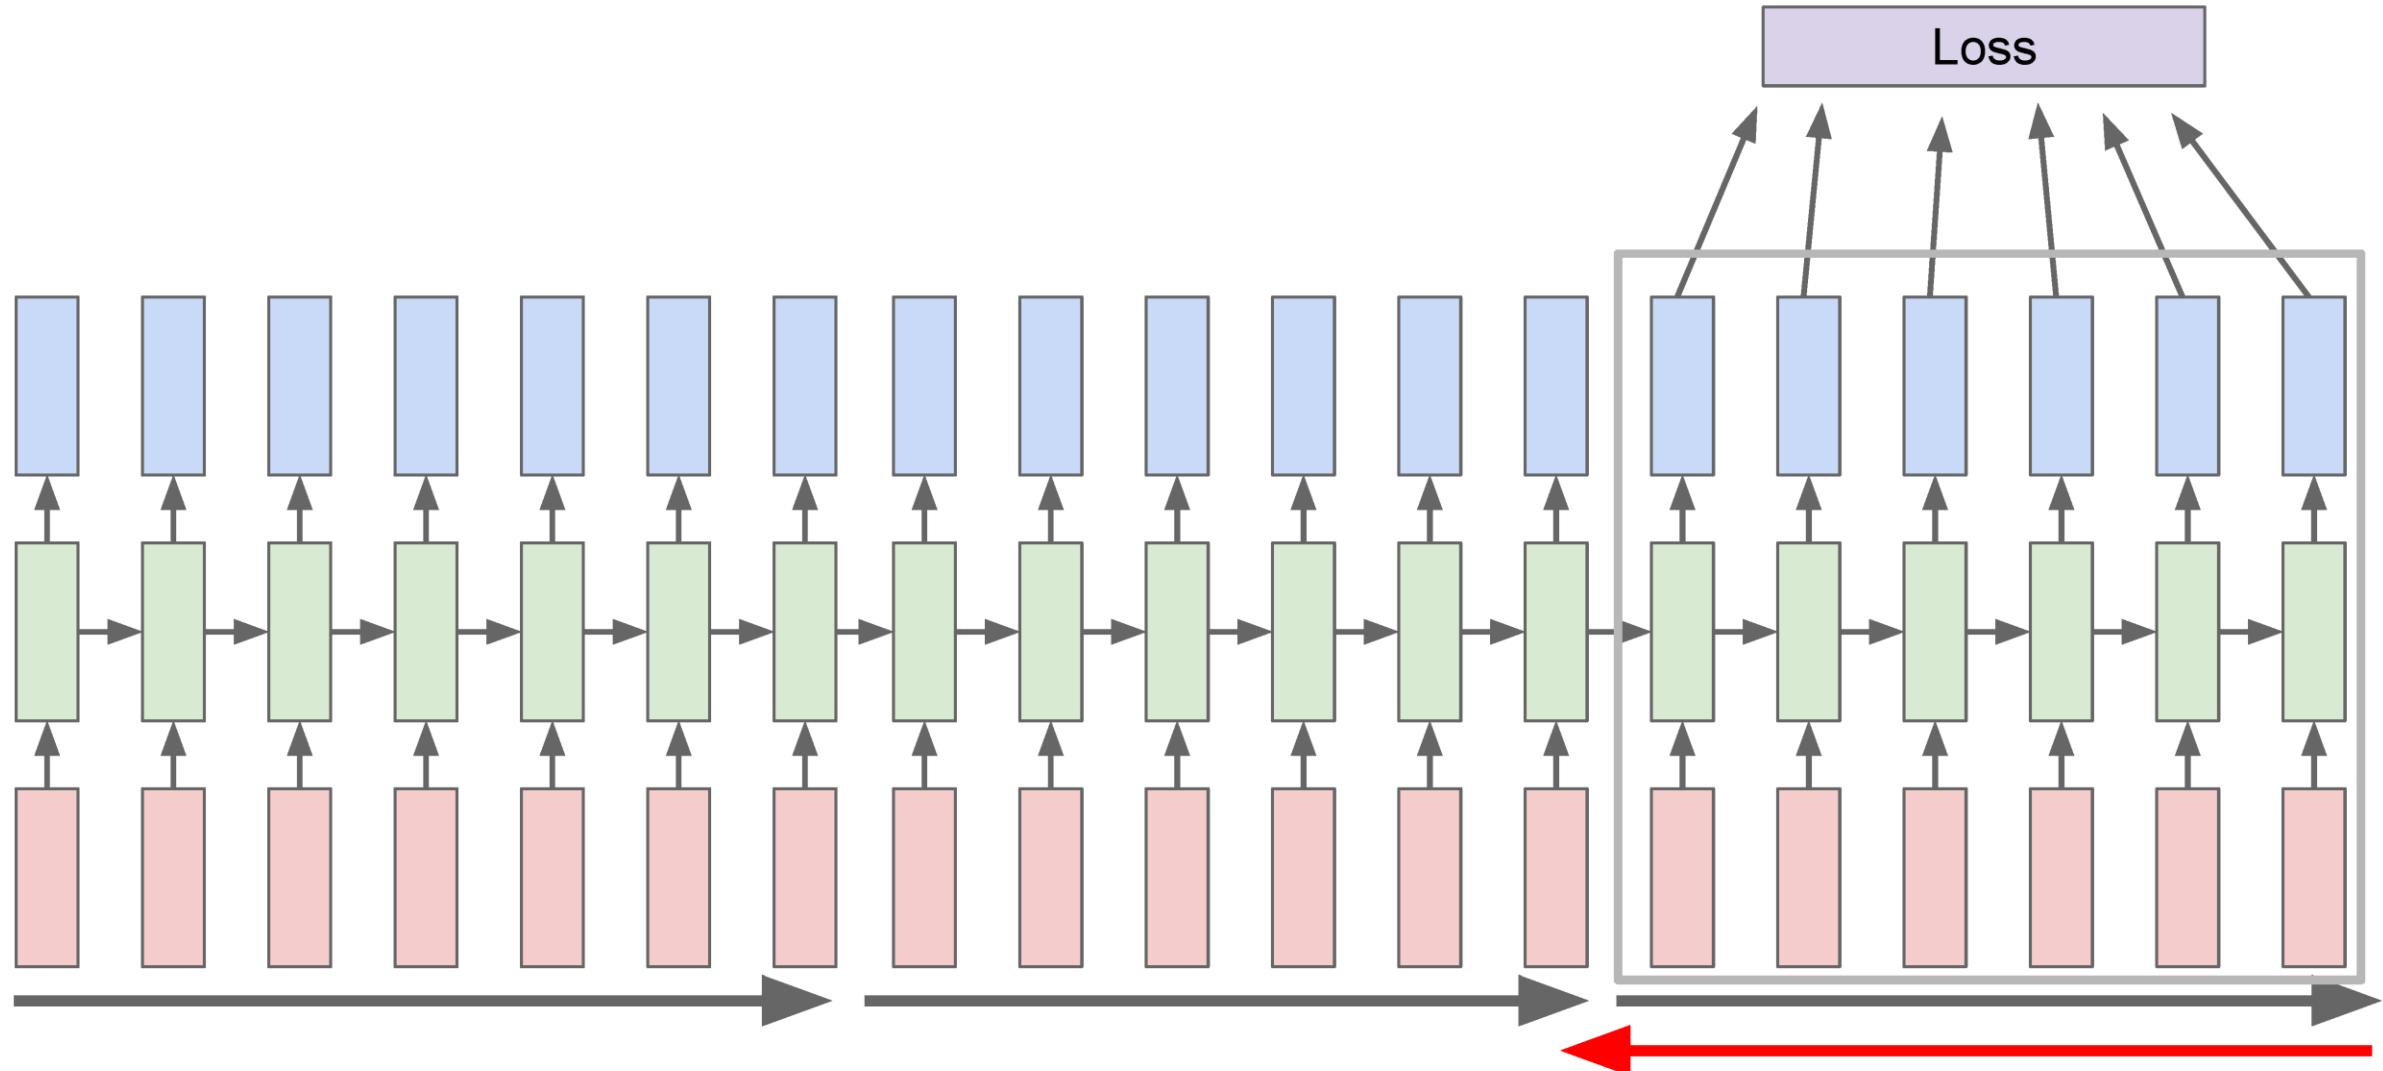
Run forward and backward
through chunks of the
sequence instead of whole
sequence

Truncated Backpropagation through time



Carry hidden states forward in time forever, but only backpropagate for some smaller number of steps

Truncated Backpropagation through time



Minimal character-level language model with a Vanilla Recurrent Neural Network

`min-char-rnn.py`

```

1 """
2 Minimal character-level Vanilla RNN model. Written
3 under BSD License
4 """
5 import numpy as np
6
7 # data I/O
8 data = open('input.txt', 'r').read() # should be one line of text
9 chars = list(set(data))
10 data_size, vocab_size = len(data), len(chars)
11 print 'data has %d characters, %d unique.' % (data_size, vocab_size)
12 char_to_ix = { ch:i for i,ch in enumerate(chars) }
13 ix_to_char = { i:ch for i,ch in enumerate(chars) }
14
15 # hyperparameters
16 hidden_size = 100 # size of hidden layer of neurons
17 seq_length = 25 # number of steps to unroll the
18 learning_rate = 1e-1
19
20 # model parameters
21 Wxh = np.random.randn(hidden_size, vocab_size)*0.01
22 Whh = np.random.randn(hidden_size, hidden_size)*0.01
23 Why = np.random.randn(vocab_size, hidden_size)*0.01
24 bh = np.zeros((hidden_size, 1)) # hidden bias
25 by = np.zeros((vocab_size, 1)) # output bias
26
27 def lossFun(inputs, targets, hprev):
28 """
29 inputs,targets are both list of integers.
30
31 hprev is Hx1 array of initial hidden state
32 returns the loss, gradients on model parameter
33 """
34
35 xs, hs, ys, ps = {}, {}, {}, {}
36 hs[-1] = np.copy(hprev)
37 loss = 0
38 # forward pass
39 for t in xrange(len(inputs)):
40     xs[t] = np.zeros((vocab_size,1)) # encode in 1-of-k
41     xs[t][inputs[t]] = 1
42     hs[t] = np.tanh(np.dot(Wxh, xs[t]) + np.dot(Whh, hs[-1]))
43     ys[t] = np.dot(Why, hs[t]) + by # unnormalized log p
44     ps[t] = np.exp(ys[t]) / np.sum(np.exp(ys[t])) # probabilities
45     loss += -np.log(ps[t][targets[t],0]) # softmax (cross-entropy loss)
46     # backward pass: compute gradients going backwards
47     dWxh, dWhh, dWhy = np.zeros_like(Wxh), np.zeros_like(Whh),
48     dbh, dby = np.zeros_like(bh), np.zeros_like(by)
49     dhnext = np.zeros_like(hs[0])
50     for t in reversed(xrange(len(inputs))):
51         dy = np.copy(ps[t])
52         dy[targets[t]] -= 1 # backprop into y. see http://cs231n.github.io/rnn-tutorial/part1/
53         dWhy += np.dot(dy, hs[t].T)
54         dby += dy
55         dh = np.dot(Why.T, dy) + dhnext # backprop into h
56         ddraw = (1 - hs[t] * hs[t]) * dh # backprop through tanh nonlinearity
57         dbh += ddraw
58         dWxh += np.dot(ddraw, xs[t].T)
59         dWhh += np.dot(ddraw, hs[t-1].T)
60         dhnext = np.dot(Why.T, ddraw)
61     for dparam in [dWxh, dWhh, dWhy, dbh, dby]:
62         np.clip(dparam, -5, 5, out=dparam) # clip to mitigate exploding gradients
63     return loss, dWxh, dWhh, dWhy, dbh, dby, hs[len(inputs):]
64
65 def sample(h, seed_ix, n):
66 """
67 sample a sequence of integers from the model
68 h is memory state, seed_ix is seed letter for first time step
69 """
70
71     x = np.zeros((vocab_size, 1))
72     x[seed_ix] = 1
73     ixes = []
74     for t in xrange(n):
75         h = np.tanh(np.dot(Wxh, x) + np.dot(Whh, h) + bh)
76         y = np.dot(Why, h) + by
77         p = np.exp(y) / np.sum(np.exp(y))
78         ix = np.random.choice(range(vocab_size), p=p.ravel())
79         x = np.zeros((vocab_size, 1))
80         x[ix] = 1
81         ixes.append(ix)
82     return ixes

```

```

81 n, p = 0, 0
82 mWxh, mWhh, mWhy = np.zeros_like(Wxh), np.zeros_like(Whh), np.zeros_like(Why)
83 mbh, mby = np.zeros_like(bh), np.zeros_like(by) # memory variables for Adagrad
84 smooth_loss = -np.log(1.0/vocab_size)*seq_length # loss at iteration 0
85 while True:
86     # prepare inputs (we're sweeping from left to right in steps seq_length long)
87     if p+seq_length+1 >= len(data) or n == 0:
88         hprev = np.zeros((hidden_size,1)) # reset RNN memory
89         p = 0 # go from start of data
90     inputs = [char_to_ix[ch] for ch in data[p:p+seq_length]]
91     targets = [char_to_ix[ch] for ch in data[p+1:p+seq_length+1]]
92
93     # sample from the model now and then
94     if n % 100 == 0:
95         sample_ix = sample(hprev, inputs[0], 200)
96         txt = ''.join(ix_to_char[ix] for ix in sample_ix)
97         print '----\n%s\n----' % (txt, )
98
99     # forward seq_length characters through the net and fetch gradient
100 loss, dWxh, dWhh, dWhy, dbh, dby, hprev = lossFun(inputs, targets, hprev)
101 smooth_loss = smooth_loss * 0.999 + loss * 0.001
102 if n % 100 == 0: print 'iter %d, loss: %f' % (n, smooth_loss) # print progress
103
104 # perform parameter update with Adagrad
105 for param, dparam, mem in zip([Wxh, Whh, Why, bh, by],
106                               [dWxh, dWhh, dWhy, dbh, dby],
107                               [mWxh, mWhh, mWhy, mbh, mby]):
108     mem += dparam * dparam
109     param += -learning_rate * dparam / np.sqrt(mem + 1e-8) # adagrad update
110
111     p += seq_length # move data pointer
112     n += 1 # iteration counter

```

<https://gist.github.com/karpathy/d4dee566867f8291f086>

Example: Image Captioning

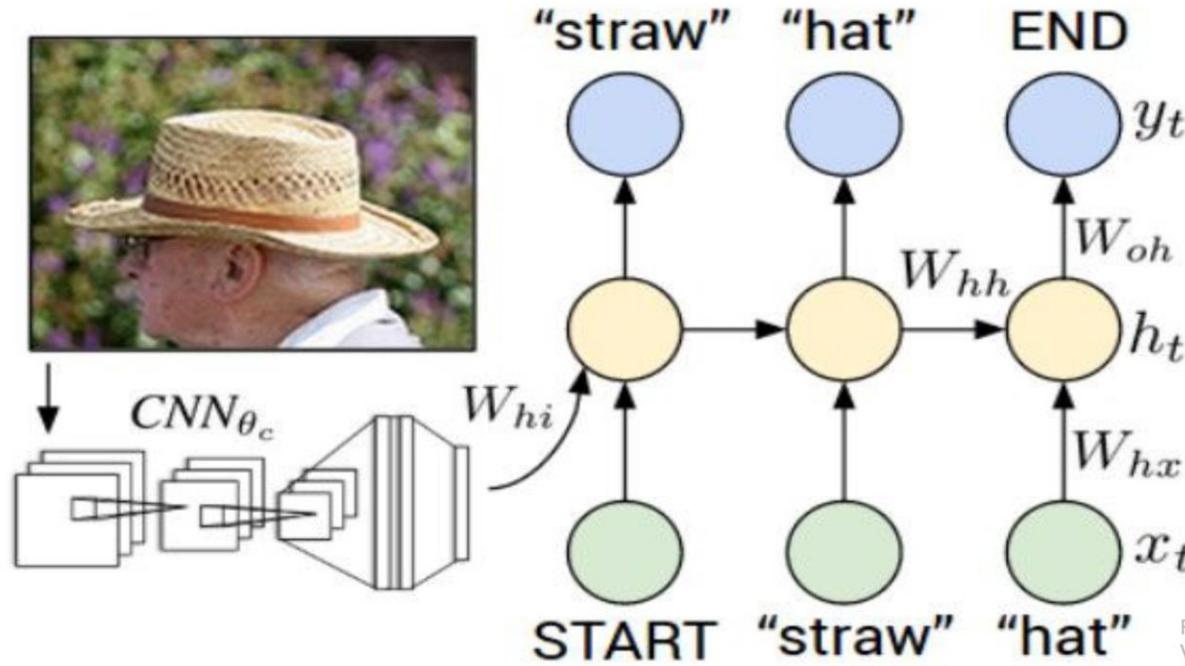
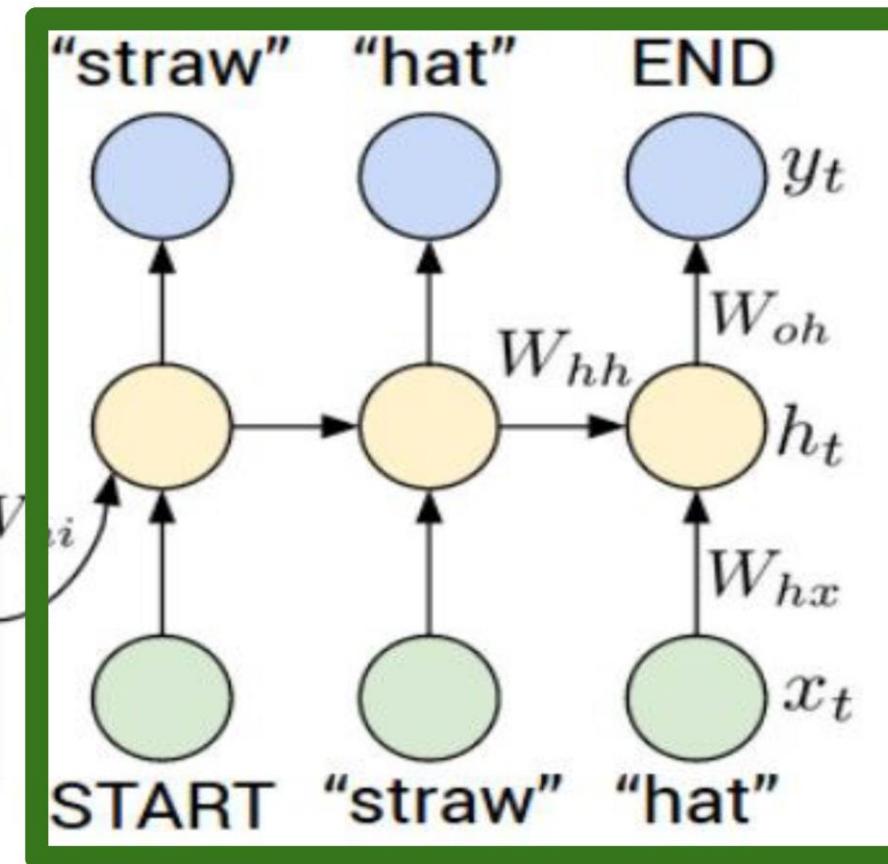
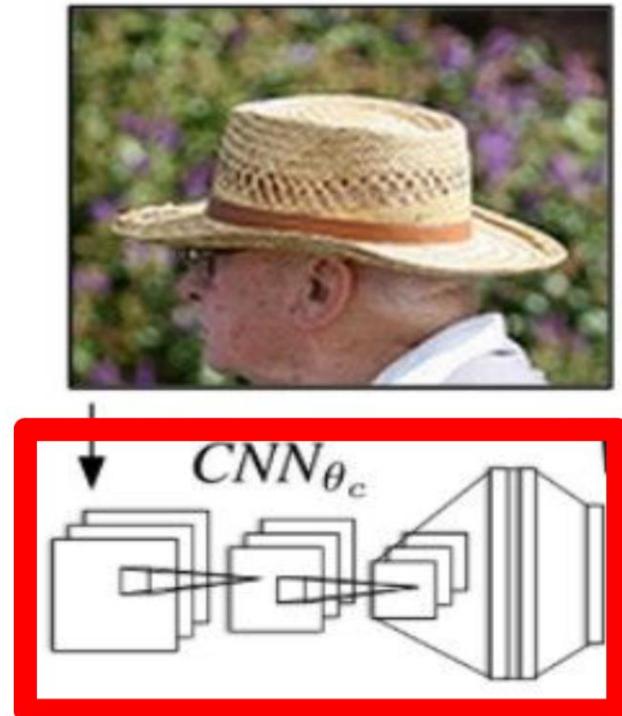


Figure from Karpathy et al, "Deep Visual-Semantic Alignments for Generating Image Descriptions", CVPR 2015; figure copyright IEEE, 2015.
Reproduced for educational purposes.

Recurrent Neural Network



Convolutional Neural Network

test image



[This image](#) is CC0 public domain



test image



test image



image

conv-64

conv-64

maxpool

conv-128

conv-128

maxpool

conv-256

conv-256

maxpool

conv-512

conv-512

maxpool

conv-512

conv-512

maxpool

FC-4096

FC-4096

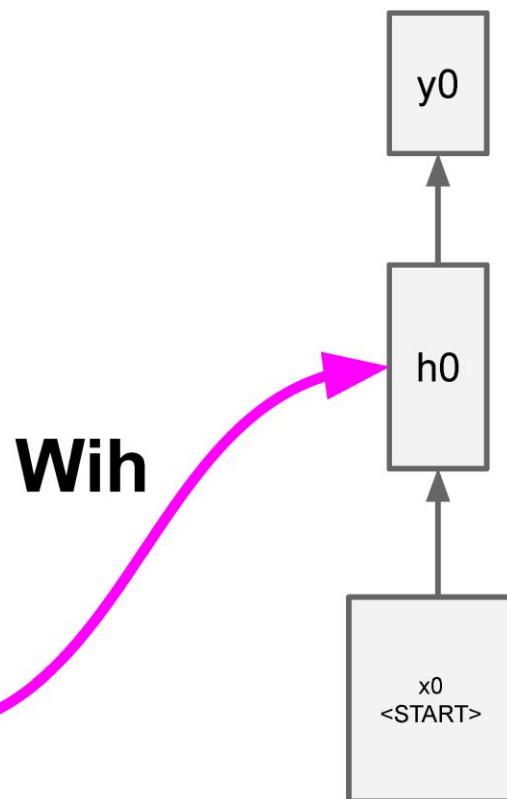


test image

x0
<START>



test image



before:

$$h = \tanh(W_{xh} * x + W_{hh} * h)$$

now:

$$h = \tanh(W_{xh} * x + W_{hh} * h + W_{ih} * v)$$

image



conv-64

conv-64

maxpool

conv-128

conv-128

maxpool

conv-256

conv-256

maxpool

conv-512

conv-512

maxpool

conv-512

conv-512

maxpool

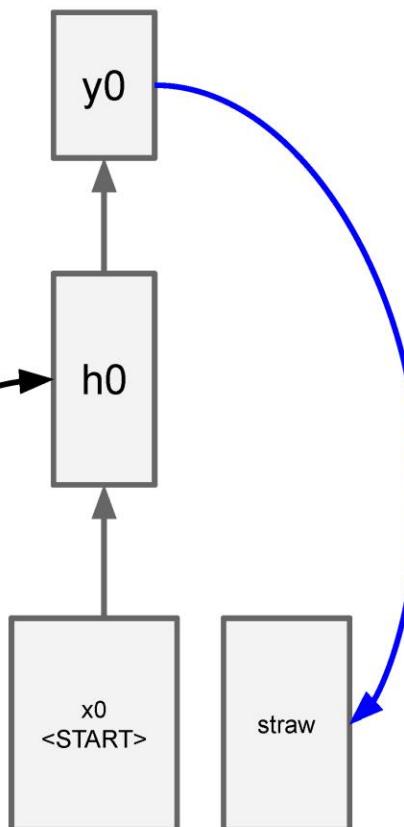
FC-4096

FC-4096



test image

sample!



image

conv-64

conv-64

maxpool

conv-128

conv-128

maxpool

conv-256

conv-256

maxpool

conv-512

conv-512

maxpool

conv-512

conv-512

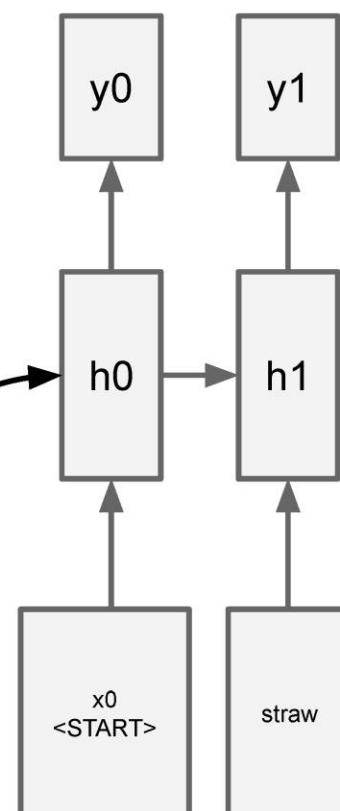
maxpool

FC-4096

FC-4096

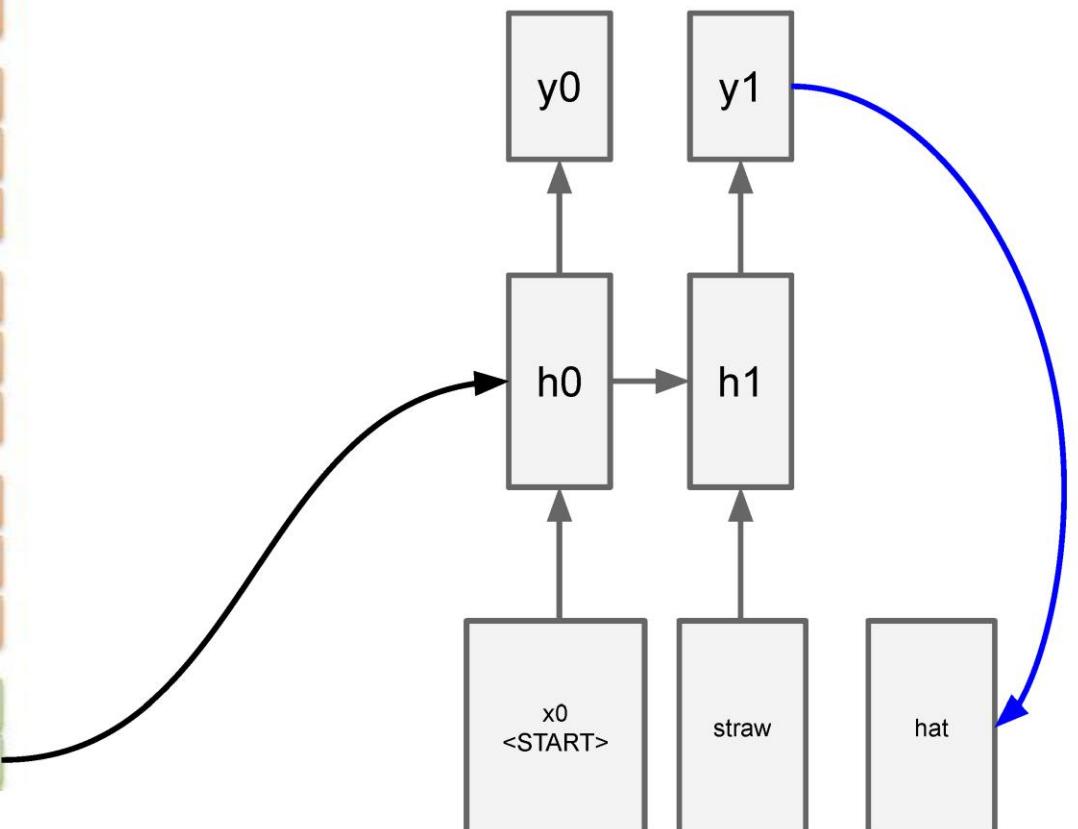


test image





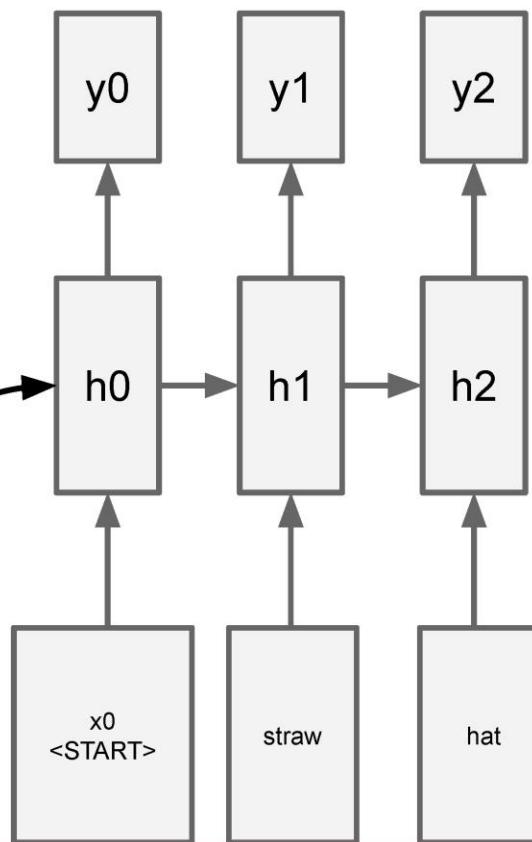
test image



sample!

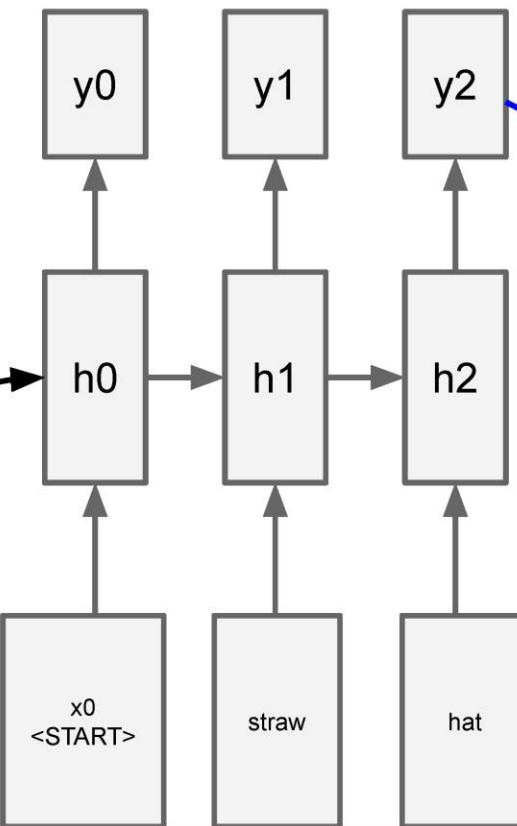


test image





test image



sample
<END> token
=> finish.

Image Captioning: Example Results



A cat sitting on a suitcase on the floor



A cat is sitting on a tree branch



A dog is running in the grass with a frisbee



A white teddy bear sitting in the grass



Two people walking on the beach with surfboards



A tennis player in action on the court



Two giraffes standing in a grassy field

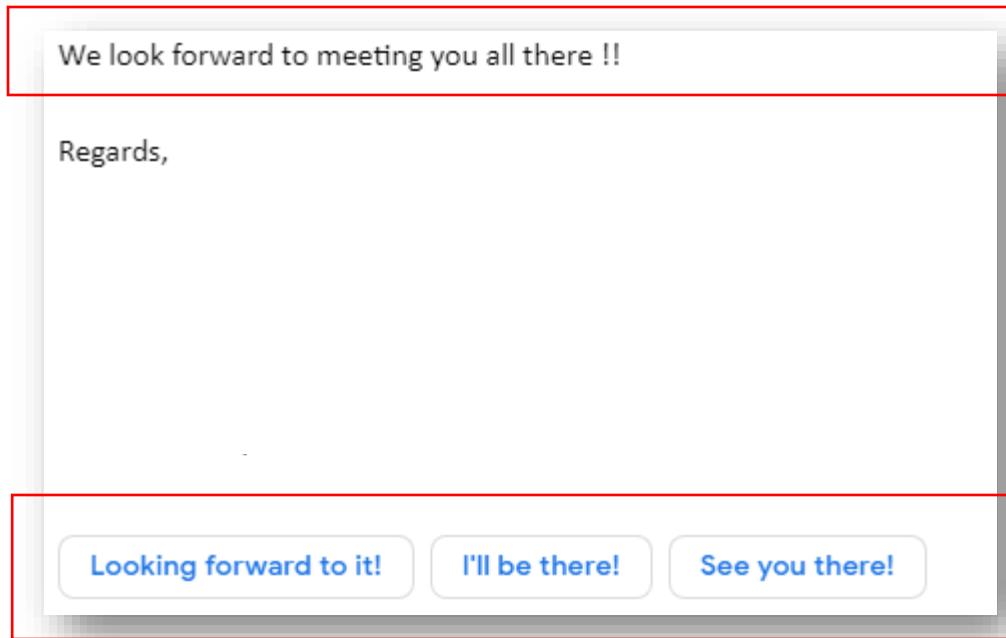


A man riding a dirt bike on a dirt track

Captions generated using [neuraltalk2](#)
All images are [CC0 Public domain](#)
[cat](#) [suitcase](#) [cat](#) [tree](#) [dog](#) [bear](#)
[surfers](#) [tennis](#) [giraffe](#) [motorcycle](#)

Example: Generating Responses to Support Fact-checkers

- Text Generation is very popular.



Responses generated by Gmail

- Users can select a response and edit it.
 - Faster typing
 - Increase users' engagement

Generating Responses to Support Fact-checkers

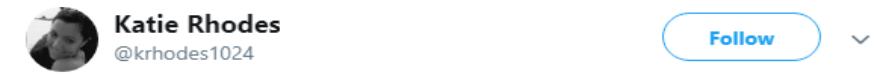
- Goal:
 - Given content of an original tweet:
 - Generating responses with **fact-checking intention**
 - Fact-checking intention:
 - Refuting
 - Supporting

Refuting



Study what General Pershing of the United States did to terrorists when caught. There was no more Radical Islamic Terror for 35 years!

11:45 AM - 17 Aug 2017



Replying to @realDonaldTrump

Stop spreading fake news. This Pershing story has long been debunked!
<http://www.politifact.com/truth-o-meter/statements/2017/aug/17/donald-trump/donald-trump-retells-pants-fire-claim-about-gen-pe>

1:20 PM - 17 Aug 2017

Generating Responses to Support Fact-checkers

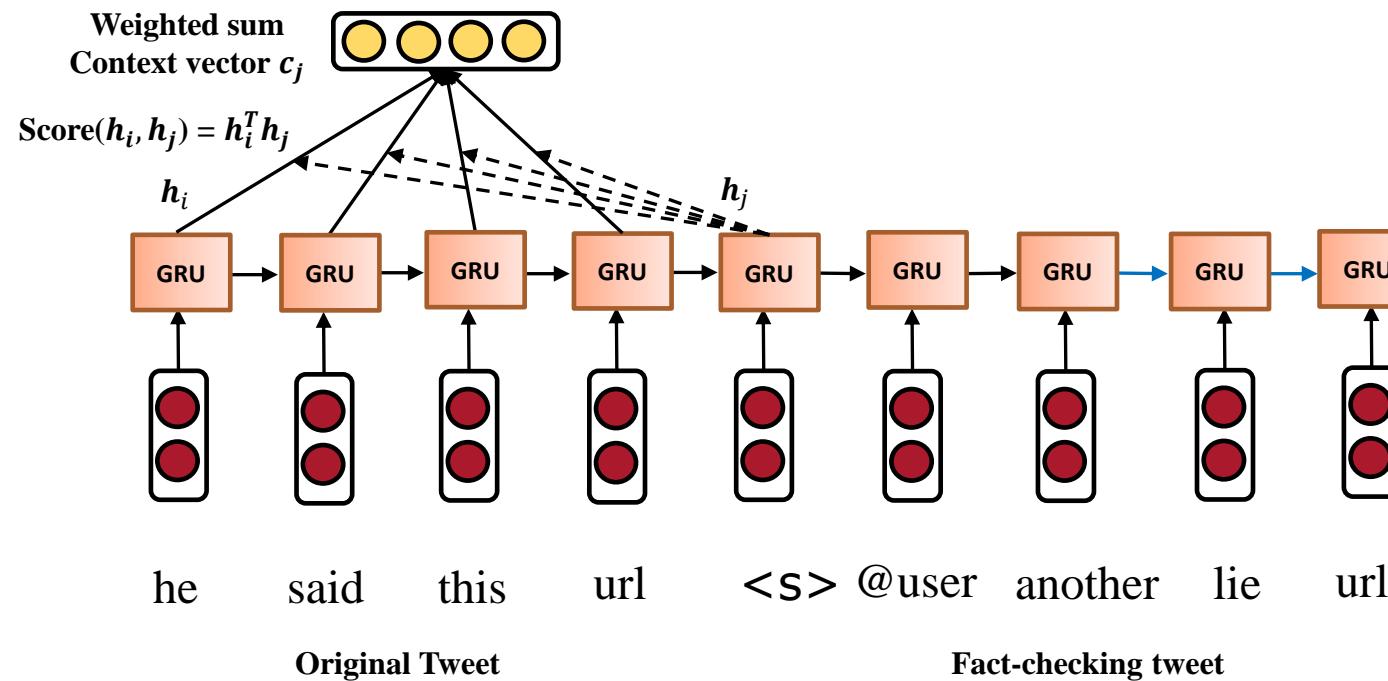
Recap: An example of fact-checking activity

The diagram shows a sequence of three tweets on a platform with user profiles, follow buttons, and a dropdown menu.

- Original tweet:** A tweet from Donald J. Trump (@realDonaldTrump) dated 11:45 AM - 17 Aug 2017. The text reads: "Study what General Pershing of the United States did to terrorists when caught. There was no more Radical Islamic Terror for 35 years!" A red arrow points from the text to the label "Original tweet".
- An online fact-checker:** A reply from Katie Rhodes (@krhodes1024) dated 1:20 PM - 17 Aug 2017, replying to @realDonaldTrump. The text reads: "Stop spreading fake news. This Pershing story has long been debunked! <http://www.politifact.com/truth-o-meter/statements/2017/aug/17/donald-trump/donald-trump-retells-pants-fire-claim-about-gen-pe>". A red arrow points from the text to the label "An online fact-checker".
- A fact-checking tweet (FC-tweet):** The same tweet from Katie Rhodes (@krhodes1024) dated 1:20 PM - 17 Aug 2017, replying to @realDonaldTrump. A red arrow points from the text to the label "A fact-checking tweet (FC-tweet)".

Generating Responses to Support Fact-checkers

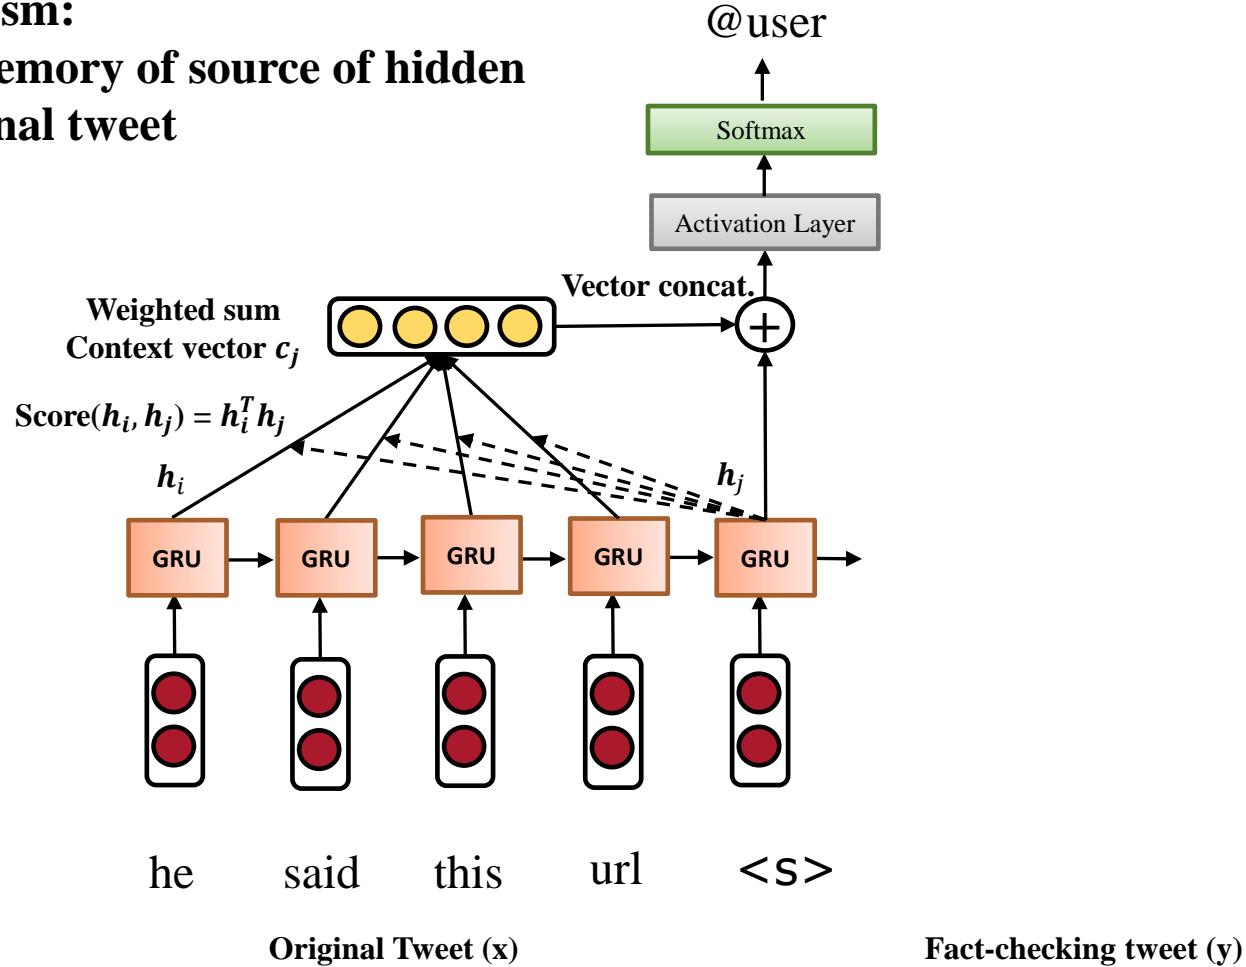
- The framework is based on Sequence to Sequence model with attention mechanism
 - Input: Original tweet
 - Output: Fact-checking tweet
- Attention mechanism:
 - Maintain a memory of source of hidden states from original tweet



Generating Responses to Support Fact-checkers

- Attention mechanism:

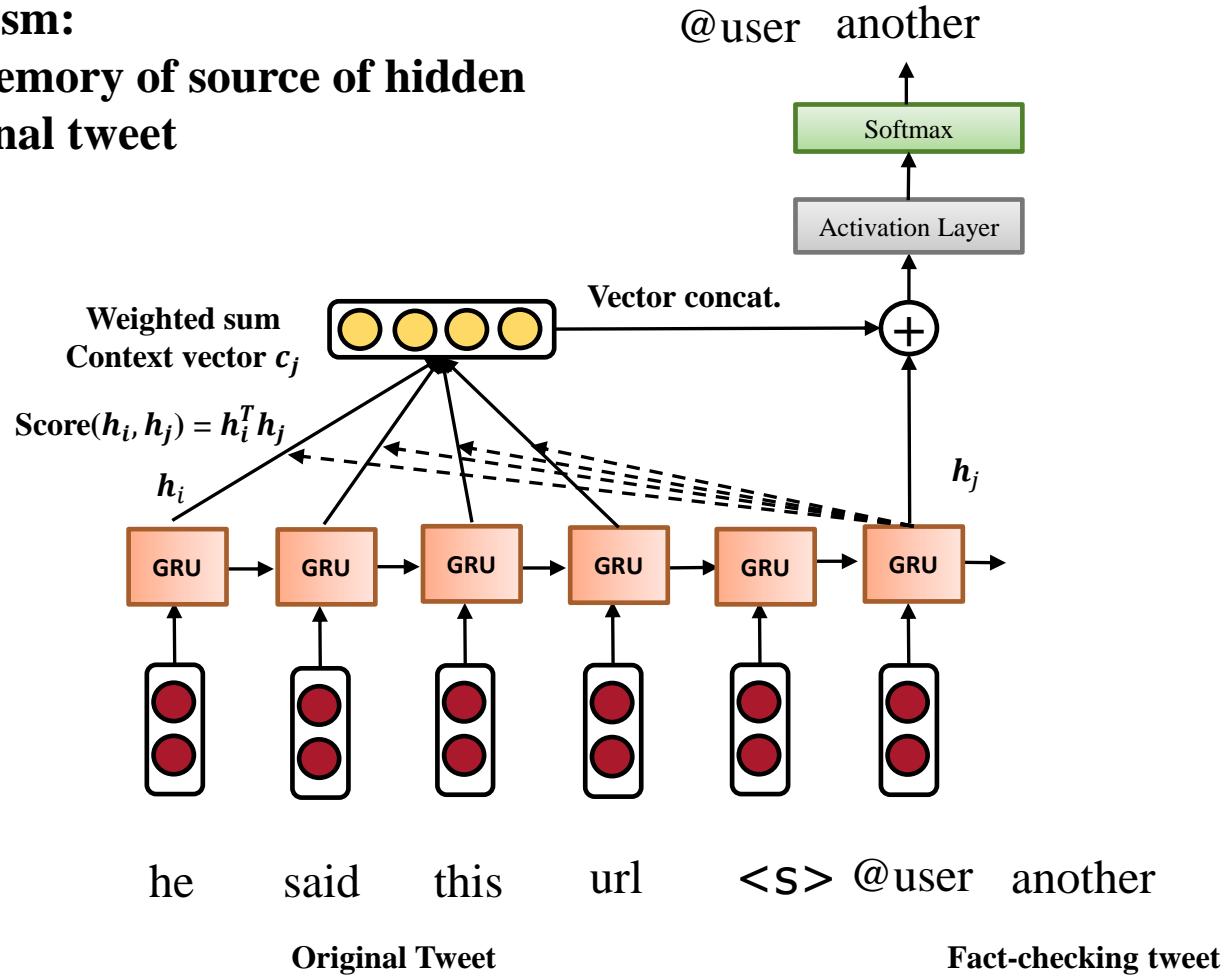
- Maintain a memory of source of hidden states in original tweet



Generating Responses to Support Fact-checkers

- Attention mechanism:

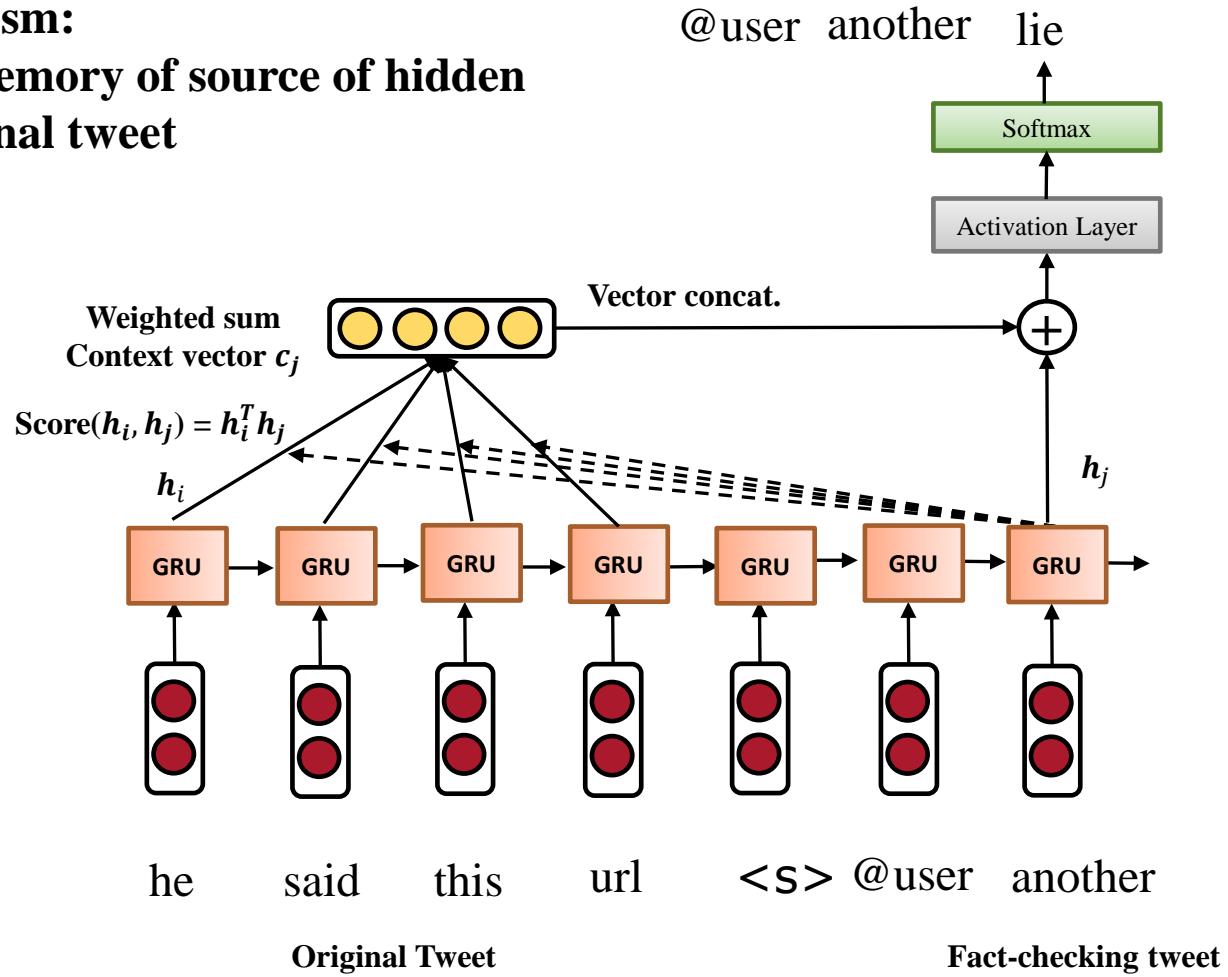
- Maintain a memory of source of hidden states in original tweet



Generating Responses to Support Fact-checkers

- Attention mechanism:

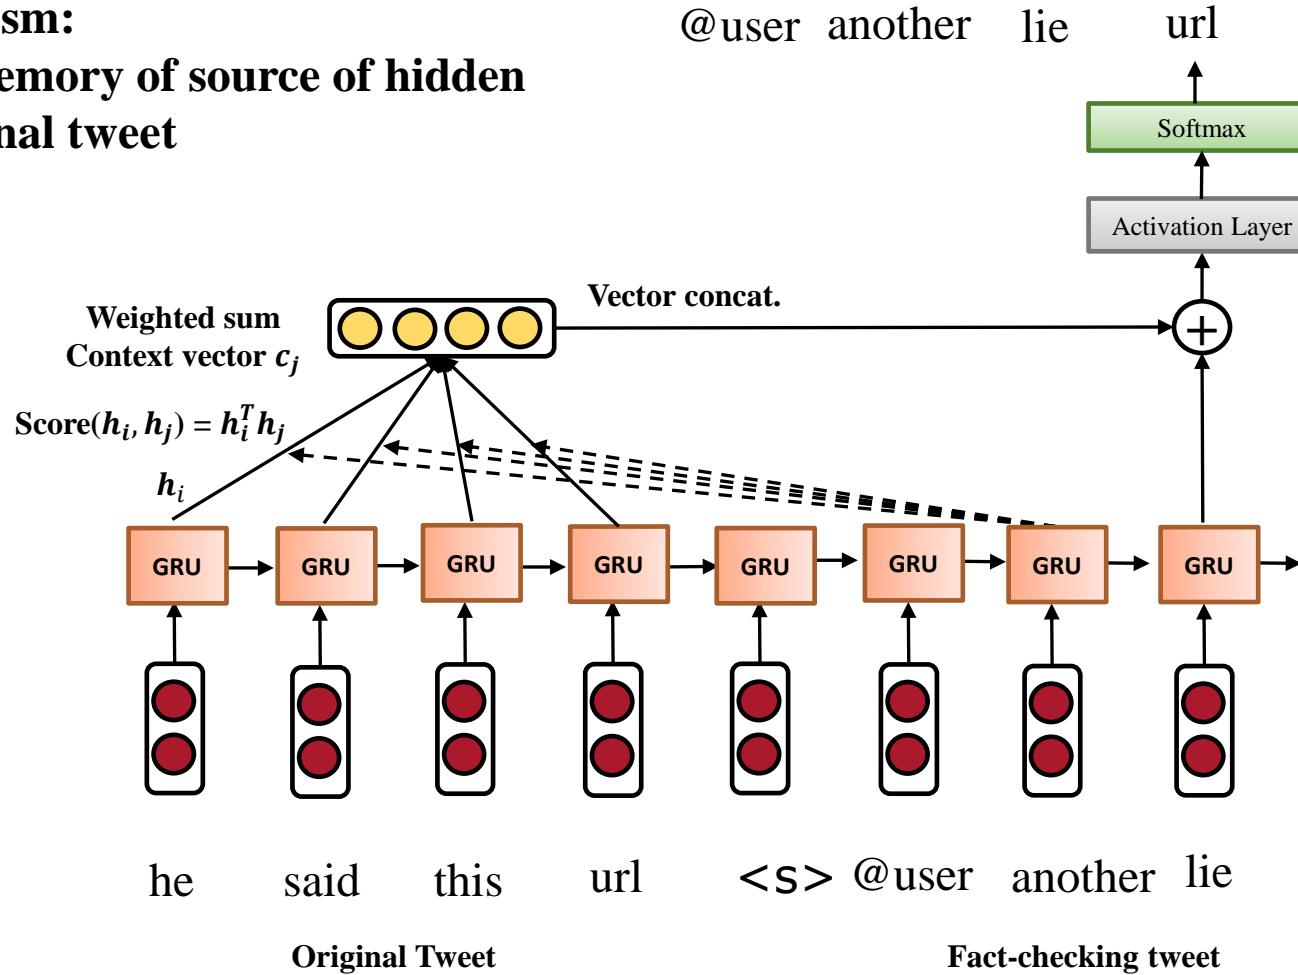
- Maintain a memory of source of hidden states in original tweet



Generating Responses to Support Fact-checkers

- Attention mechanism:

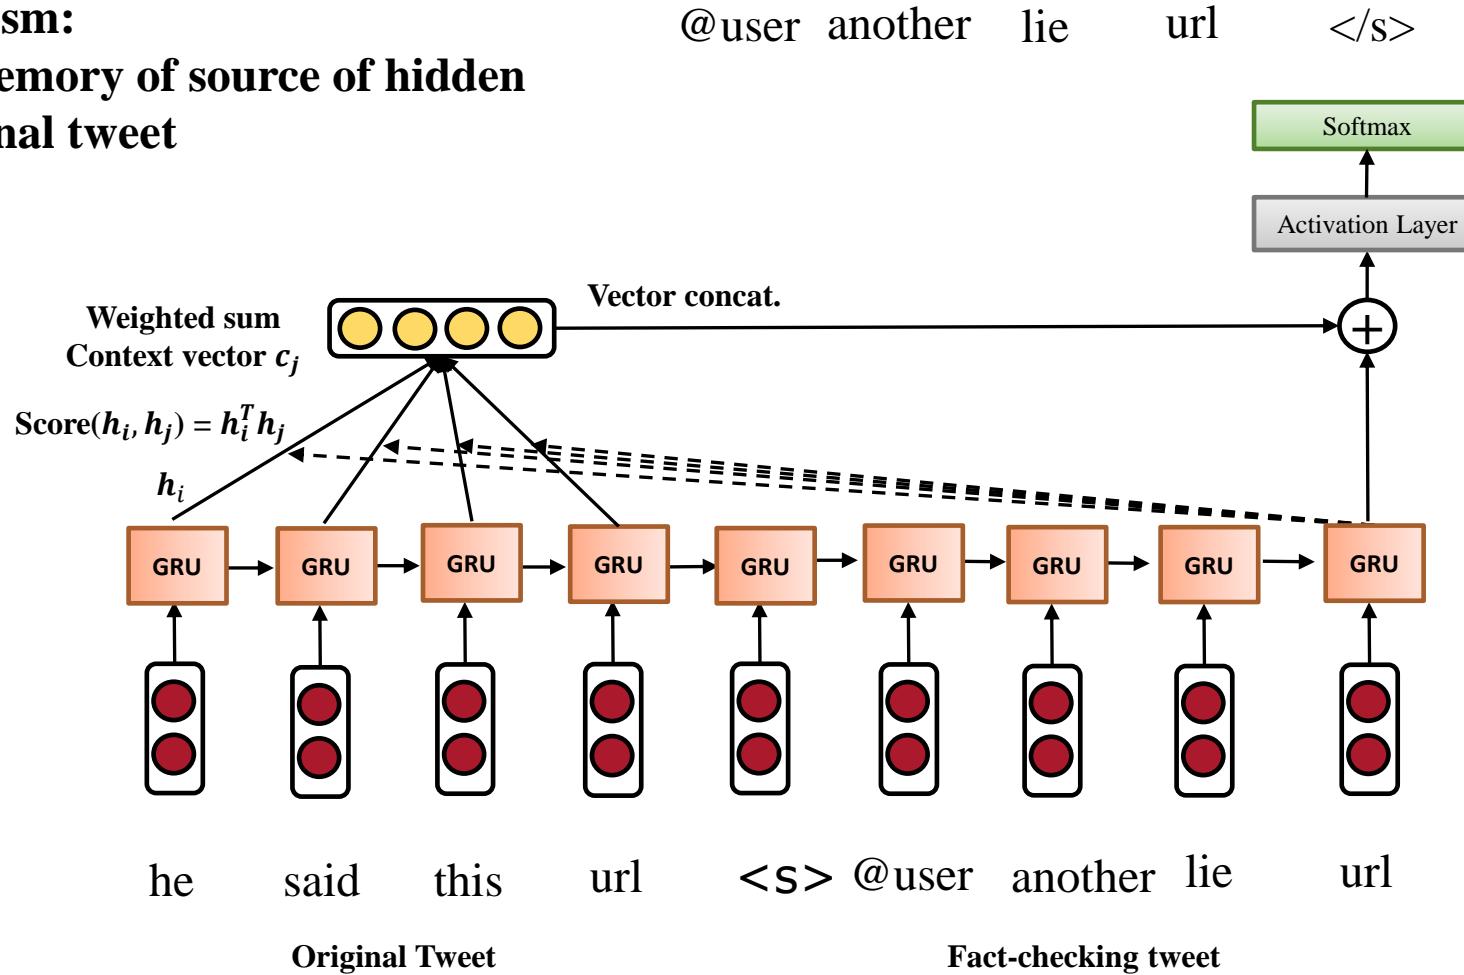
- Maintain a memory of source of hidden states in original tweet



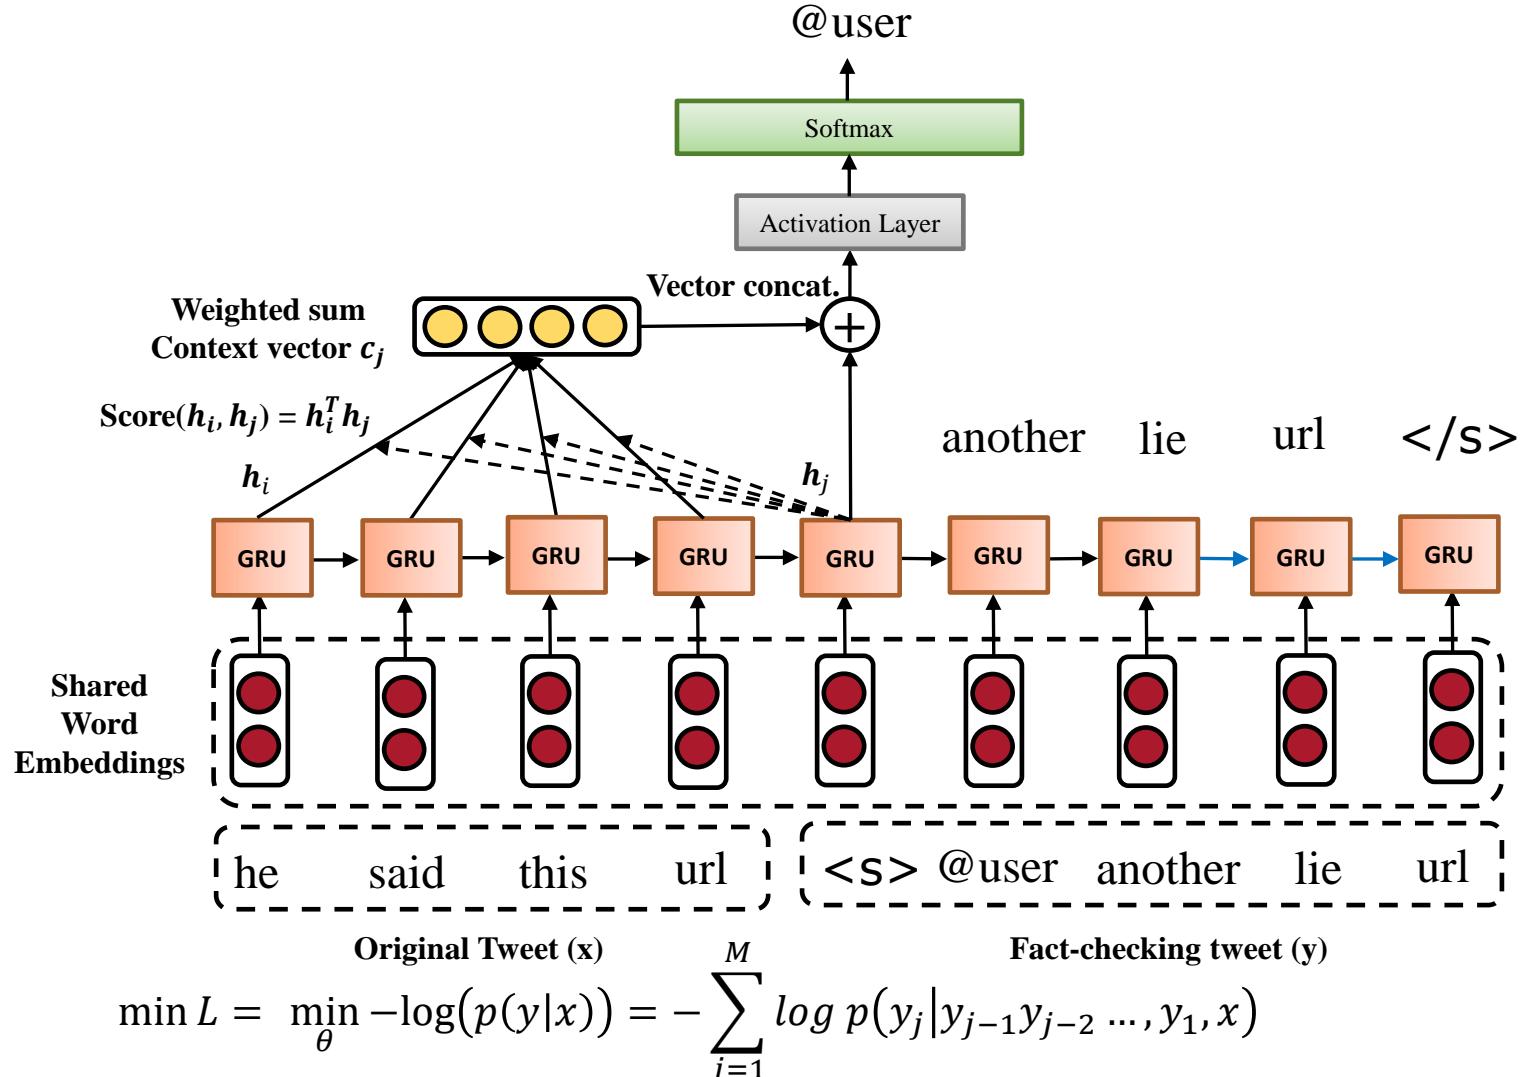
Generating Responses to Support Fact-checkers

- Attention mechanism:

- Maintain a memory of source of hidden states in original tweet



Generating Responses to Support Fact-checkers



What's next?

- Study Deep Learning
 - [Deep Learning](#) by Yoshua Bengio, Ian Goodfellow, and Aaron Courville. (Available online)
 - [Deep Learning with Python](#) by Francois Chollet
 - <http://neuralnetworksanddeeplearning.com/index.html>
 - <https://keras.io/examples/>
 - <https://www.tensorflow.org/tutorials>
 - <https://pytorch.org/tutorials/>
- Take Deep Learning course
 - CS 541. DEEP LEARNING
 - Online course: <https://www.coursera.org/specializations/deep-learning>